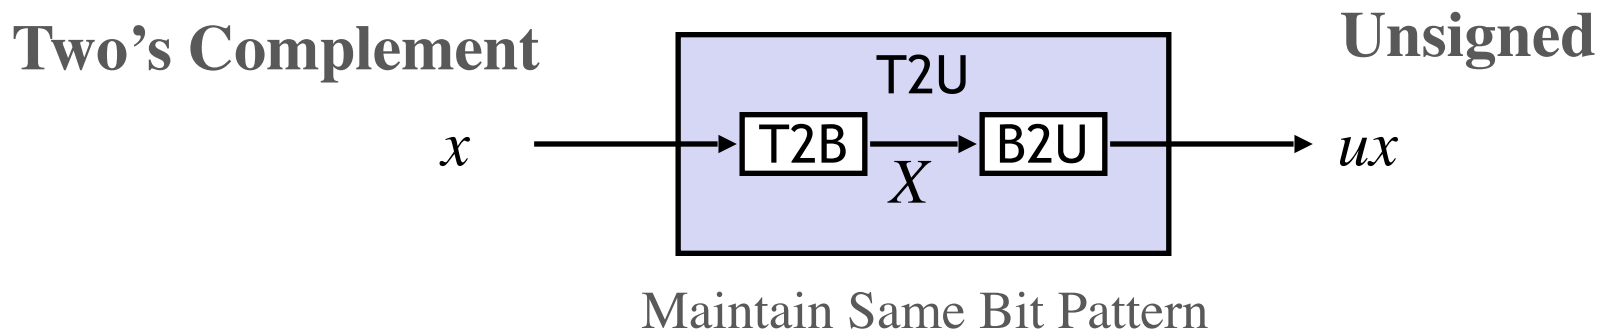




Interpretation of Bit Vectors


+ Mapping signed \leftrightarrow unsigned

- The computer itself has no idea if a given bit pattern at a particular location in memory “signed” or “unsigned”.
- The program interprets some given bit pattern according to the *type* that value has been assigned.
- Moreover, mappings between unsigned and two’s complement numbers keep the same bit representations but are interpreted differently depending on type, *which may yield a different value in your program*.






Mapping signed \leftrightarrow unsigned *con't*



| Bits | Signed | | Unsigned |
|------|--------|-------------------------------------------------------------|----------|
| 0000 | 0 | \longleftrightarrow = \longleftrightarrow | 0 |
| 0001 | 1 | | 1 |
| 0010 | 2 | | 2 |
| 0011 | 3 | | 3 |
| 0100 | 4 | | 4 |
| 0101 | 5 | | 5 |
| 0110 | 6 | | 6 |
| 0111 | 7 | | 7 |
| 1000 | -8 | $\xrightarrow{\text{T2U}}$ $\xleftarrow{\text{U2T}}$ | 8 |
| 1001 | -7 | | 9 |
| 1010 | -6 | | 10 |
| 1011 | -5 | | 11 |
| 1100 | -4 | \longleftrightarrow +/- 2^w \longleftrightarrow | 12 |
| 1101 | -3 | | 13 |
| 1110 | -2 | | 14 |
| 1111 | -1 | | 15 |

+ Insights into overflow




- Lets say you have a signed char with the bit pattern...

01111111

- What is its value in two's complement in decimal? How about unsigned?

+ Insights into overflow



- Lets say you have a signed char with the bit pattern...

01111111


- What is its value in two's complement in decimal? How about unsigned?

t: 127

u: 127

- Lets say 1 is added to 127. What is the bit pattern for 128?

+ Insights into overflow



- Lets say you have a signed char with the bit pattern...

01111111

- What is its value in two's complement in decimal? How about unsigned?

t: 127


u: 127

- Lets say 1 is added to 127. What is the bit pattern for 128?

10000000

- What is this bit pattern's value in two's complement in decimal? How about unsigned?

+ Insights into overflow



- Lets say you have a signed char with the bit pattern...

01111111

- What is its value in two's complement in decimal? How about unsigned?

t: 127

u: 127

- Lets say 1 is added to 127. What is the bit pattern for 128?

10000000

- What is this bit pattern's value in two's complement in decimal? How about unsigned?


t: -128

u: 128

- See *overflow.c*




It's all a matter of interpretation



- The key idea so far here is that a bit pattern is just a bit pattern!!
 - It has no intrinsic value or semantics.
- How that bit pattern is '*interpreted*' determines its value in your program.
- Ok, so how are bit patterns interpreted in programs?



It's all a matter of interpretation




- The key idea so far here is that a bit pattern is just a bit pattern!!
 - It has no intrinsic value or semantics.
- How that bit pattern is '*interpreted*' determines its value in your program.
- Ok, so how are bit patterns interpreted in programs?

Datatypes!



+ Conversion & Casting with Integers

+ Signed vs. unsigned in C



- **Constants**

- By default are considered to be signed integers
- If you want unsigned you must add a “U” suffix

```
unsigned int x = 0U;  
unsigned int y = 4294967259U;
```

- **Casting**

- *Explicit* casting between signed & unsigned


```
int tx, ty;  
unsigned int ux, uy;  
tx = (int) ux;  
uy = (unsigned) ty;
```

- *Implicit* casting also occurs during assignments and function calls

```
tx = ux;  
uy = ty;
```




Casting surprises



- If there is a mix of unsigned and signed in single expression, signed values are *implicitly cast to unsigned*
 - Includes expressions with comparison operators: `<`, `>`, `==`, `<=`, `>=`
 - See *casting_surprise.c*
- There can also be unexpected results when working with array indices
 - See *array_surprise.c* and *array_surprise2.c*




Casting signed \leftrightarrow unsigned: summary



- When the coercion takes place the bit pattern is *maintained*
 - However the program will *reinterpret* its value!
 - Can have unexpected effects if not careful, as we just observed.
- Again, expressions containing signed and unsigned int...
 - signed integral is coerced to an unsigned integral!!



Signed 'extension'



- When we do a 'widening conversion' of a value via casting, what happens?
- In other words, given w -bit signed typed integer value x , convert it to $w+k$ -bit typed integer with same value.
 - w is the number of bits in the type of x
 - ex. short = 16
 - k is the number of bits difference between the two types
 - ex. k of short vs int = 16
- Moreover, what happens in cases like this?


```
short x    = 15213;
int  ix   = (int) x;
short y    = -15213;
int  iy   = (int) y;
```

+ Signed 'extension' *con't*

Solution: make k copies of the sign bit

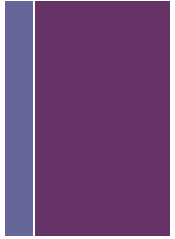
$$X = x_{w-1} x_{w-2} \dots x_1 x_0$$

$$X' = \underbrace{x_{w-1} \dots x_{w-1}}_{\leftarrow k \text{ times} \rightarrow} x_{w-1} x_{w-2} \dots x_1 x_0$$



- Unsigned: zeros added
- Signed: sign bit extension
- Both yield intuitive and expected result

+ Signed 'extension' *con't*




- Therefore, converting from smaller to larger integer data type C *automatically* performs sign extension
- Therefore, this code...

```
short x    = 15213;  
int  ix    = (int) x;  
short y    = -15213;  
int  iy    = (int) y;
```

- ...has the values....

| | Decimal | Hex | Binary |
|-----------|---------|-------------|-------------------------------------|
| x | 15213 | 3B 6D | 00111011 01101101 |
| ix | 15213 | 00 00 3B 6D | 00000000 00000000 00111011 01101101 |
| y | -15213 | C4 93 | 11000100 10010011 |
| iy | -15213 | FF FF C4 93 | 11111111 11111111 11000100 10010011 |

+ Truncation



- When we do a ‘narrowing conversion’ of a value via coercion or casting, what happens? (i.e. from 32-bit int to 16-bit short)
- Higher-order bits are *truncated*. Value is altered, will be reinterpreted.
- Might yield reasonable result if value is ‘small enough’ to fit in smaller type...


```
int i = 1;  
short s = (short) i;
```

- But what about something like this?

```
short s = 256;  
char c = (char) s;
```

- This non-intuitive behavior can lead to buggy code!
- See *coercion.c*

+ Summary




- **Extension** (e.g. short to int)
 - Unsigned: zeroes added
 - Signed: sign extension
 - Both yield expected results
- **Truncation** (e.g. unsigned short to unsigned int)
 - Unsigned/signed: Higher weighted bits are lopped off
 - Result must be reinterpreted
 - For ‘small numbers’ (e.g. int w/ value 16 into short), ok
 - For ‘large numbers’ (e.g. int w/ value 2^{20} into short), problematic.



+

Negation & Addition

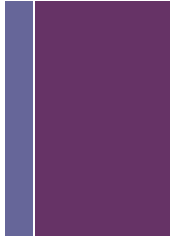
+ Negation



- **Task:** given a bit-vector X compute $-X$
- **Solution:** $-X = \sim X + 1$
 - Negating a value is done by computing its complement and adding 1
- **Example:**
$$X = 011001_2 = 25_{10}$$
$$\sim X = 100110_2 = -26_{10}$$
$$\sim X + 1 = 100111_2 = -25_{10}$$
- Notice, therefore, that for any signed integral type x , $\sim x + x = -1$
 - See *negation.c*



Addition in base 2



- Very simple, works same as base 10, just remember..

- $0 + 0 = 0$
- $0 + 1 = 1$
- $1 + 0 = 1$
- $1 + 1 = 10$


- Examples:

$$\begin{array}{r} 101 \\ +101 \\ \hline 1010 \end{array}$$

$$\begin{array}{r} 1011 \\ +1011 \\ \hline 10110 \end{array}$$

- Note in the second example that in the 2^1 column, we have $1 + (1 + 1)$, where the first 1 is "carried" from the 2^0 column.


+ Unsigned addition




Operands: w bits



True Sum: $w+1$ bits



Discard Carry: w bits



- However since types have a limited number of bits, any carry bits after the MSB simply get truncated.

$$\begin{array}{rcl} 10010_2 & = & 18_{10} \\ + 11011_2 & = & 27_{10} \\ \hline 101101_2 & = & 45_{10} \\ \hline 01101_2 & = & 13_{10} \end{array}$$

- See *unsigned_addition_overflow.c*

+ Signed addition

Operands: w bits


u 

+ v 

True Sum: $w+1$ bits

$u + v$ 

Discard Carry: w bits


$\text{TAdd}_w(u, v)$ 

- **TAdd** and **UAdd** have identical bit-level behavior. If true sum requires $w+1$ bits, any carry bits after the MSB simply get truncated.

$$\begin{array}{rcl}
 10010_2 & = & -14_{10} \\
 + 11011_2 & = & -5_{10} \\
 \hline
 101101_2 & = & -19_{10} \\
 \hline
 01101_2 & = & 13_{10}
 \end{array}$$

+ Signed addition *con't*

- One important notable difference!
 - If $\text{sum} \geq 2^{w-1}$, value becomes negative (overflow)
 - If $\text{sum} < -2^{w-1}$, value becomes positive (underflow)
- An now you can explain integer overflow to all your friends!!
- See *signed_addition_overflow.c*






+

Multiplication & Division

+ Multiplication



- **Task:** Computing exact product of w -bit numbers x, y (either signed or unsigned)

- **Range of Results:**

- Unsigned multiplication requires *up to* $2*w$ bits to store result

$$0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$$

- Two's complement **min** possible value requires up to $2*w-1$ bits

$$x * y \geq (-2^{w-1}) * (2^{w-1}-1) = -2^{2w-2} + 2^{w-1}$$

- Two's complement **max** possible value requires up to $2*w$ bits


$$x * y \leq (-2^{w-1})^2 = 2^{2w-2}$$

- Therefore, maintaining exact results would need to keep expanding size with each product computed.

+ Multiplication *con't*




Operands: w bits



True product: $2*w$ bits




Discard w bits: w bits



- Similar to overflow in addition as standard multiplication function drops higher order bits
- Yields same non-intuitive results as overflow in addition.

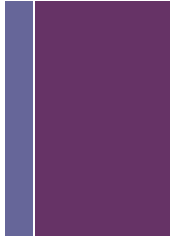
+ Multiplication *con't*



- Maintaining exact results would require...
 - Keep expanding memory footprint with each product computed
 - This is expensive!
- Therefore, must be done in software
 - e.g., by “arbitrary precision” arithmetic packages
 - ex. Java’s BigDecimal



Power-of-2 multiply with shifting



- Multiplication by a power of two is equivalent to the left shift operation.

$$u \ll k == u * 2^k$$

- Examples

$$u \ll 3 == u * 8$$


$$(u \ll 5) - (u \ll 3) == u * 24$$

$$(u + (u \ll 1)) \ll 2 == u * 12$$

- Most machines shift and add faster than multiply.



Generated multiplication code




- Compiler will convert some multiplication to shift operations during compilation process. Example, for something like this...

```
int multi_by_12(int x)
{
    return x * 12;
}
```

- The compiler generates something like this...

```
int multi_by_12(int x)
{
    int t = x + x*2
    return t << 2;
}
```

+ Unsigned power-of-2 division with shifting



- Unsigned integer division by a power of two is equivalent to right shift

$$-u \gg k == \lfloor u / 2^k \rfloor$$

- Uses logical shift.
- With signed integers, when u is negative the results are rounded incorrectly.

| | Division | Computed | Hex | Binary |
|---------------------|-------------------|--------------|--------------|--------------------------|
| x | 15213 | 15213 | 3B 6D | 00111011 01101101 |
| x >> 1 | 7606.5 | 7606 | 1D B6 | 00011101 10110110 |
| x >> 4 | 950.8125 | 950 | 03 B6 | 00000011 10110110 |
| x >> 8 | 59.4257813 | 59 | 00 3B | 00000000 00111011 |

+ Generated division code


- Again, a C compiler automatically generates shift/add code when dividing by constant.
- Example, for something like this...

```
unsigned int udiv(unsigned x)
{
    return x/8;
}
```

- The compiler generates something like this...

```
unsigned int udiv(unsigned x)
{
    return x >> 3;
}
```


+ Summary



- **Addition:**
 - Unsigned/signed: Normal addition followed by truncate, same operation on bit level
 - Unsigned: addition mod 2^w
 - Mathematical addition + possible subtraction of 2^w
 - Signed: modified addition mod 2^w (result in proper range)
 - Mathematical addition + possible addition or subtraction of 2^w
- **Multiplication:**
 - Unsigned/signed: Normal multiplication followed by truncate, same operation on bit level