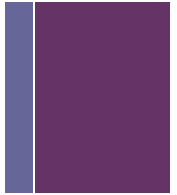




+

Procedures *con't*

# + Register Saving Conventions



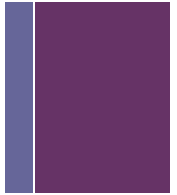
- When procedure yoo calls who:
  - `yoo()` is the *caller*
  - `who()` is the *callee*
- Can register be used for temporary storage?

```
yoo:
    . . .
    movq $15213, %rdx
    call who
    addq %rdx, %rax
    . . .
    ret
```

```
who:
    . . .
    subq $18213, %rdx
    . . .
    ret
```

- Contents of register `%rdx` overwritten by `who`
- Machine-Level programmer needs to solve for this.

# + Register Saving Conventions

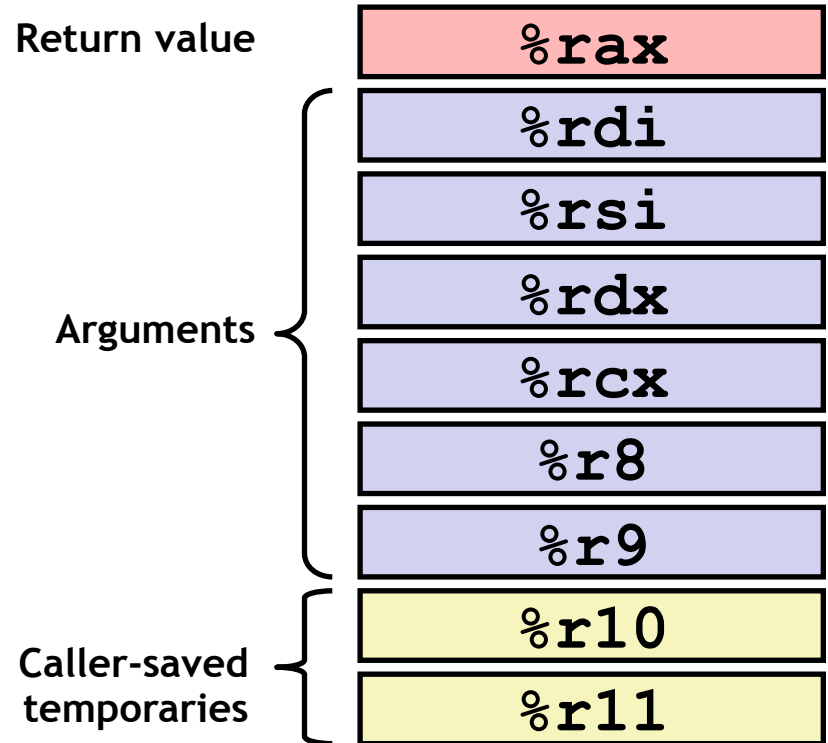


- When procedure yoo calls who:
  - **yoo ()** is the *caller*
  - **who ()** is the *callee*
- Can register be used for temporary storage?
- **Conventions**
  - “*Caller Saved*”
    - Caller saves temporary values in its frame before the call
  - “*Callee Saved*”
    - Callee saves temporary values in its frame before using
    - Callee restores them before returning to caller

# + x86-64 Linux Caller-saved Registers




- **%rax**
  - Return value
  - Also caller-saved
  - Can be modified by procedure
- **%rdi, ..., %r9**
  - Arguments
  - Also caller-saved
  - Can be modified by procedure
- **%r10, %r11**
  - Caller-saved
  - Can be modified by procedure



# + x86-64 Linux Callee-saved Registers



- **%rbx, %r12, %r13, %r14**
  - Callee-saved
  - Callee must save & restore
- **%rbp**
  - Callee-saved
  - Callee must save & restore
  - May be used as frame pointer
- **%rsp**
  - Special form of callee save
  - Restored to original value upon exit from procedure



# + Callee-Saved Example




```
long incr(long *p, long val) {  
    long x = *p;  
    long y = x + val;  
    *p = y;  
    return x;  
}
```

```
long call_incr(long x) {  
    long v1 = 15213;  
    long v2 = incr(&v1, 3000);  
    return x+v2;  
}
```


```
call_incr:  
    pushq    %rbx           # save %rbx  
    subq     $8, %rsp       # allocate  
    movq     $15213, (%rsp) # "push"  
    movq     %rdi, %rbx  
    movq     $3000, %rsi  
    leaq     (%rsp), %rdi  
    call     incr  
    addq     %rbx, %rax  
    addq     $8, %rsp  
    popq     %rbx  
    ret
```

## call\_incr's stack

### Initial Stack



### Resulting Stack



# + Callee-Saved Example *con't*




```
long incr(long *p, long val) {  
    long x = *p;  
    long y = x + val;  
    *p = y;  
    return x;  
}
```

```
long call_incr(long x) {  
    long v1 = 15213;  
    long v2 = incr(&v1, 3000);  
    return x+v2;  
}
```


```
call_incr:  
    pushq    %rbx  
    subq     $8, %rsp  
    movq     %rdi, %rbx  
    movq     $15213, (%rsp)  
    movq     $3000, %rsi  
    leaq     (%rsp), %rdi  
    call     incr  
    addq     %rbx, %rax  
    addq     $8, %rsp          # deallocate  
    popq     %rbx             # restore %rbx  
    ret
```

## call\_incr's stack

### Resulting Stack



### Pre-return Stack





## Data: Arrays




# + Array Allocation



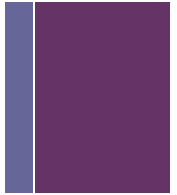
- Basic Principle

$T$   $A[N]$  ;

- Array of data type  $T$  and length  $N$
- Contiguously allocated region of  $N * \text{sizeof}(T)$  bytes in memory




# + Array Access



## ▪ Basic Principle


$T$   $A[N]$  ;

- Array of data type  $T$  and length  $N$
- Identifier  $A$  can be used as a pointer to array element 0: type  $T^*$



Reference	Type	Value
<code>val[4]</code>	<code>int</code>	3
<code>val</code>	<code>int*</code>	$x$
<code>val+1</code>	<code>int*</code>	$x + 4$
<code>&amp;val[2]</code>	<code>int*</code>	$x + 8$
<code>val[5]</code>	<code>int</code>	??
<code>*(val+1)</code>	<code>int</code>	0
<code>val + i</code>	<code>int*</code>	$x + 4 i$

# + Array Accessing Example

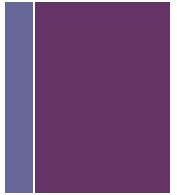


```
int get_digit(int[] z, int digit){  
    return z[digit];  
}
```

```
# %rdi = z  
# %rsi = digit  
movq (%rdi,%rsi,4), %rax # z[digit]
```

- Register `%rdi` contains starting address of array
- Register `%rsi` contains array index
- Desired digit at `%rdi + 4*%rsi`
- Use memory reference `(%rdi,%rsi,4)`

# + Array Loop Example



```
void zincr(int[] z) {  
    int i;  
    for (i = 0; i < 5; i++)  
        z[i]++;  
}
```

```
# %rdi = z, %rax = i  
    movl    $0, %rax          # i = 0  
    jmp     .L3               # goto middle  
.L4:                          # loop:  
    addl    $1, (%rdi,%rax,4)  # z[i]++  
    addq    $1, %rax          # i++  
.L3:                          # middle  
    cmpq    $4, %rax          # i:4  
    jbe     .L4               # if <=, goto loop  
    rep; ret
```

# + Multidimensional Arrays



- **Declaration**

`T A[R][C];`


- 2D array of data type  $T$
- $R$  rows,  $C$  columns
- Type  $T$  element requires  $K$  bytes

- **Array Size**


- $R * C * K$  bytes

- **Arrangement**

- Row-Major Ordering




`int A[R][C];`



# + Nested Array Example



```
int[4][5] zips =  
    {{1, 5, 2, 0, 6},  
     {1, 0, 0, 0, 3 },  
     {1, 5, 2, 1, 7 },  
     {1, 5, 2, 2, 1 }};
```



- Variable **zips**: array of 4 elements, allocated contiguously
- Each element is an array of 5 **int**'s, allocated contiguously
- “Row-Major” ordering of all elements in memory

# + Nested Array Row Access



- Declaration


$T \ A[R][C];$

- Row Vectors

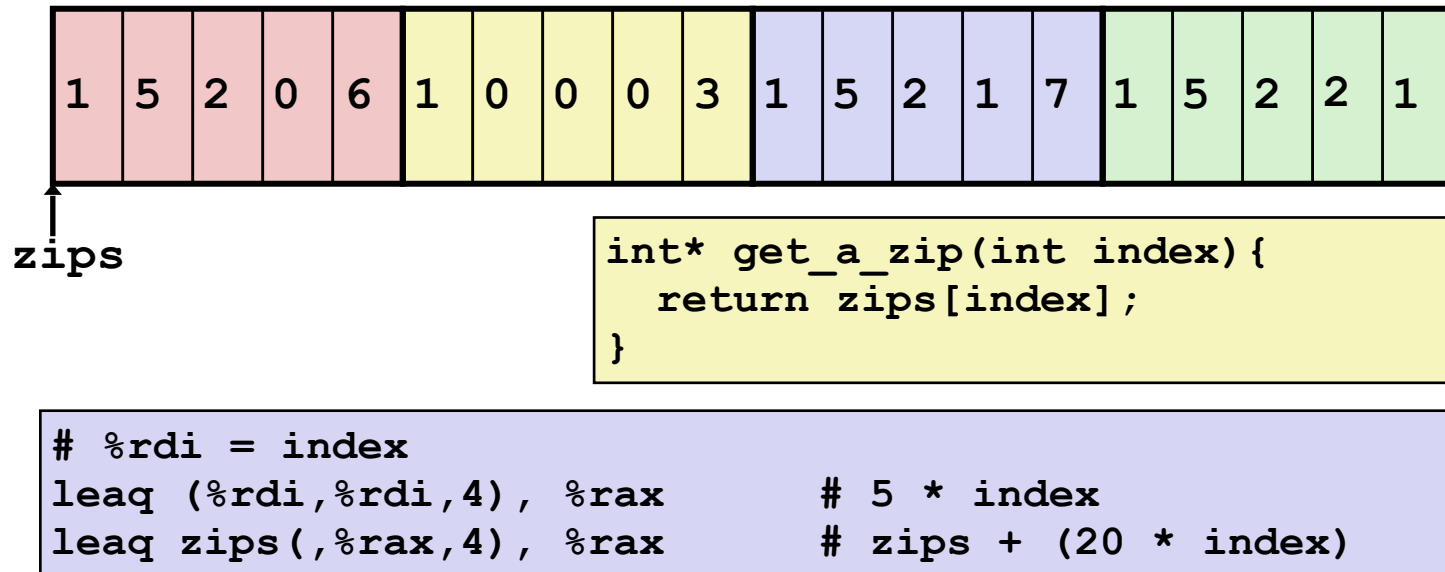
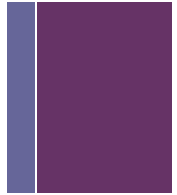
- $A[i]$  is an array of  $C$  elements, e.g. a “column”
- Each element of type  $T$  requires  $K$  bytes
- Therefore the starting address of column is  $A + i * (C * K)$

- Example

`int A[R][C];`



# + Nested Array Row Access Code



- **Row Vector**
  - **zips[index]** is array of 5 int's
  - Starting address **zips + 20 \* index**
- **Machine Code**
  - Computes and returns address
  - Compute as **zips + 4 \* (index + 4 \* index)**




# + Nested Array Element Access



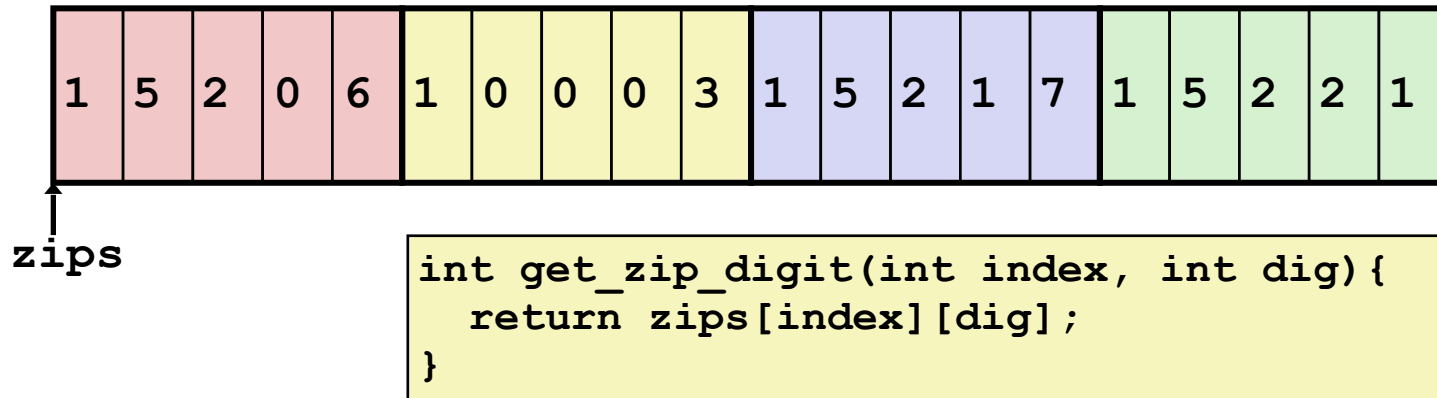
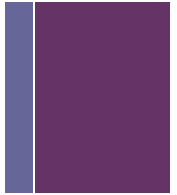
- Array Elements

- $A[i][j]$  is element of type  $T$ , which requires  $K$  bytes
- Address:  $A + i * (C * K) + j * K$

```
int A[R][C];
```



# + Nested Array Element Access Code



```
leaq    (%rdi,%rdi,4), %rax    # 5 * index (%rdi is index)  
addl    %rax, %rsi             # 5 * index + dig  
movl    zips(,%rsi,4), %rax    # M[zips + 4*(5 * index + dig)]
```

- Array Elements

- `zips[index][dig]` is an int
- Address: `zips + 20*index + 4*dig`
  - Expressed in assembly as `zips + 4*(5 * index + dig)`


# + Multi-Level Array Example



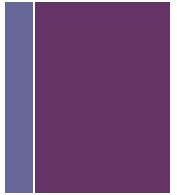
```
int* nyu = { 1, 0, 0, 0, 3 };  
int* mit = { 0, 2, 1, 3, 9 };  
int* ucb = { 9, 4, 7, 2, 0 };
```

```
#define UCOUNT 3  
int *univ[UCOUNT] = {mit, nyu, ucb};
```


- Variable **univ** denotes an array of 3 elements
- Each element is a pointer (8 bytes)
- Each pointer points to array of int's



# + Element Access in Multi-Level Array



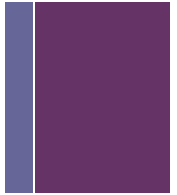
```
int get_uni_digit
(size_t index, size_t digit){
    return univ[index][digit];
}
```



```
salq    $2, %rsi          # 4*digit
addq    univ(,%rdi,8), %rsi # p = univ[index] + 4*digit
movl    (%rsi), %eax       # return *p
ret
```


- Computation
  - Element access **Mem[Mem[univ+8\*index]+4\*digit]**
  - Must do two memory reads
    - First get pointer to row array
    - Then access element within array

# + Array Element Accesses




## Nested array

```
int get_zip_digit
(size_t index, size_t digit)
{
    return zips[index][digit];
}
```



## Multi-level array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



Accesses looks similar in C, but address computations very different:

`Mem[zip+20*index+4*digit]`

`Mem[Mem[univ+8*index]+4*digit]`




+

Data: Structs

# + Structure Representation




```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



- **Structure represented as block of memory**
  - Big enough to hold all of the fields
- **Fields ordered according to declaration**
  - Even if another ordering could yield a more compact representation
- **Compiler determines overall size + positions of fields**
  - Machine-level program has no understanding of the structures in the source code

# + Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    int i;  
    struct rec* next;  
};
```



```
int* get_array_ptr(struct rec* r, int idx){  
    return &r->a[idx];  
}
```

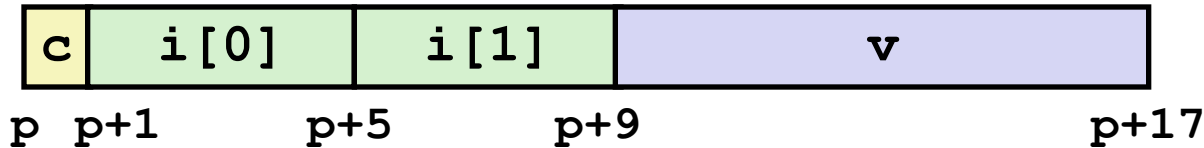
```
leaq  (%rdi,%rsi,4), %rax    # r in %rdi, idx in %rsi  
ret                                # move ptr into %rax, return
```

- Generating pointer to array element
  - Offset of each structure member determined at compile time
  - Compute as  $r + 4 * idx$



# + Structures & Alignment

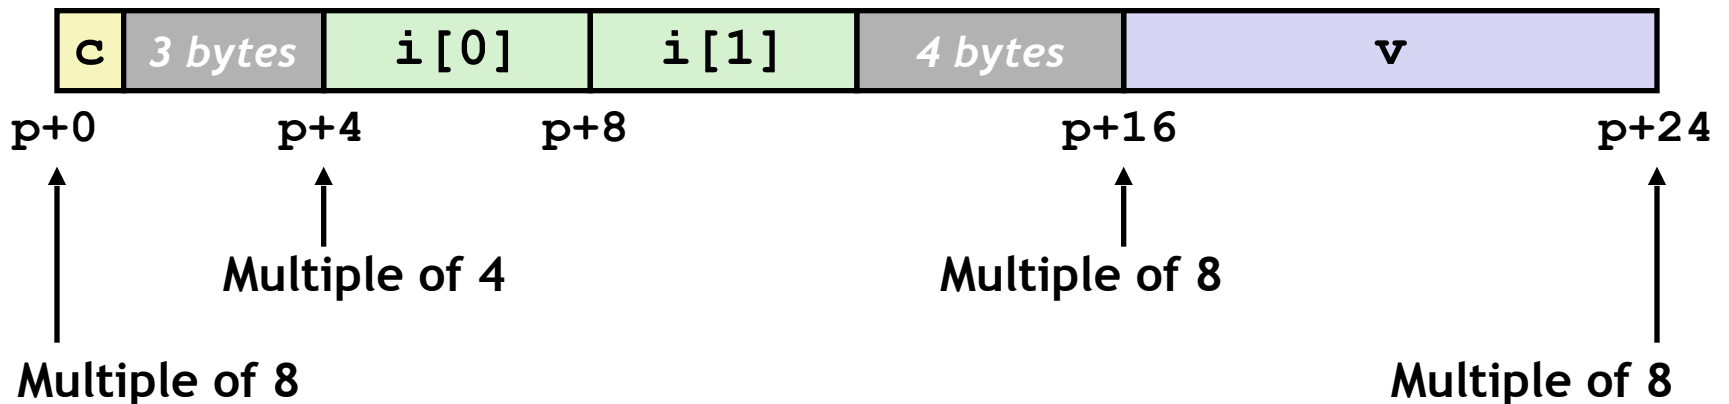
## ■ Unaligned Data




```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

## ■ Aligned Data

- Primitive data type requires  $K$  bytes
- Address must be multiple of  $K$



# + Alignment Principles



- **Aligned Data**

- Primitive data type requires  $K$  bytes
- Address must be multiple of  $K$
- Required on some machines; advised on x86-64

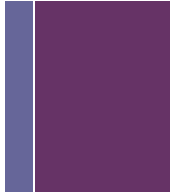
- **Motivation for Aligning Data**

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
  - Inefficient to load or store datum that spans word boundaries
  - Virtual memory trickier when you allow unaligned data.

- **Compiler**

- Inserts gaps in structure to ensure correct alignment of fields

# + Specific Cases of Alignment (x86-64)



- 1 byte: **char**, ...
  - no restrictions on address
- 2 bytes: **short**, ...
  - lowest 1 bit of address must be  $0_2$
- 4 bytes: **int**, **float**, ...
  - lowest 2 bits of address must be  $00_2$
- 8 bytes: **double**, **long**, **char \***, ...
  - lowest 3 bits of address must be  $000_2$

# + Satisfying Alignment with Structures

- **Within structure:**

- Must satisfy each element's alignment requirement


- **Overall structure placement**

- Each structure has alignment requirement **K**
  - **K** = Largest alignment of any element
- Initial address & structure length must be multiples of **K**

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

- **Example:**


- **K** = 8, due to **double** element




# + Meeting Overall Alignment Requirement

- For largest alignment requirement K
- Overall structure must be multiple of K

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



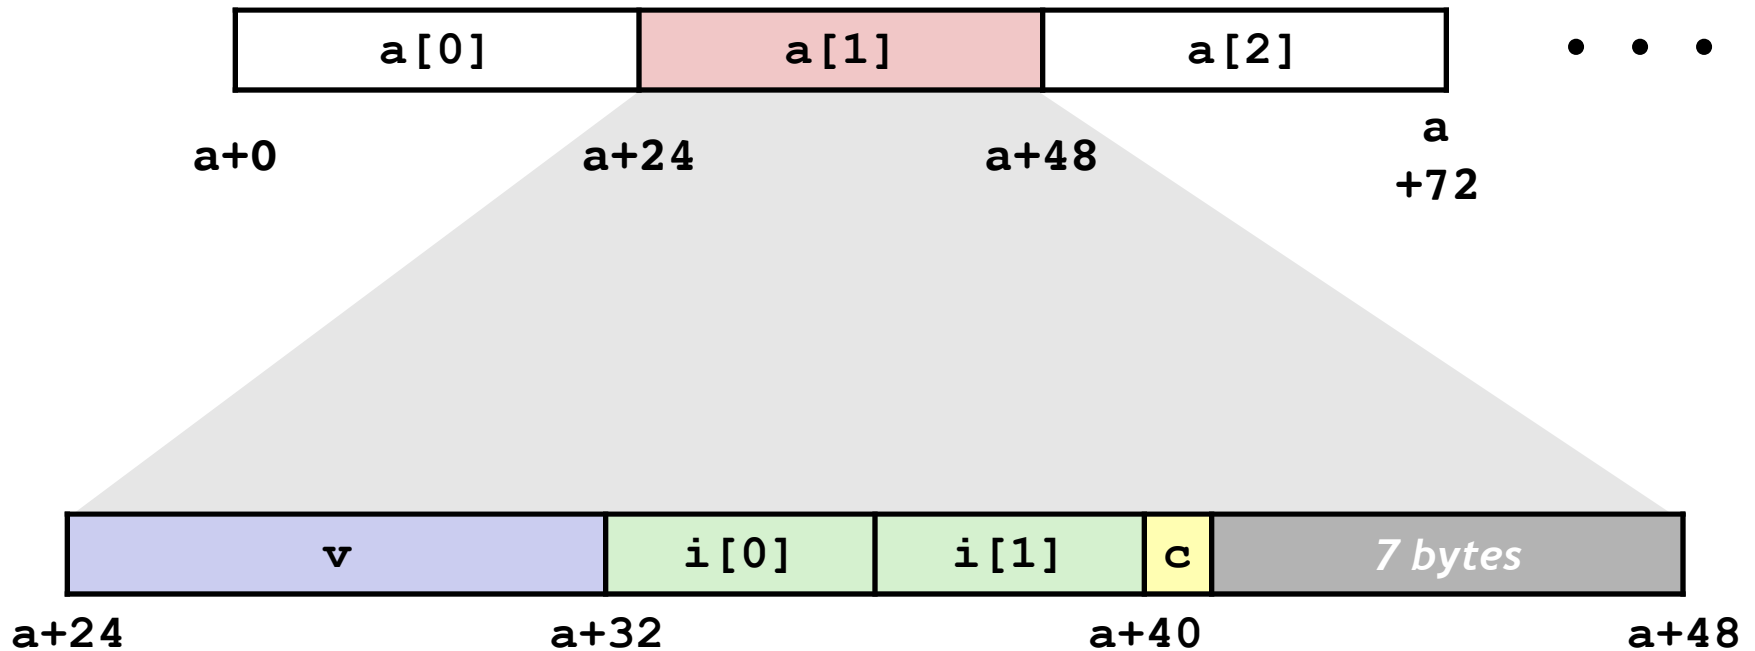
Multiple of K=8



# + Arrays of Structures

- Overall structure length multiple of K
- Satisfy alignment requirement for every element

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```




# + Accessing Members in Array of Struct



- Compute array offset  $12 * \text{idx}$ 
  - $\text{sizeof}(i) + \text{sizeof}(v) + \text{sizeof}(j) + 4 \text{ bytes}$   
 $== \text{sizeof}(S3) == 12 \text{ bytes}$  (4 bytes is the padding)
- Element  $j$  is at an offset of 8 bytes within structure
  - $\text{sizeof}(i) + 2 \text{ bytes} + \text{sizeof}(\text{float})$   
 $== 8 \text{ bytes}$


```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```



# + Saving Space (Struct packing)

- Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

- Effect (K=12 -> K=4)

