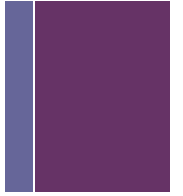





Exceptions

+ Control Flow



- **Processors do only one thing:**
 - From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
 - This sequence is the CPU's control flow (or flow of control)




+ Altering the Control Flow



- **Up to now: two mechanisms for changing control flow:**
 - Jumps and branches
 - Call and return
 - Reactions to changes in *program state*
- **Insufficient for a useful system:**
Difficult to react to changes in *system state*
 - Data arrives from a disk or a network adapter
 - Instruction divides by zero
 - System timer expires
- **System needs mechanisms for “exceptional control flow”**

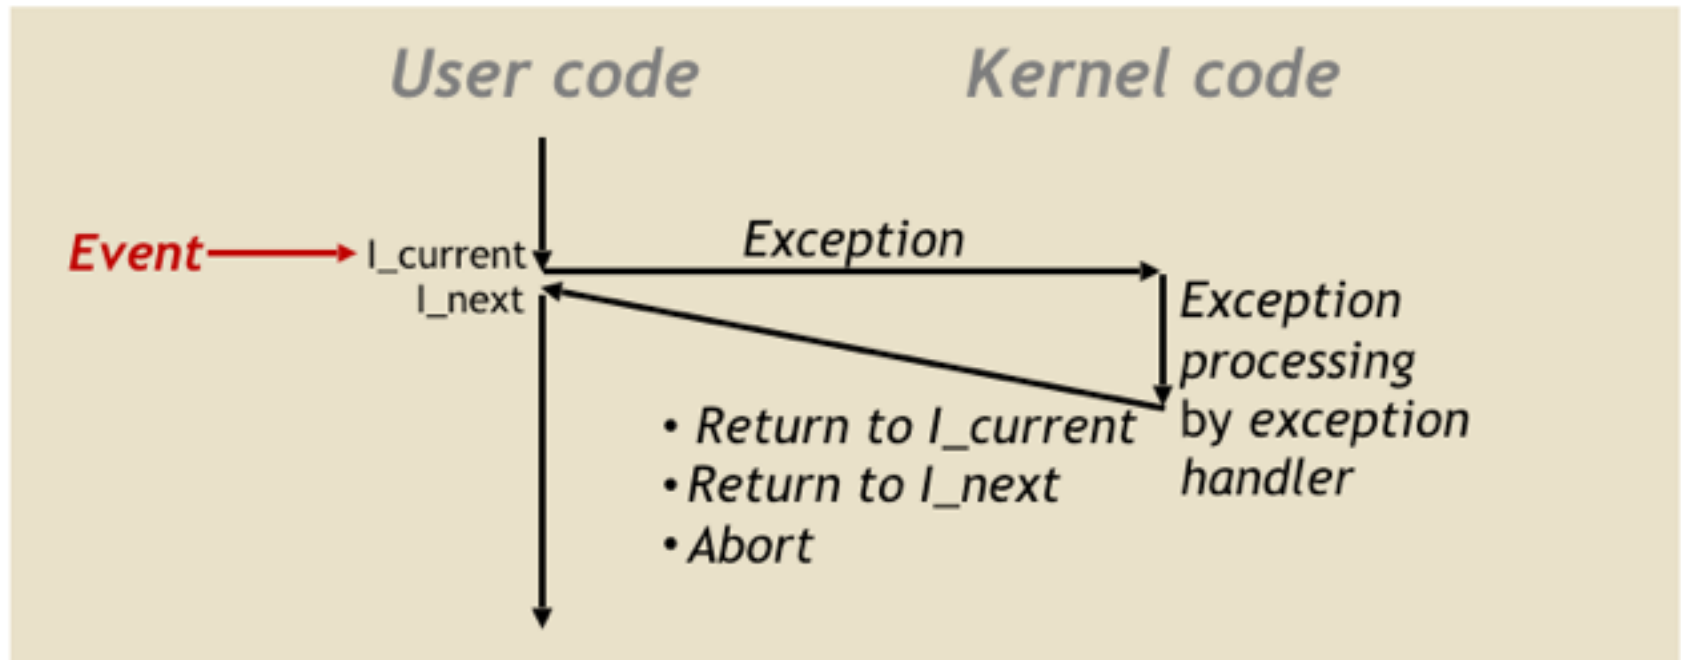
+ Exceptional Control Flow




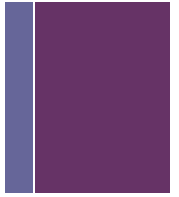
- **Exists at all levels of a computer system**
- **Low level mechanisms**
 - 1. Exceptions
 - Change in control flow in response to a system event (i.e., change in system state)
 - Implemented using combination of hardware and OS software
- **Higher level mechanisms**
 - 2. Process context switch
 - Implemented by OS software and hardware timer
 - 3. Signals
 - Implemented by OS software (next lecture)
 - 4. Nonlocal jumps: `setjmp()` and `longjmp()`
 - Implemented by C runtime library

+ Exceptions

- An exception is a transfer of control to the OS kernel in response to some event (i.e., change in processor state)
 - Kernel is the memory-resident part of the OS
 - Examples of events: Divide by 0, page fault, I/O request completes, typing Ctrl-C

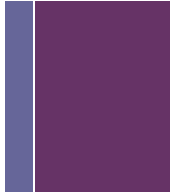


+ Exception Tables



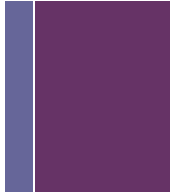
- Each type of event has a unique exception number k
- k = index into exception table (a.k.a. interrupt vector)
- Handler k is called each time exception k occurs

+ Asynchronous Exceptions (Interrupts)



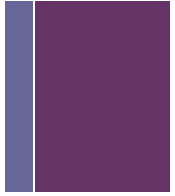
- **Caused by events external to the processor**
 - Indicated by setting the processor's *interrupt pin*
 - Handler returns to “next” instruction
- **Examples:**
 - Timer interrupt
 - Every few ms, an external timer chip triggers an interrupt
 - Used by the kernel to take back control from user programs
 - I/O interrupt from external device
 - Arrival of a packet from a network
 - Arrival of data from a disk

+ Synchronous Exceptions



- **Caused by events that occur as a result of executing an instruction:**
 - Traps
 - Intentional
 - Example: *system calls*
 - Returns control to “next” instruction
 - Faults
 - Unintentional but possibly recoverable
 - Example: *page fault*
 - Either re-executes faulting (“current”) instruction or aborts
 - Aborts
 - Unintentional and unrecoverable
 - Example: *illegal memory access*
 - Aborts current program

+ System Calls



Each x86-64 system call has a unique ID number

Examples:

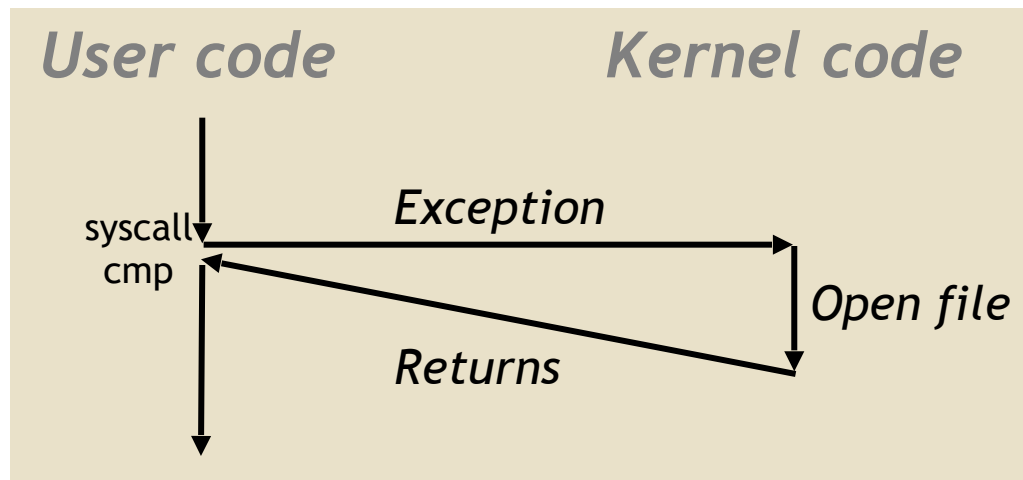
<i>Number</i>	<i>Name</i>	<i>Description</i>
0	read	Read file
1	write	Write file
2	open	Open file
3	close	Close file
4	stat	Get info about file
57	fork	Create process
59	execve	Execute a program
60	_exit	Terminate process
62	kill	Send signal to process

+ System Call Example: Opening File



- User calls: `open(filename, options)`
- Calls `__open` function, which invokes system call instruction `syscall`

```
0000000000e5d70 <__open>:  
...  
e5d79:  b8 02 00 00 00      mov  $0x2,%rax  # open is syscall #2  
e5d7e:  0f 05               syscall         # Return value in %rax  
e5d80:  48 3d 01 f0 ff ff   cmp  $0xffffffffffffffff001, %rax  
...  
e5dfa:  c3                 retq
```



- `%rax` contains syscall number
- Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
- Return value in `%rax`
- Negative value is an error corresponding to negative `errno`

+ Fault Example: Page Fault


- User writes to memory location
- That portion (page) of user's memory is currently on disk

```
int a[10000];  
main ()  
{  
    a[5000] = 13;  
}
```

80483b7:	c7 05 10 9d 04 08 0d	movq	\$0xd,0x8049d10
----------	----------------------	------	-----------------

User code

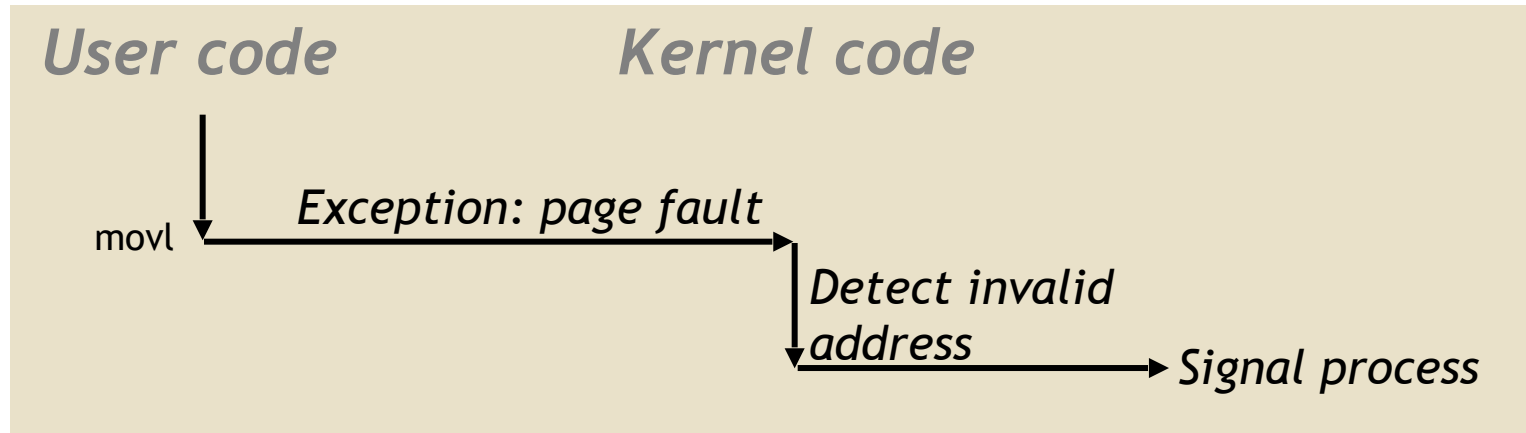
Kernel code



+ Fault Example: Invalid Memory Reference

```
int a[1000];  
main () {  
    a[5000] = 13;  
}
```

80483b7: c7 05 60 e3 04 08 0d movq \$0xd,0x804e360



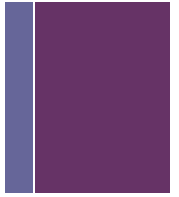
- Sends SIGSEGV signal to user process
- User process exits with “segmentation fault”




+

Processes


+ Processes



- A *process* is an instance of a running program.
 - One of the most successful ideas in computer science
 - Not the same as “program”
- Process provided with two key abstractions by OS:
 - *Logical control flow*
 - Each program seems to have exclusive use of the CPU
 - Provided by kernel mechanism called *context switching*
 - *Private address space*
 - Each program seems to have exclusive use of main memory.
 - Provided by kernel mechanism called *virtual memory*

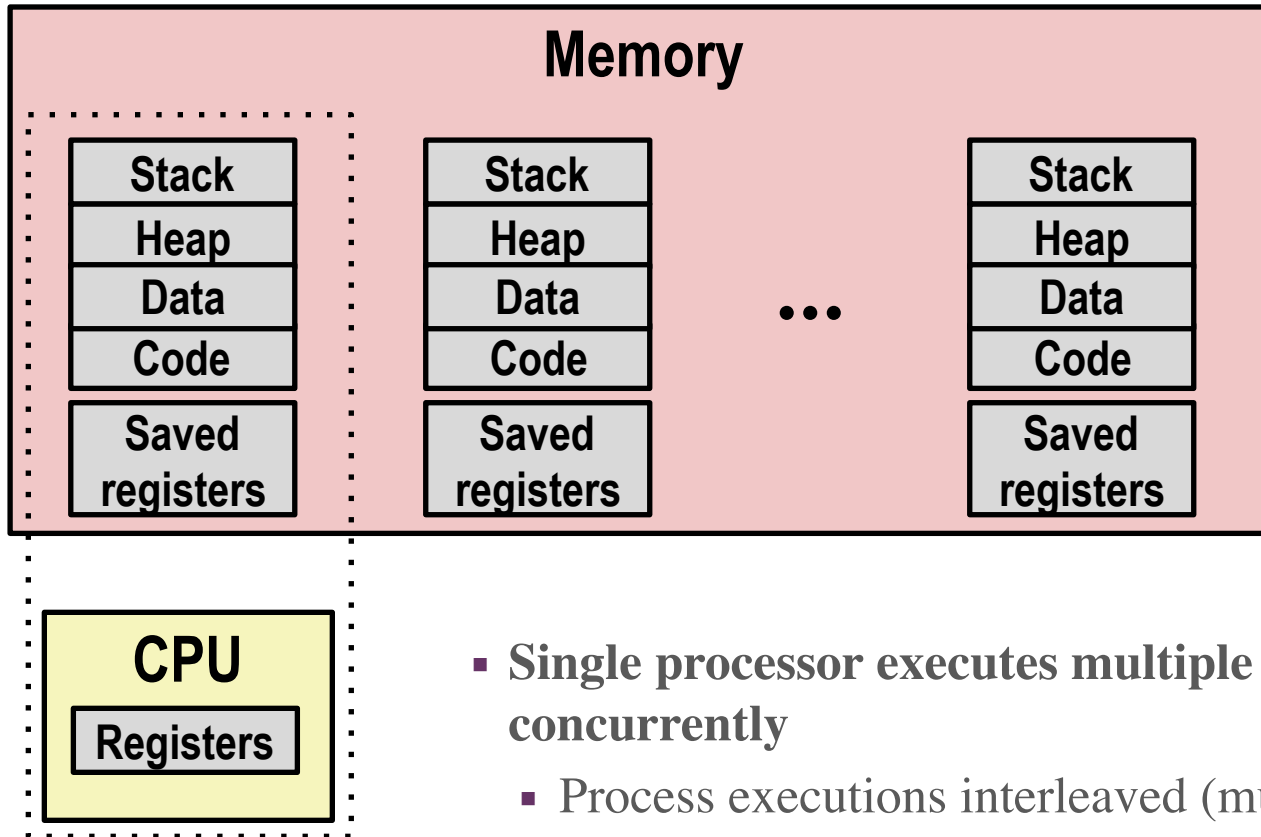



+ Multiprocessing: The Illusion



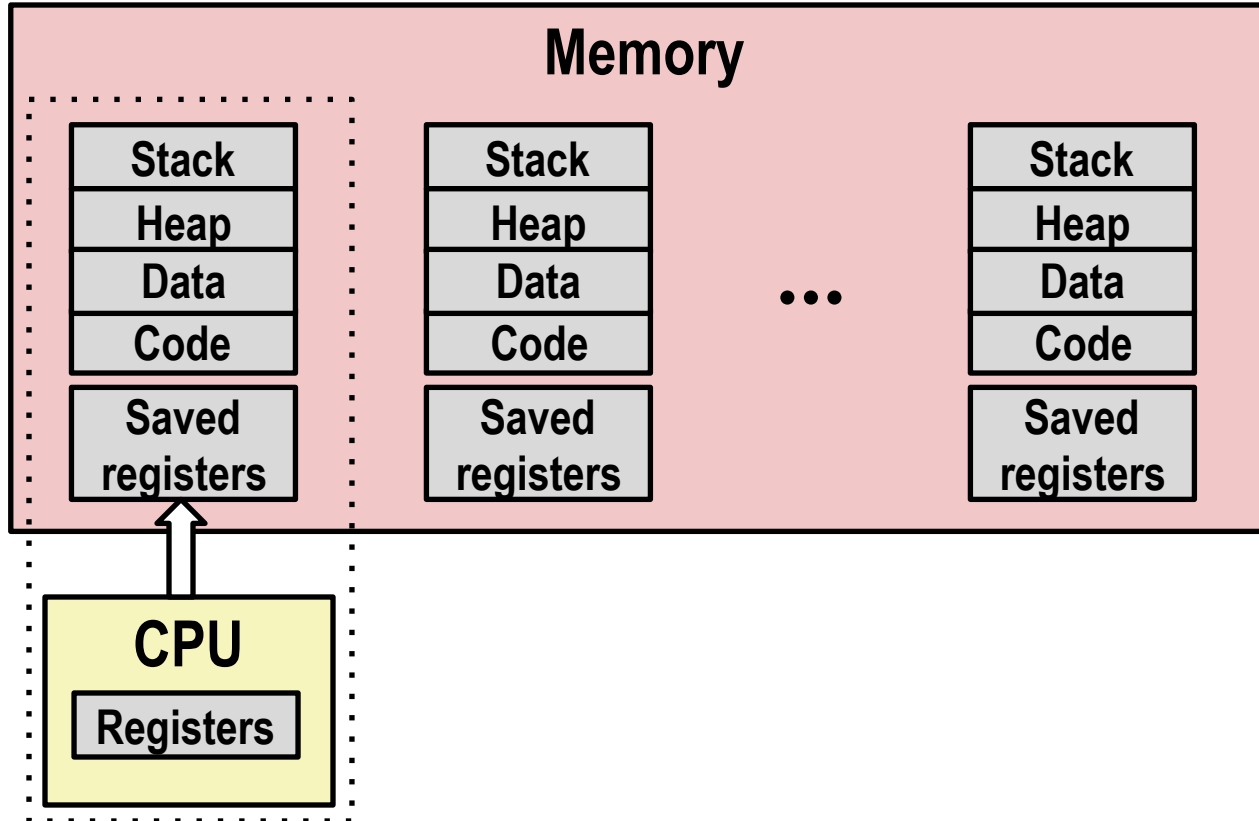
- **Computer runs many processes simultaneously**
 - Applications for one or more users
 - Web browsers, email clients, editors, ...
 - Background tasks
 - Monitoring network & I/O devices

+ Multiprocessing: The (Traditional) Reality



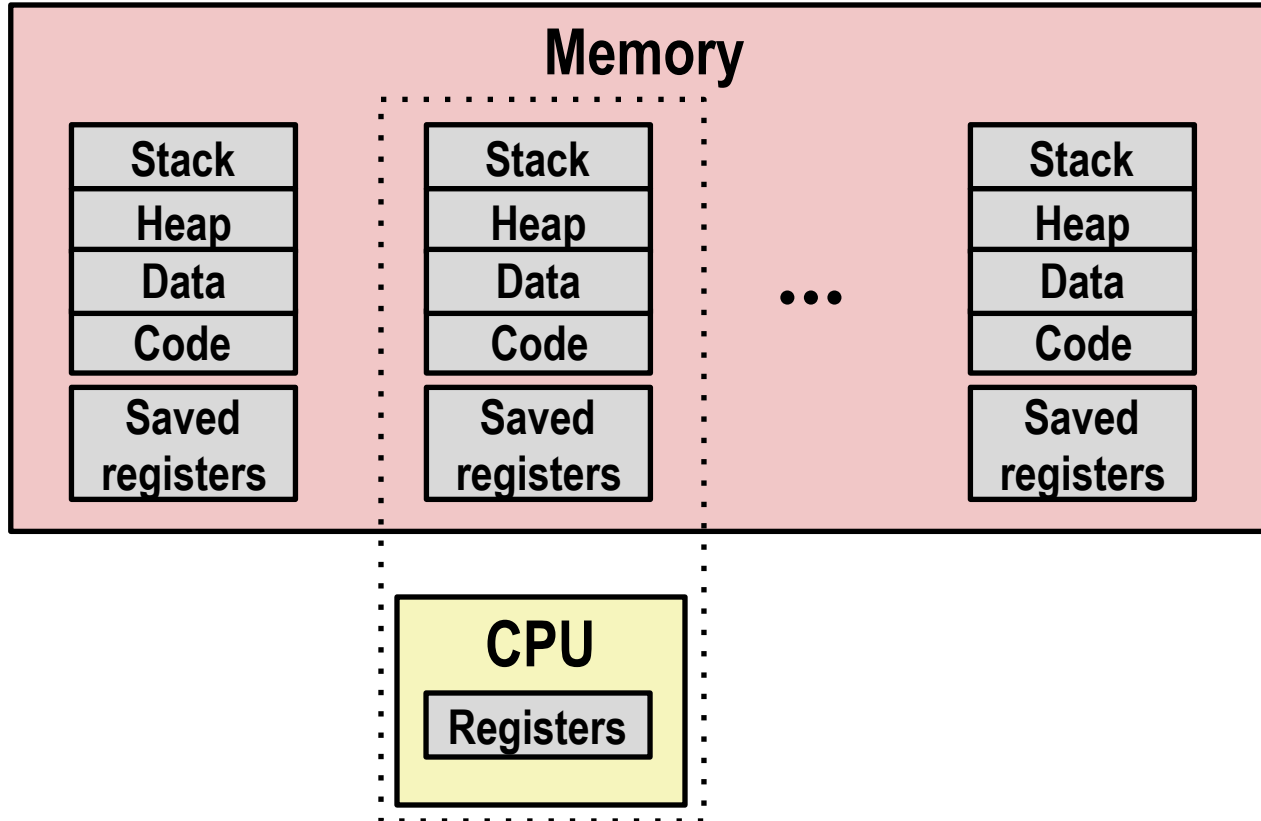
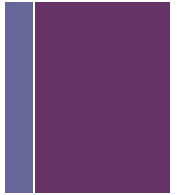
- **Single processor executes multiple processes concurrently**
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system (later in course)
 - Register values for non-executing processes saved in memory (context)

+ Multiprocessing: The (Traditional) Reality



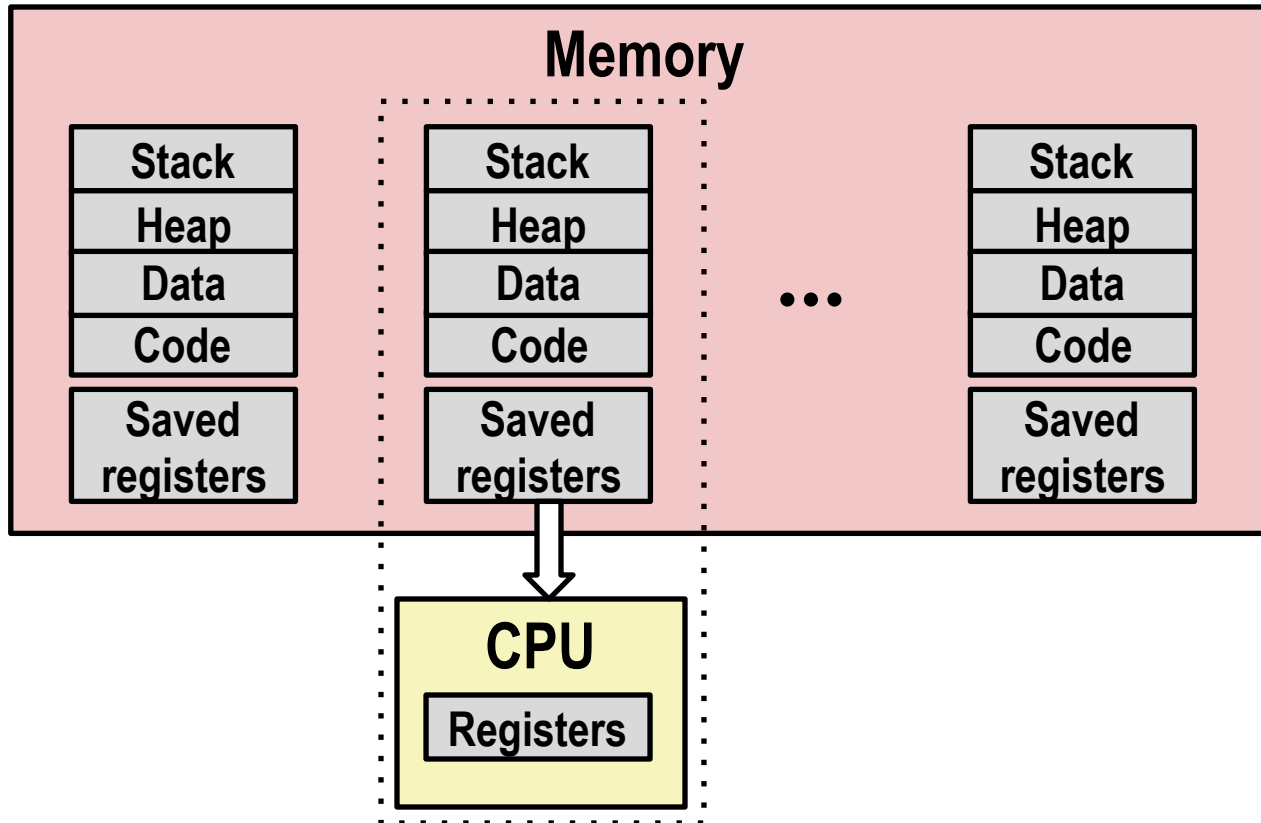
- Save current registers in memory

+ Multiprocessing: The (Traditional) Reality



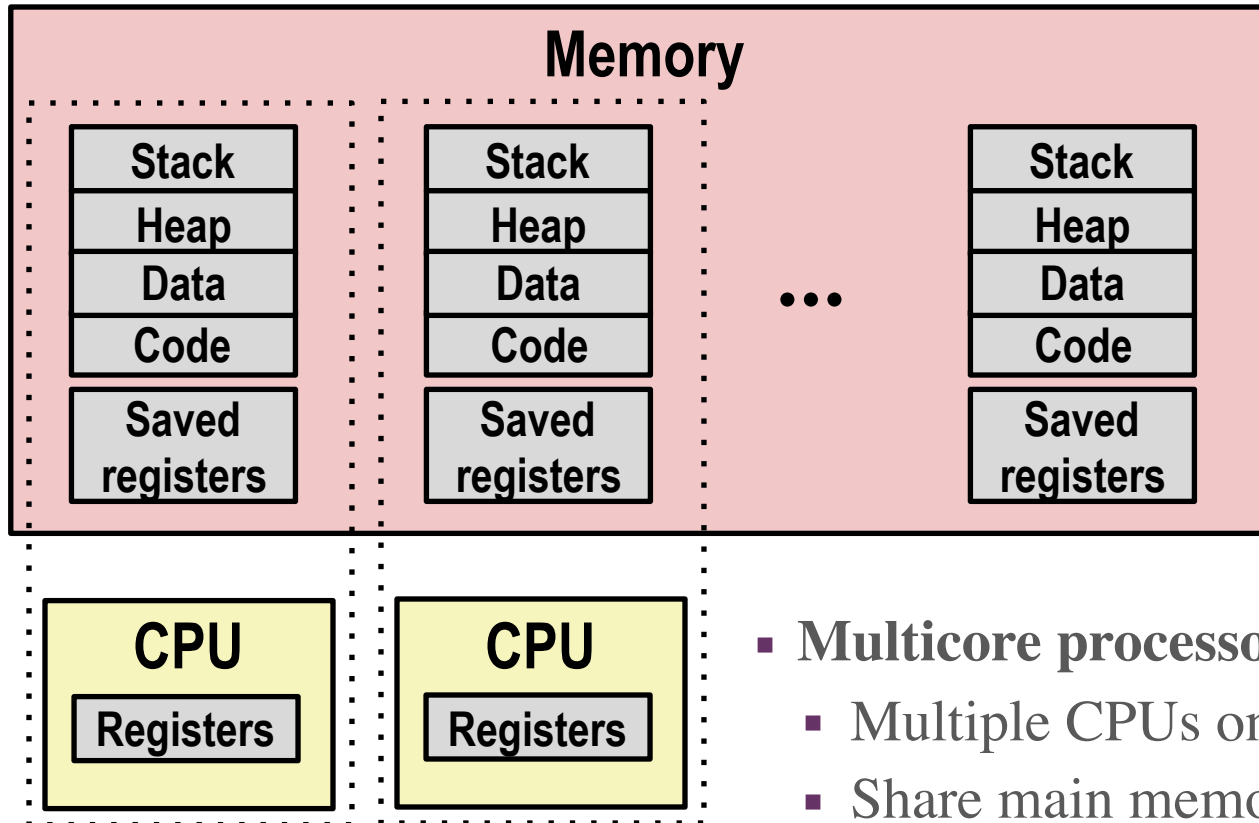

- Schedule next process for execution

+ Multiprocessing: The (Traditional) Reality



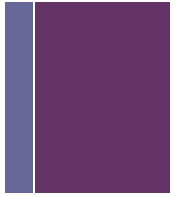
- Load saved registers and switch address space (context switch)

+ Multiprocessing: The (Modern) Reality




- **Multicore processors**
 - Multiple CPUs on single chip
 - Share main memory (and some of the caches)
 - Each can execute a separate process (Scheduling of processors onto cores by kernel)

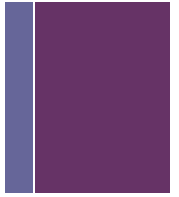
+ Concurrent Processes



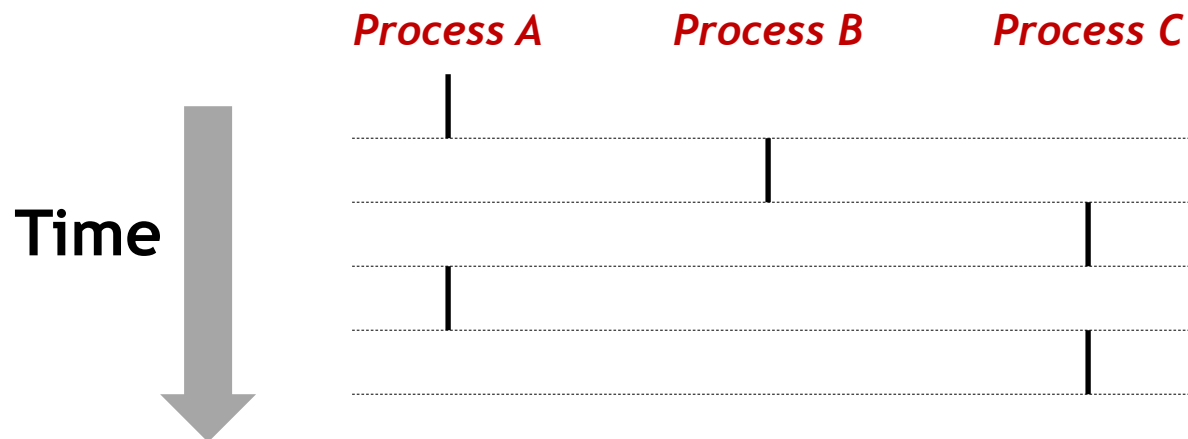
- Each process is a logical control flow.
- Two processes run concurrently if their flows overlap in time
 - Otherwise, they are sequential
- Examples (running on single core):
 - Concurrent: ??
 - Sequential: ??



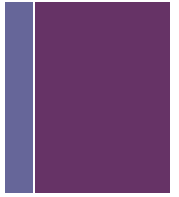
+ Concurrent Processes




- Each process is a logical control flow.
- Two processes run concurrently if their flows overlap in time
 - Otherwise, they are sequential
- Examples (running on single core):
 - Concurrent: A & B, A & C
 - Sequential: B & C



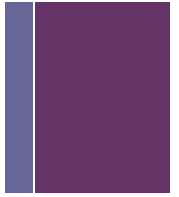
+ Concurrent vs Parallel Processes



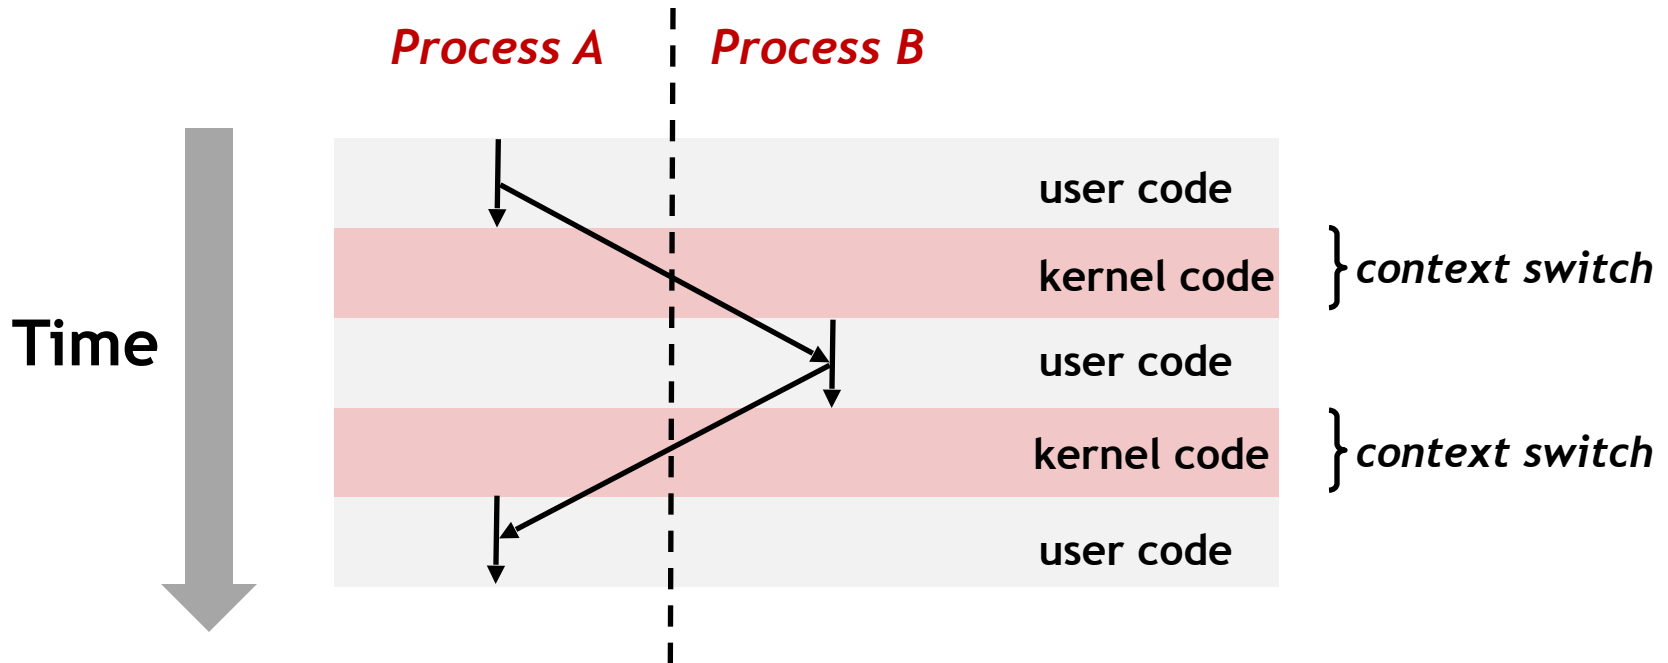
- Control flows for concurrent processes are physically disjoint in time
- We can *think* of concurrent processes as running in parallel with each other, however, this is not necessarily the case.
- If two processes are executing simultaneously on different cores, these are true parallel processes.



+ Context Switching



- Processes are managed by a shared chunk of memory-resident OS code called the kernel
 - Important: the kernel is not a separate process, but rather runs as part of all processes.
- Control flow passes from one process to another via a context switch






+

Process Control

+ Creating and Terminating Processes



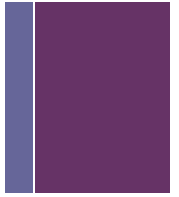
- **From a programmer's perspective, we can think of a process as being in one of three states**
- **Running**
 - Process is either executing, or waiting to be executed and will eventually be scheduled (i.e., chosen to execute) by the kernel
- **Stopped**
 - Process execution is suspended and will not be scheduled until further notice
- **Terminated**
 - Process is stopped permanently

+ Terminating Processes



- **Process becomes terminated for one of three reasons:**
 - Receiving a signal whose default action is to terminate (next lecture)
 - Returning from the **main** routine
 - Calling the **exit** function
- **void exit(int status)**
 - Terminates with an exit status of status
 - Convention: normal return status is 0, nonzero on error
 - Another way to explicitly set the exit status is to return an integer value from the main routine

+ Creating Processes



- Parent process creates a new running child process by calling `fork`
- `int fork(void)`
 - Returns 0 to the child process, child's PID to parent process
 - Child is almost identical to parent:
 - Child get an identical (but separate) copy of the parent's virtual address space.
 - Child gets identical copies of the parent's open file descriptors
 - Child has a different PID than the parent
- `fork` is interesting (and often confusing) because it is called once but returns twice

+ fork Example



```
int main()
{
    pid_t pid;
    int x = 1;


    pid = fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```

- Call once, return twice
- Concurrent execution
 - Can't predict execution order of parent and child
- Duplicate but separate address space
 - x has a value of 1 when fork returns in parent and child
 - Subsequent changes to x are independent

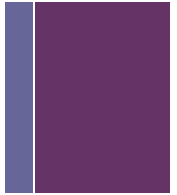
```
linux> ./fork
parent: x=0
child : x=2
```

+ Modeling `fork` with Process Graphs




- A *process graph* is a useful tool for capturing the partial ordering of statements in a concurrent program:
 - Each vertex is the execution of a statement
 - $a \rightarrow b$ means a happens before b
 - Edges can be labeled with current value of variables
 - `printf` vertices can be labeled with output
 - Each graph begins with a vertex with no in-edges
- Any topological sort of the graph corresponds to a feasible total ordering.
 - Total ordering of vertices where all edges point from left to right

+ Process Graph Example




```
int main() {  
    pid_t pid;  
    int x = 1;  
  
    pid = fork();  
    if (pid == 0) { /* Child */  
        printf("child : x=%d\n", ++x);  
        exit(0);  
    }  
  
    /* Parent */  
    printf("parent: x=%d\n", --x);  
    exit(0);  
}
```




+ Interpreting Process Graphs




- Original graph:




- Re-labeled graph:



Feasible total ordering:




Infeasible total ordering:



+ fork Example: Two consecutive forks

```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```



Feasible output:


L0
L1
Bye
Bye
L1
Bye
Bye

Infeasible output:

L0
Bye
L1
Bye
L1
Bye
Bye

+ fork Example: Nested forks in parent

```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```



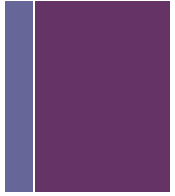
Feasible output:

L0
L1
Bye
Bye
L2
Bye

Infeasible output:

L0
Bye
L1
Bye
Bye
L2

+ Reaping Child Processes



- **Idea**

- When process terminates, it still consumes system resources
- Called a “**zombie**”

- **Reaping**

- Performed by parent on terminated child (using wait or waitpid)
- Parent is given exit status information
- Kernel then deletes zombie child process

- **What if parent doesn't reap?**

- If any parent terminates without reaping a child, then the orphaned child will be reaped by init process (pid == 1)
- So, only need explicit reaping in long-running processes
 - e.g., shells and servers

+Zombie Example

```
void fork7() {  
    if (fork() == 0) {  
        /* Child */  
        printf("Terminating Child, PID = %d\n", getpid());  
        exit(0);  
    } else {  
        printf("Running Parent, PID = %d\n", getpid());  
        while (1)  
            ; /* Infinite loop */  
    }  
}
```

```
linux> ./forks 7 &  
[1] 6639  
Running Parent, PID = 6639  
Terminating Child, PID = 6640  
linux> ps  
  PID TTY          TIME CMD  
 6585 ttyp9      00:00:00 tcsh  
 6639 ttyp9      00:00:03 forks  
 6640 ttyp9      00:00:00 forks <defunct>  
 6641 ttyp9      00:00:00 ps  
linux> kill 6639  
[1]      Terminated  
linux> ps  
  PID TTY          TIME CMD  
 6585 ttyp9      00:00:00 tcsh  
 6642 ttyp9      00:00:00 ps
```

- **ps** shows child process as “defunct” (i.e., a zombie)
- Killing parent allows child to be reaped by init

+ Non-Terminating Child Example

```
void fork8()
{
    if (fork() == 0) {
        /* Child */
        printf("Child, PID = %d\n", getpid());
        while (1)
            ; /* Infinite loop */
    } else {
        printf("Parent, PID = %d\n", getpid());
        exit(0);
    }
}
```

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9        00:00:00 tcsh
 6676 tttyp9        00:00:06 forks
 6677 tttyp9        00:00:00 ps
linux> kill 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9        00:00:00 tcsh
 6678 tttyp9        00:00:00 ps
```

- Child process still active even though parent has terminated
- Must kill child explicitly, or else will keep running indefinitely