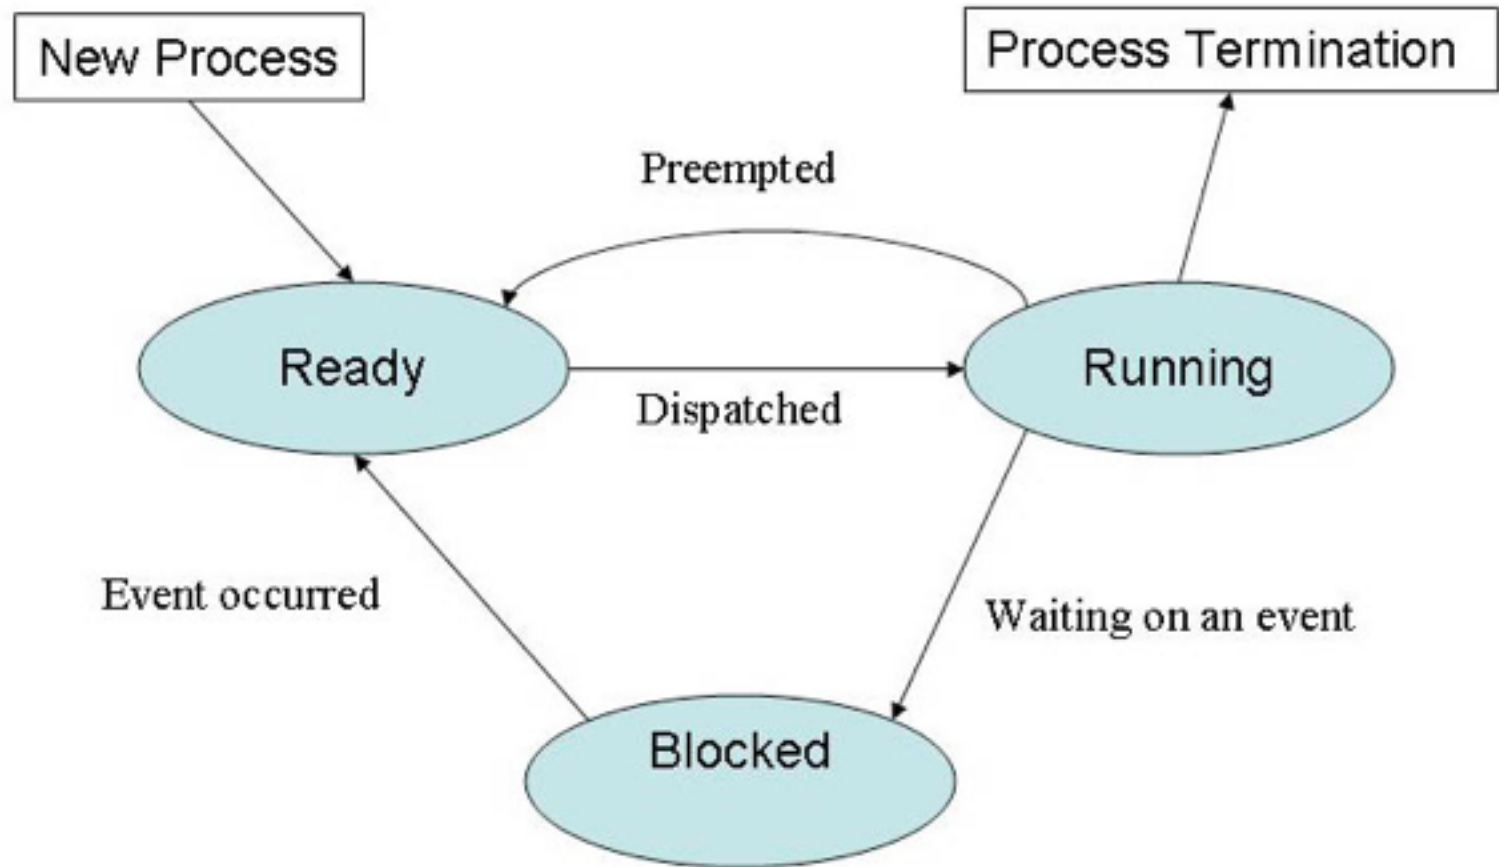
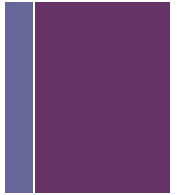




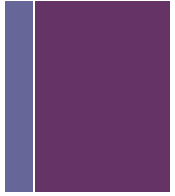
+

Process Control *con't*

+ Review: Process States

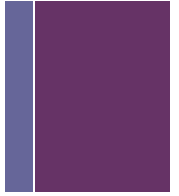


+ Review: `fork()`



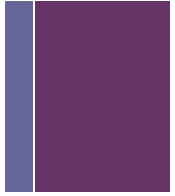
- Parent process creates a new running child process by calling `fork`
- `int fork(void)`
 - Returns 0 to the child process, child's PID to parent process
 - Child is *almost* identical to parent:
 - Child gets identical copy of the parent's virtual address space.
 - Child gets identical copies of the parent's open file descriptors
 - Child gets a different PID than the parent

+ Review: Process Graphs



- A *process graph* visualizes the partial ordering of statements in a concurrent program:
 - Each vertex is the execution of a statement
 - $a \rightarrow b$ means a happens *before* b
 - Edges can be labeled with current value of variables
 - `printf` vertices can be labeled with output
- Any sort of the graph corresponds to a feasible total ordering.

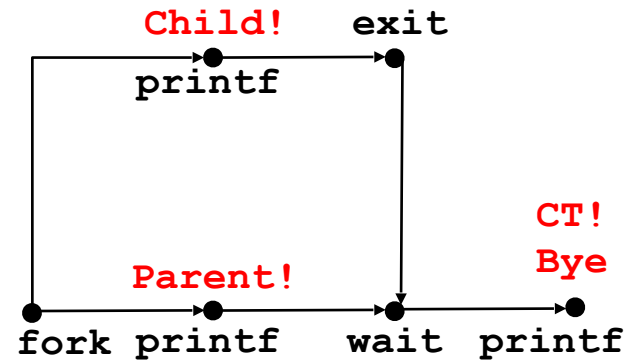
+ `wait()` : Synchronizing with Children



- Parent reaps a child by calling the wait function
- `int wait(int *child_status)`
 - Suspends current process until one of its children terminates
 - Return value is the pid of the child process that terminated
 - If `child_status != NULL`, then the integer it points to will be set to a value that indicates reason the child terminated and the exit status.

+wait(): Synchronizing with Children

```
void fork9() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("Child!");  
        exit(0);  
    } else {  
        printf("Parent!");  
        wait(&child_status);  
        printf("CT!");  
    }  
    printf("Bye\n");  
}
```



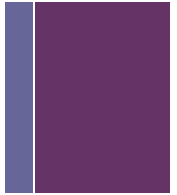
Feasible output:

Parent!
Child!
CT
Bye

Infeasible output:

Parent!
CT
Bye
Child!

+ Another wait() Example




- If multiple children completed, will take in arbitrary order
- Use WIFEXITED and WEXITSTATUS to get exit status

```
void fork10() {
    int pid[N];
    int i, child_status;

    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */

    for (i = 0; i < N; i++) { /* Parent */
        int child_pid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                  child_pid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", child_pid);
    }
}
```

+ waitpid(): Waiting for a Specific Process



- `int waitpid(pid_t pid, int &child_status, int options)`
 - Suspends current process until specific process terminates
 - Various options (see textbook)

```
void fork11() {
    int pid[N];
    int i, child_status;

    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */


    for (i = N-1; i >= 0; i--) {
        int child_pid = waitpid(pid[i], &child_status, 0);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                  child_pid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", child_pid);
    }
}
```


+ `execve()` : Loading and Running Programs

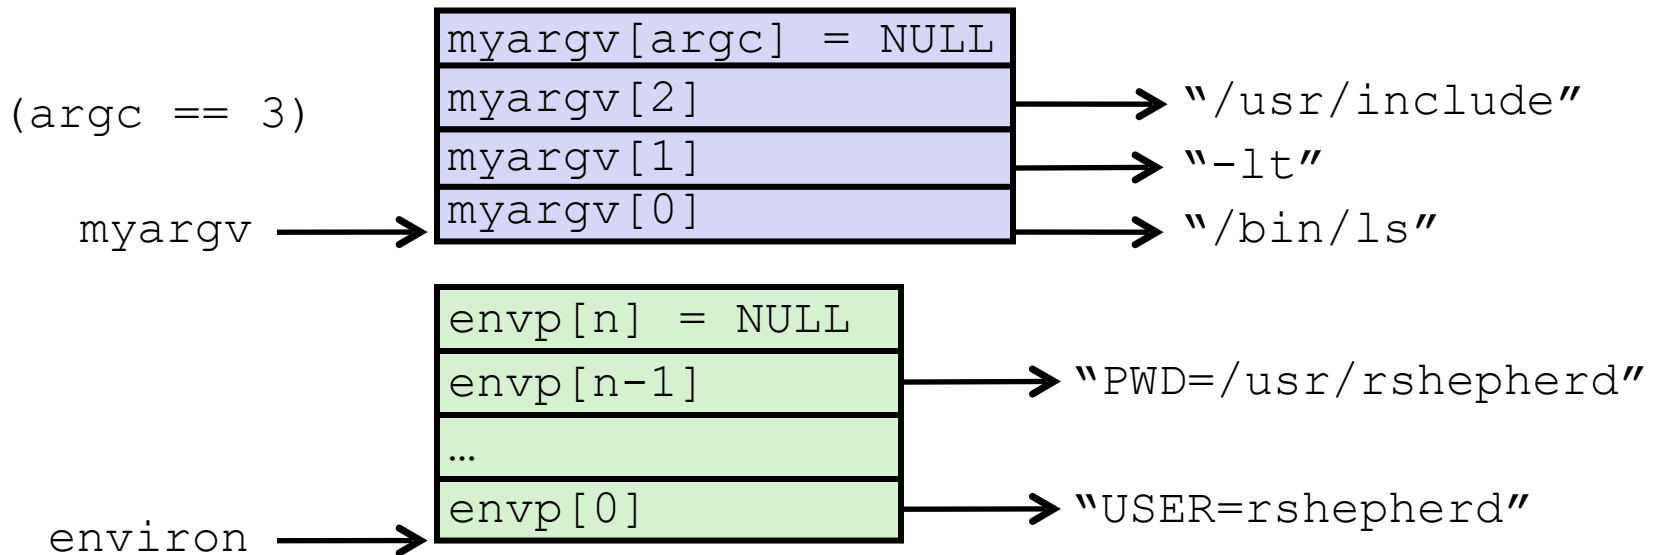
- `int execve(char* filename, char* argv[], char* envp[])`
- **Loads and runs in the current process:**
 - Executable file `filename`
 - ...with argument list `argv`
 - ...and environment variable list `envp`
 - “name=value” strings (e.g., `USER=rshpherd`)
 - `getenv()`, `putenv()`, `printenv()`
- **Overwrites code, data, and stack**
 - Retains PID, open files and signal context
- **Called once and never returns**
 - ...except if there is an error



+ execve Example

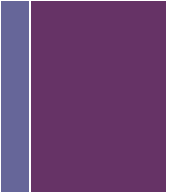


- Executes “/bin/l^s -lt /usr/include” in child process using current environment:



```
if ((pid = Fork()) == 0) {    /* Child runs program */
    if (execve(myargv[0], myargv, environ) < 0) {
        printf("%s: Command not found.\n", myargv[0]);
        exit(1);
    }
}
```


+ Summary



- **Processes**

- At any given time, system has multiple active processes
- Only one can execute at a time on a single core, though
- Each process appears to have total control of processor + private memory space

+ Summary *con't*




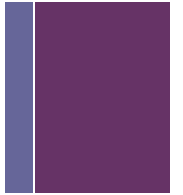
- **Creating processes**
 - Call `fork`
 - One call, two returns
- **Process completion**
 - Call `exit`
 - One call, no return
- **Reaping and waiting for processes**
 - Call `wait` or `waitpid`
- **Loading and running programs**
 - Call `execve` (or variant)
 - One call, (normally) no return



+

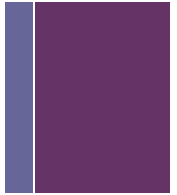
Signals

+ Linux Process Hierarchy



Note: you can view the hierarchy using the Linux `ps tree` command

+ Shell Programs



- A shell is an application program that runs programs on behalf of the user.
 - **sh** Original Unix shell (Stephen Bourne, AT&T Bell Labs, 1977)
 - **csch/tcsch** BSD Unix C shell
 - **bash** “Bourne-Again” Shell (default Linux shell)


```
int main()
{
    char cmdline[MAXLINE]; /* command line */

    while (1) {
        /* read */
        printf("> ");
        Fgets(cmdline, MAXLINE, stdin);
        if (feof(stdin))
            exit(0);

        /* evaluate */
        eval(cmdline);
    }
}
```

*Execution is a
sequence of read/
evaluate steps*

+ Simple Shell eval Function



```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list for execve() */
    int bg;               /* Should the job run in bg or fg? */
    int pid;              /* Process id */

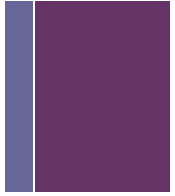
    bg = parseline(cmdline, argv); /* Extract arguments and set bg */

    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = Fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

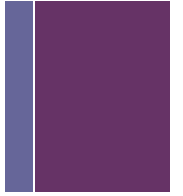
        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
        else
            printf("%d %s", pid, cmdline);
    }
    return;
}
```


+ Problem with Simple Shell Example



- **Our example shell correctly waits for and reaps foreground jobs**
- **But what about background jobs?**
 - Will become zombies when they terminate
 - Will never be reaped because shell (probably) will not terminate
 - Will create a memory leak that could run the kernel out of memory

+ Solution: Exceptional control flow



- **We can leverage exceptional control flow from our programs**
 - The kernel will interrupt regular processing to alert us when a background process completes
 - In Unix, the mechanism is called a **signal**


+ Signals



- A signal is a small message that notifies a process that an event of some type has occurred in the system
 - Akin to exceptions and interrupts
 - Sent from the kernel (sometimes at the request of another process)
 - Signal type is identified by integer ID's (1-30)
 - Only information in a signal is its ID and the fact that it arrived

<i>ID</i>	<i>Name</i>	<i>Default Action</i>	<i>Corresponding Event</i>
2	SIGINT	Terminate	User typed ctrl-c
9	SIGKILL	Terminate	Kill program (cannot override or ignore)
11	SIGSEGV	Terminate	Segmentation violation
14	SIGALRM	Terminate	Timer signal
17	SIGCHLD	Ignore	Child stopped or terminated

+ Signal Concepts: Sending a Signal

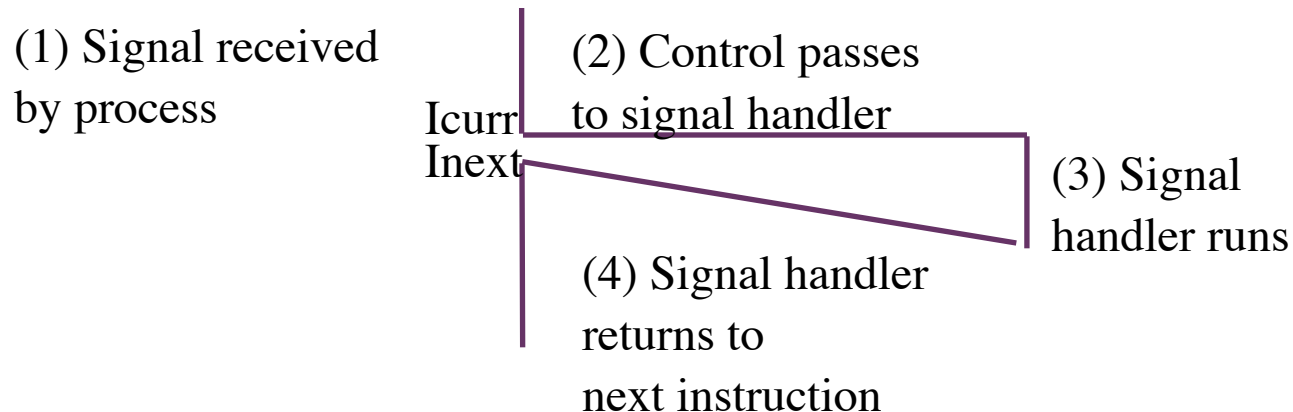


- **Kernel sends (delivers) a signal to a destination process by updating some state in the context of the destination process**
- **Kernel sends a signal for one of the following reasons:**
 - Kernel has detected a system event such as the termination of a child process (SIGCHLD)
 - Another process has invoked the `kill` system call to explicitly request the kernel to send a signal to the destination process


+ Signal Concepts: Receiving a Signal



- A destination process receives a signal when it is forced by the kernel to react in some way to the delivery of the signal
- Some possible ways to react:
 - Ignore the signal (do nothing)
 - Terminate the process (with optional core dump)
 - Catch the signal by executing a user-level function called *signal handler*
 - Akin to a hardware exception handler being called in response to an asynchronous interrupt

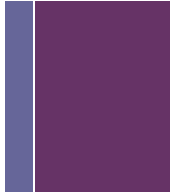


+ Signal Concepts: Pending & Blocked



- A signal is *pending* if *sent* but not yet *received*
 - There can be at most one pending signal of any particular type
 - Important: Signals are not queued
 - If a process has a pending signal of type *k*, then subsequent signals of type *k* that are sent to that process are discarded
- A process can *block* the receipt of certain signals
 - Blocked signals can be delivered, but will not be received until the signal is unblocked
- A *pending* signal is received at most once

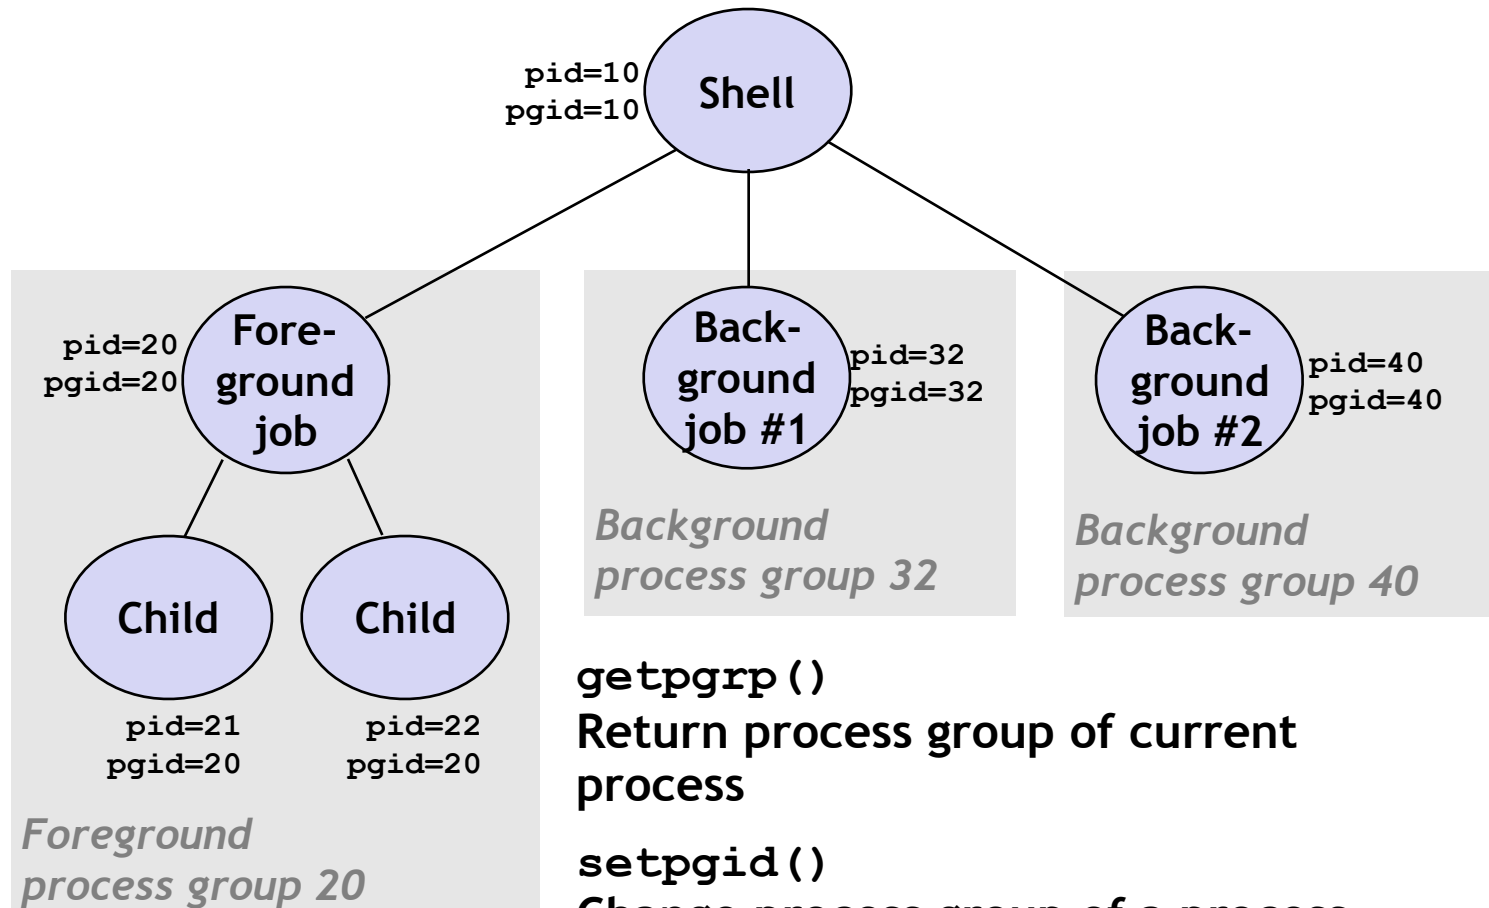
+ Signal Concepts: Pending/Blocked Bits



- Kernel maintains **pending** and **blocked** bit vectors in the context of each process
 - **pending**: represents the set of pending signals
 - Kernel sets bit `k` in pending when a signal of type `k` is delivered
 - Kernel clears bit `k` in pending when a signal of type `k` is received
 - **blocked**: represents the set of blocked signals
 - Can be set and cleared by using the `sigprocmask` function
 - Also referred to as the `signal mask`.

+ Sending Signals: Process Groups

- Every process belongs to exactly one process group




`getpgrp()`

Return process group of current process

`setpgid()`

Change process group of a process (see text for details)

+ Sending Signals with /bin/kill Program



- `/bin/kill` program sends specified signal to a process or process group
- **Examples**
 - `/bin/kill -9 24818`
Send SIGKILL to process 24818
 - `/bin/kill -9 -24817`
Send SIGKILL to every process in process group 24817

```
linux> ./forks 16
Child1: pid=24818 pgrp=24817
Child2: pid=24819 pgrp=24817

linux> ps
  PID TTY          TIME CMD
 24788 pts/2        00:00:00 tcsh
 24818 pts/2        00:00:02 forks
 24819 pts/2        00:00:02 forks
 24820 pts/2        00:00:00 ps
linux> /bin/kill -9 -24817
linux> ps
  PID TTY          TIME CMD
 24788 pts/2        00:00:00 tcsh
 24823 pts/2        00:00:00 ps
linux>
```

+ Sending Signals with `kill` Function


```
void fork12()
{
    pid_t pid[N];
    int i;
    int child_status;

    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0) {
            /* Child: Infinite Loop */
            while(1)
                ;
        }

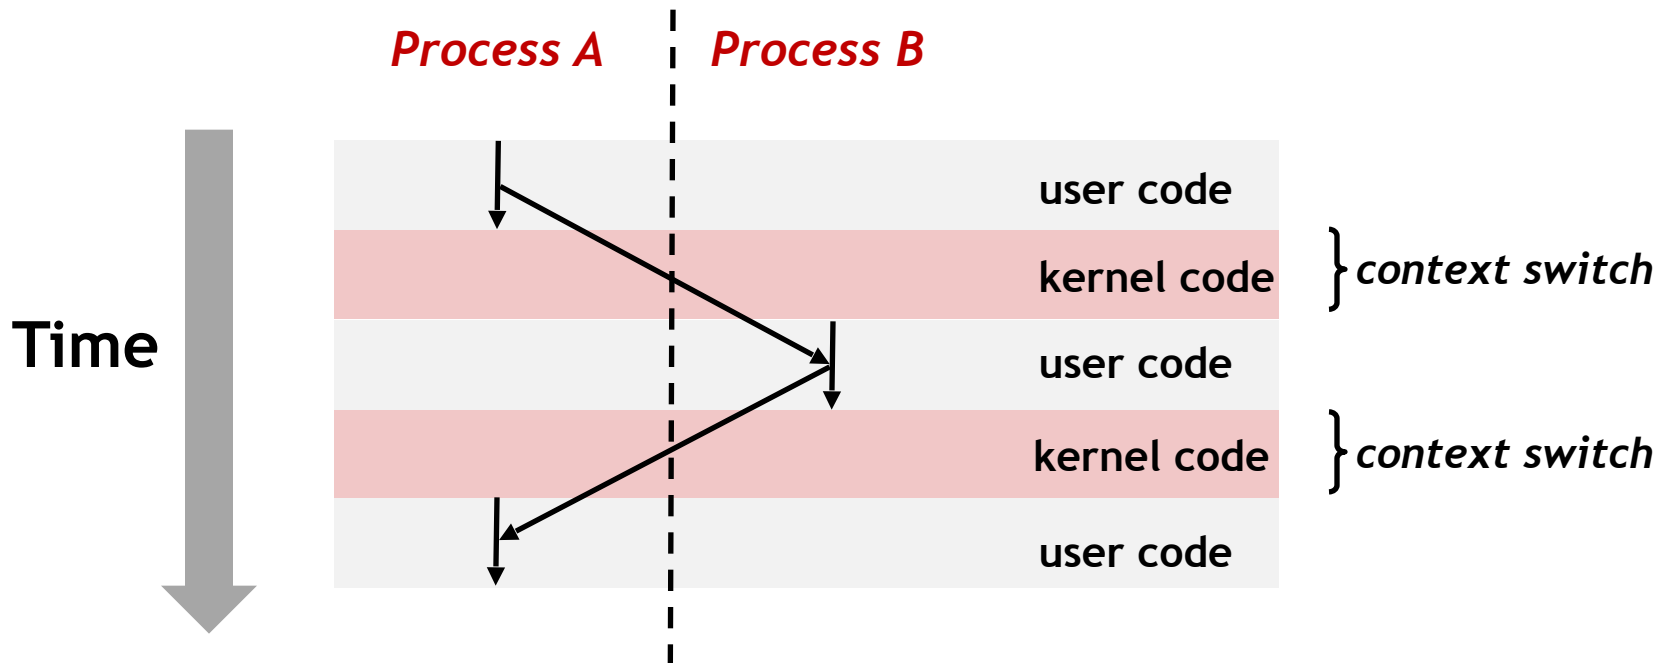
    for (i = 0; i < N; i++) {
        printf("Killing process %d\n", pid[i]);
        kill(pid[i], SIGINT);
    }

    for (i = 0; i < N; i++) {
        pid_t wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```

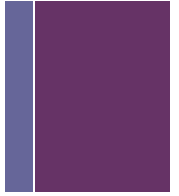
+ Receiving Signals



- Suppose kernel is returning from an exception handler and is ready to pass control to process *B*



+ Receiving Signals



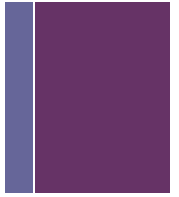
- Suppose kernel is returning from an exception handler and is ready to pass control to process B
- Kernel computes $pnb = \text{pending} \ \& \ \sim\text{blocked}$
 - The set of pending nonblocked signals for process B
- If $(pnb == 0)$
 - Pass control to next instruction in the logical flow for B
- Else
 - Choose nonzero bit k in pnb and force process B to receive signal k
 - The receipt of the signal triggers some action by B
 - Repeat for all nonzero bits in pnb
 - Pass control to next instruction in logical flow for B

+ Default Actions




- **Each signal type has a predefined default action, which is one of:**
 - The process terminates
 - The process stops until restarted by a SIGCONT signal
 - The process ignores the signal
- **What if we do not like the default action?**
 - Signal handlers

+ Installing Signal Handlers



- **The signal function modifies the default action associated with the receipt of signal `signum`:**
 - `handler_t* signal(int signum, handler_t* handler)`
- **Different values for handler:**
 - `SIG_IGN`: ignore signals of type `signum`
 - `SIG_DFL`: revert to the default action for signals of type `signum`
 - Otherwise, `handler` is the address of a user-level *signal handler*
 - Called when process receives signal of type `signum`
 - When the handler executes its return statement, control passes back to instruction in the control flow of the process that was interrupted by receipt of the signal.
- **Returns the previous value of the signal handler, or `SIG_ERR` on error.**

+ Signal Handling Example



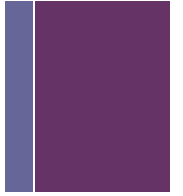
```
void sigint_handler(int sig) /* SIGINT handler */ {
    printf("So you think you can stop the bomb with ctrl-c?\n");
    sleep(2);
    printf("Well...");
    fflush(stdout);
    sleep(1);
    printf("OK. :-)\n");
}

int main() {
    /* Install the SIGINT handler */
    if (signal(SIGINT, sigint_handler) == SIG_ERR)
        unix_error("signal error");

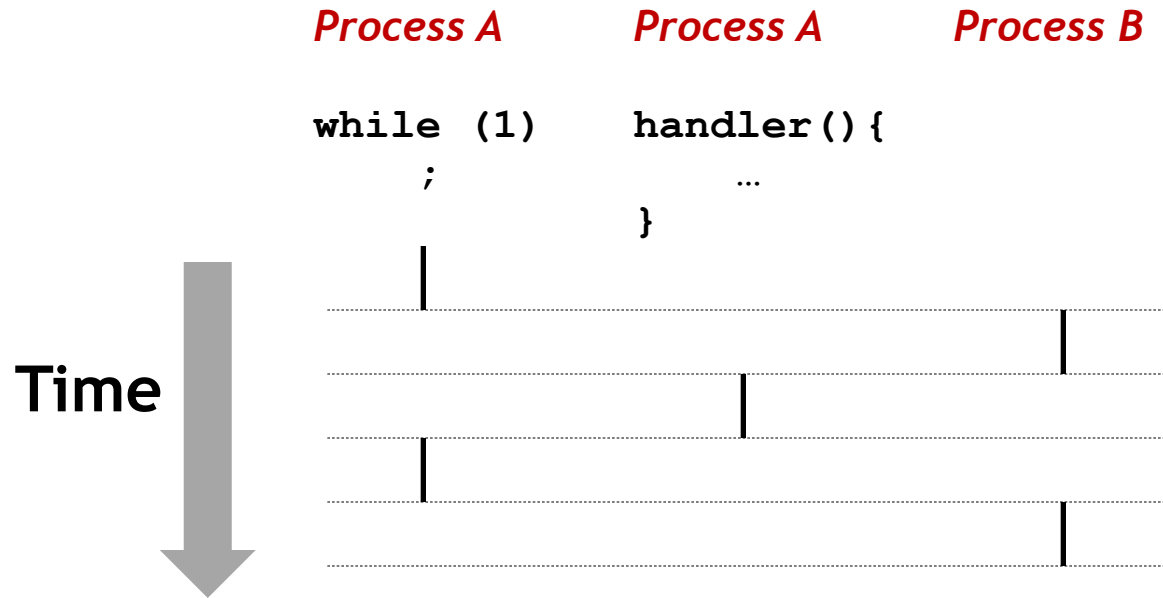
    /* Wait for the receipt of a signal */
    pause();

    return 0;
}
```

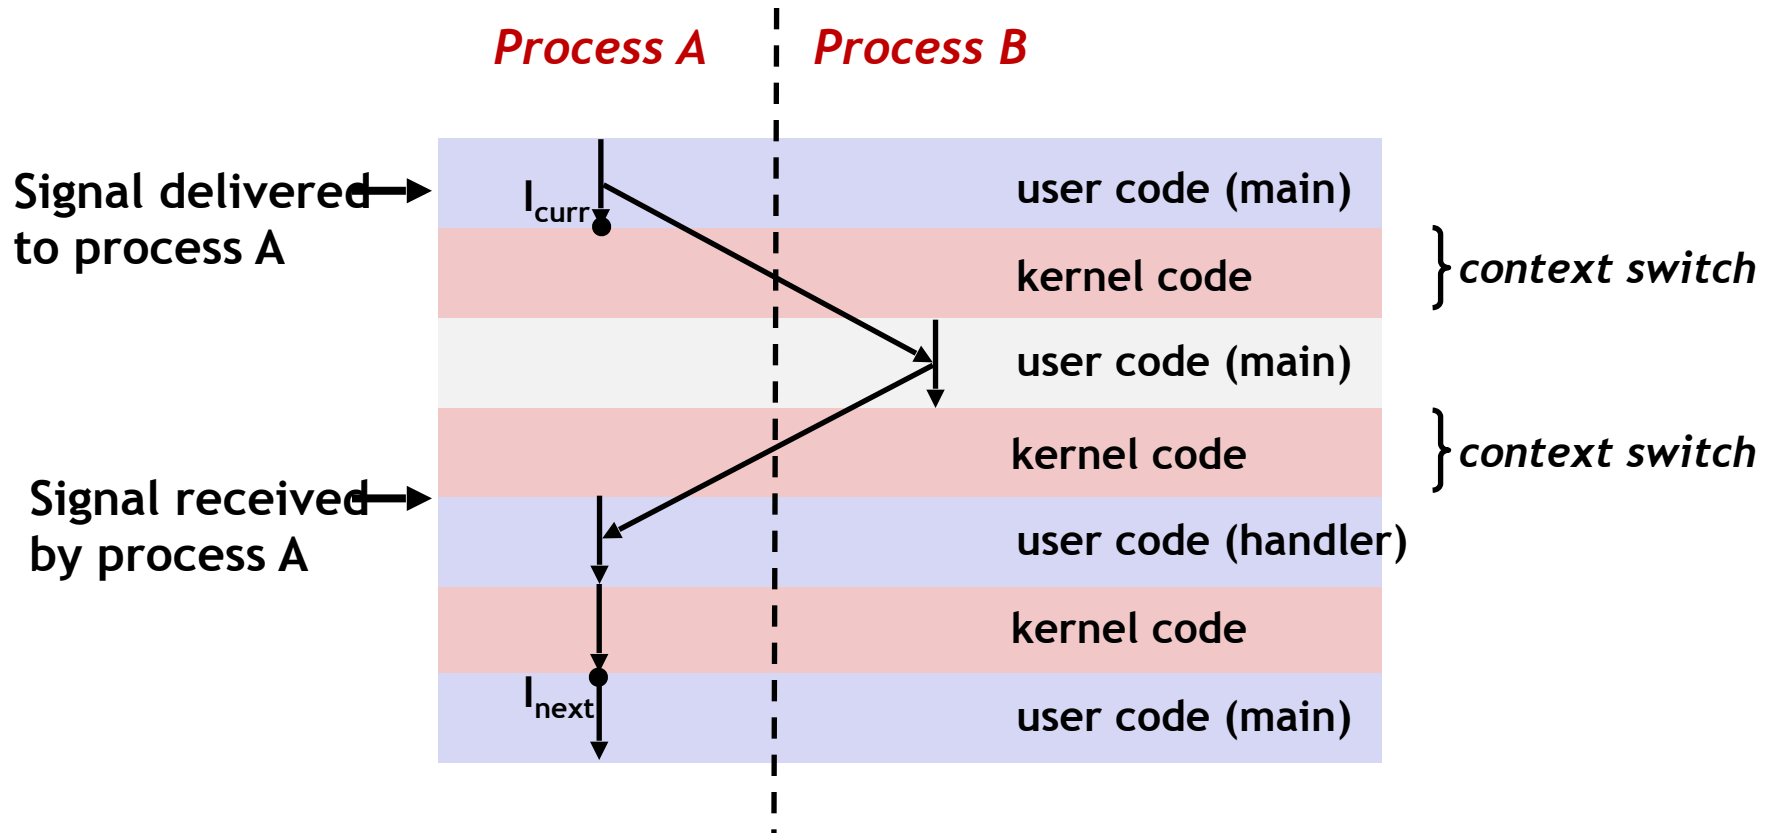
+ Signals Handlers as Concurrent Flows



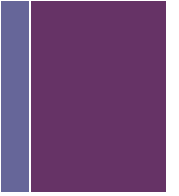
- A signal handler is a separate logical flow (not process) that runs concurrently with the main program



+ Signal Handlers as Concurrent Flows *con't*



+ Concurrency Issue with Signal Handlers



```
int N = 999;

void sigint_handler(int sig) /* SIGINT handler */ {
    x = rand() % 2 == 0 ? x : x+=1 ;
}

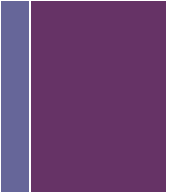
int main() {
    int array[N];
    if (signal(SIGINT, sigint_handler) == SIG_ERR)
        unix_error("signal error");

    pause();

    printf("%d", array[N]);

    return 0;
}
```

+ Summary



- **Signals provide process-level exception handling**
 - Can generate from user programs
 - Can define effect by declaring signal handler
- **Handlers are tricky because they are concurrent with main program and share the same global data structures.**
 - Shared data structures can become corrupted.
- **We'll explore concurrency issues later in the term.**