



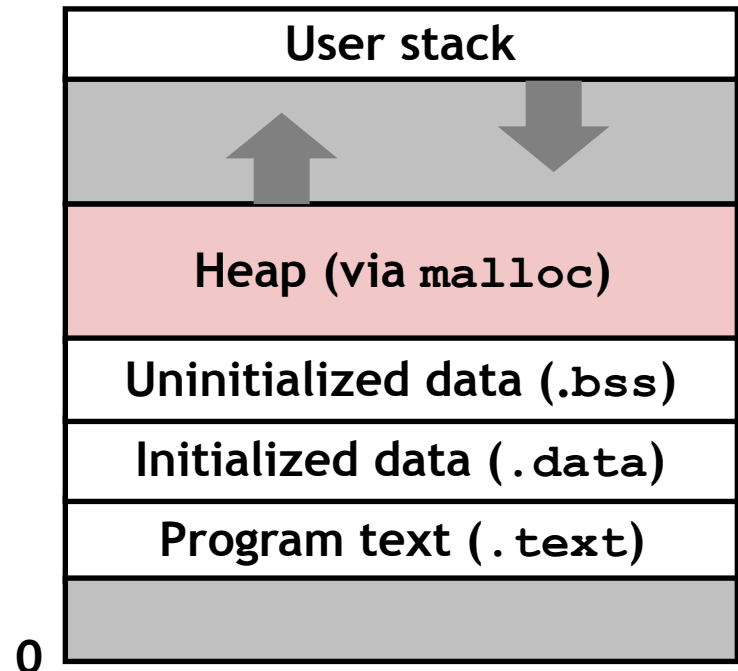

+

Dynamic Allocation

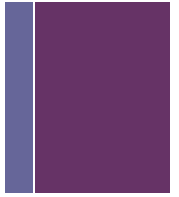
+ Dynamic Memory Allocation



- **Programmers use dynamic memory allocators (such as malloc) to acquire VM at run time.**
 - For data structures whose size is only known at runtime.
- **Dynamic memory allocators manage an area of process virtual memory known as the heap.**

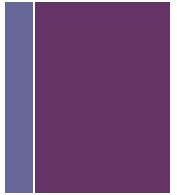


+ Dynamic Memory Allocation



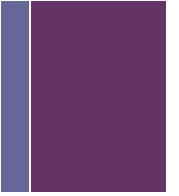
- **Allocator maintains heap as collection of variable sized *blocks*, which are either *allocated* or *free***
- **Types of allocators**
 - *Explicit* allocator: application allocates and frees space
 - E.g., `malloc` and `free` in C
 - *Implicit* allocator: application allocates, but does not free space
 - E.g. garbage collection in Java, ML, and Lisp
- **Will discuss simple explicit memory allocation today**

+ The malloc Package (*review*)



- `#include <stdlib.h>`
- `void *malloc(size_t size)`
 - Successful:
 - Returns a pointer to a memory block of at least size bytes aligned to a 16-byte boundary
 - If size == 0, returns NULL
 - Unsuccessful: returns NULL (0) and sets errno
- `void free(void *p)`
 - Releases the block pointed at by p to pool of available memory
 - p must come from a previous call to malloc or calloc
- **Other functions**
 - calloc: Version of malloc that initializes allocated block to zero.
 - realloc: Changes the size of a previously allocated block.
 - sbrk: Used internally by allocators to grow or shrink the heap

+ malloc Example (*review*)



```
#include <stdio.h>
#include <stdlib.h>

void foo(int n) {
    int i, *p;

    /* Allocate a block of n ints */
    p = (int *) malloc(n * sizeof(int));
    if (p == NULL) {
        perror("malloc");
        exit(0);
    }


    /* Initialize allocated block */
    for (i=0; i<n; i++)
        p[i] = i;

    /* Return allocated block to the heap */
    free(p);
}
```

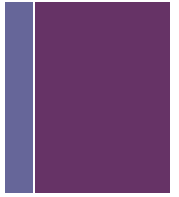
+ Assumptions Made in This Lecture



- Memory is word addressed.
- Words are int-sized.




+ Allocation Example




`p1 = malloc(4)`




`p2 = malloc(5)`



`p3 = malloc(6)`



`free(p2)`

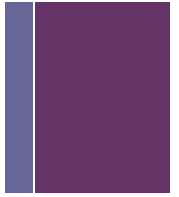


`p4 = malloc(2)`



(Arguments to malloc are in words for simplification, i.e. malloc(4) allocates 4 words.)

+ Constraints



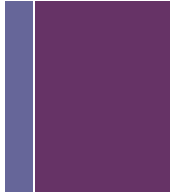
■ Applications

- Can issue arbitrary sequence of `malloc` and `free` requests
- `free` request must be to a `malloc`'d block

■ Allocators


- Can't control number or size of allocated blocks
- Must respond immediately to `malloc` requests (can't reorder or buffer requests)
- Must allocate blocks from *free* memory
- Can manipulate and modify only *free* memory
- Can't move the allocated blocks once they are `malloc`'d
- Must align blocks so they satisfy all alignment requirements
 - 16-byte alignment on Linux boxes

+ Performance Goal: Throughput



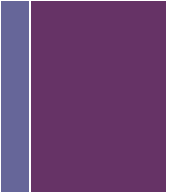
- **Given some sequence of `malloc` and `free` requests:**
 - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- **Goals: maximize throughput and peak memory utilization**
 - These goals are often conflicting
- **Throughput:**
 - Number of completed requests per unit time
 - Example:
 - 5,000 `malloc` calls and 5,000 `free` calls in 10 seconds
 - Throughput is 1,000 operations/second

+ Performance Goal: Peak Memory Utilization



- **Given some sequence of `malloc` and `free` requests:**
 - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- **Definition: *Aggregate payload* P_k**
 - `malloc(p)` results in a block with a payload of p bytes
 - After request R_k has completed, the aggregate payload P_k is the sum of currently allocated payloads
- **Definition: *Current heap size* H_k**
 - Assume H_k is monotonically nondecreasing
 - i.e., heap grows when allocator uses `sbrk`
- **Definition: *Peak memory utilization* (after k requests)**
 - $U_k = P_k / H_k$

+ Fragmentation




- **Poor memory utilization caused by fragmentation**
 - internal fragmentation
 - external fragmentation

+ Internal Fragmentation



- For a given block, internal fragmentation occurs if payload is smaller than block size



- Caused by
 - Overhead of maintaining heap data structures
 - Padding for alignment purposes
 - Explicit policy decisions
(e.g., to return a big block to satisfy a small request)
- Depends only on the pattern of previous requests
 - Thus, easy to measure

+ External Fragmentation




- Occurs when there is enough aggregate heap memory, but no single free block is large enough


```
p1 = malloc(4)
```



```
p2 = malloc(5)
```



```
p3 = malloc(6)
```



```
free(p2)
```

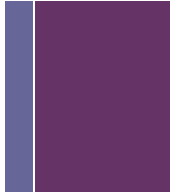


```
p4 = malloc(6)
```

Oops! (what would happen now?)

- Depends on the pattern of future requests
 - Thus, difficult to measure

+ Implementation Issues




- How do we know how much memory to free given just a pointer?
- How do we keep track of the free blocks?
- What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in? Do we split blocks?
- How do we pick a block to use for allocation -- many might fit?
- How do we reinsert freed block?

+ Knowing How Much to Free



- **Standard method**


- Keep the length of a block in the word preceding the block.
 - This word is often called the *header field* or *header*
- Requires an extra word for every allocated block




+ Keeping Track of Free Blocks



- **Method 1: Implicit list using length—links all blocks**



- **Method 2: Explicit list among the free blocks using pointers**

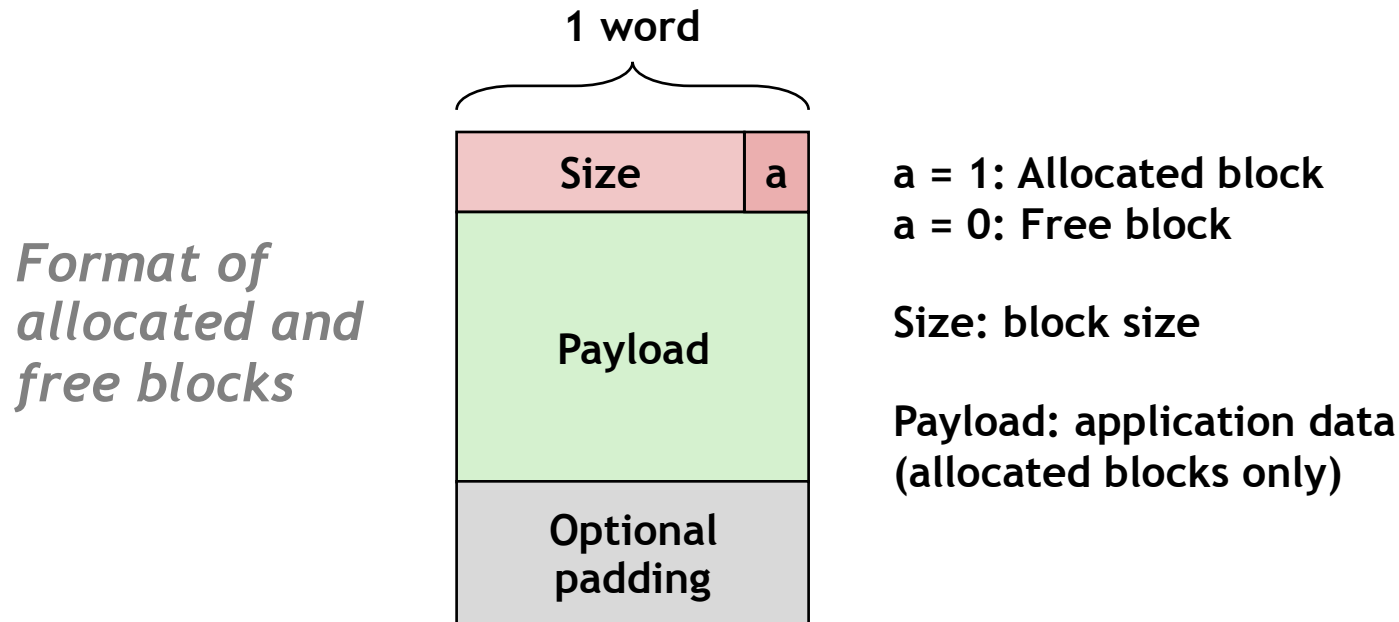


- **Method 3: Segregated free list**
 - Different free lists for different size classes


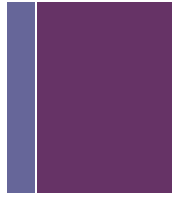
+ Method 1: Implicit List



- For each block we need both size and allocation status
 - Could store this information in two words: wasteful!
- Standard trick
 - If blocks are aligned, some low-order address bits are always 0
 - Instead of storing an always-0 bit, use it as a allocated/free flag
 - When reading size word, must mask out this bit



+ Detailed Implicit Free List Example



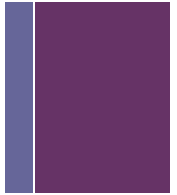
Double-word
aligned

Allocated blocks: shaded

Free blocks: unshaded

Headers: labeled with size in
bytes/allocated bit

+ Implicit List: Finding a Free Block



- **First fit:**

- Search list from beginning, choose first free block that fits:

```
p = start;
while ((p < end) &&          \\ not passed end
      ((*p & 1) ||          \\ already allocated
      (*p <= len)))        \\ too small
  p = p + (*p & -2);        \\ goto next block (word addressed)
```

- Can take linear time in total number of blocks (allocated and free)
- In practice it can cause “splinters” at beginning of list

- **Next fit:**

- Like first fit, but search list starting where previous search finished
- Should often be faster than first fit: avoids re-scanning unhelpful blocks
- Some research suggests that fragmentation is worse


- **Best fit:**

- Search the list, choose the best free block: fits, with fewest bytes left over
- Keeps fragments small—usually improves memory utilization
- Will typically run slower than first fit


+ Implicit List: Allocating in Free Block



- **Allocating in a free block: splitting**
 - Allocated space might be smaller than free space, maybe split block



`addblock(p, 4)`



```
void addblock(ptr p, int len) {  
    int newsize = ((len + 1) >> 1) << 1; // round up to even  
    int oldsize = *p & -2;                // mask out low bit  
    *p = newsize | 1;                     // set new length  
    if (newsize < oldsize)  
        *(p+newsize) = oldsize - newsize; // set length in remaining  
}
```


+ Implicit List: Freeing a Block

- Simplest implementation:

- Need only clear the “allocated” flag

- `void free_block(ptr p) { *p = *p & -2 }`

- But can lead to “false fragmentation”




Oops! There is enough free space, but the allocator won't be able to find it

+ Implicit List: Coalescing



- Join (coalesce) with next/previous blocks, if they are free
 - Coalescing with next block




- But how do we coalesce with previous block?


+ Implicit List: Bidirectional Coalescing



- **Boundary tags [Knuth '73]** https://en.wikipedia.org/wiki/Donald_Knuth
 - Replicate size/allocated word at “bottom” (end) of free blocks
 - Allows us to traverse the “list” backwards, but requires extra space
 - Important and general technique!



*Format of
allocated and
free blocks*

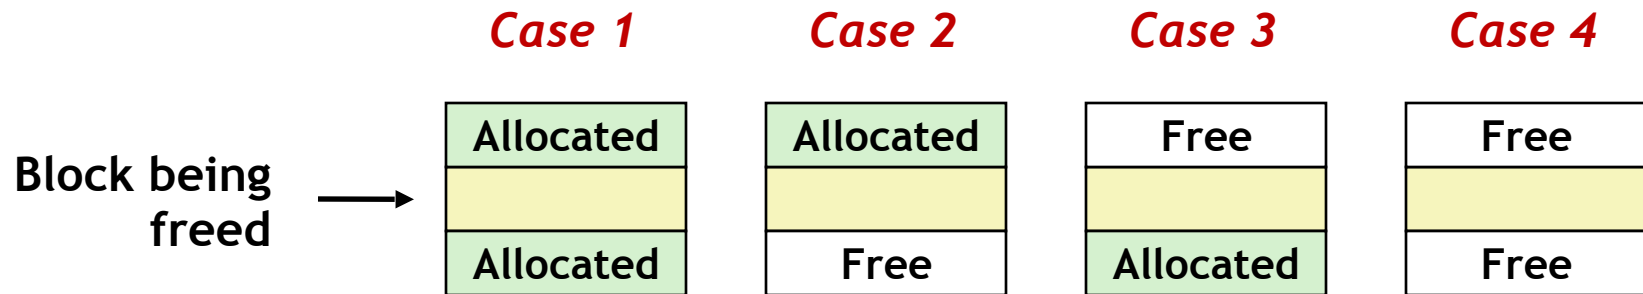


a = 1: Allocated block
a = 0: Free block


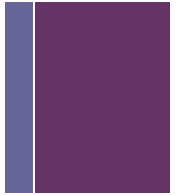
Size: Total block size

Payload: Application data
(allocated blocks only)


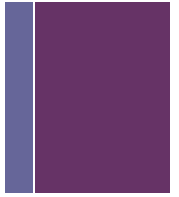
+ Constant Time Coalescing





+ Constant Time Coalescing (Case 1)





+ Constant Time Coalescing (Case 2)



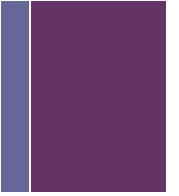
+ Constant Time Coalescing (Case 3)



+ Constant Time Coalescing (Case 4)

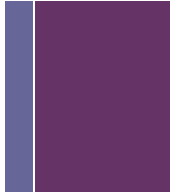


+ Disadvantages of Boundary Tags




- **Internal fragmentation**
- **Can it be optimized?**
 - Which blocks need the footer tag?

+ Disadvantages of Boundary Tags




- **Internal fragmentation**
- **Can it be optimized?**
 - Which blocks need the footer tag?
 - Only free blocks!
- **So how do we know if the last word in the previous block is a boundary tag or not, after all its just bits back there!**

+ Disadvantages of Boundary Tags



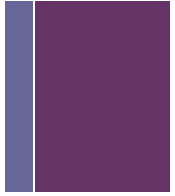
- **Internal fragmentation**
- **Can it be optimized?**
 - Which blocks need the footer tag?
 - Only free blocks!
- **So how do we know if the last word in the previous block is a boundary tag or not, after all its just bits back there!**
 - We can use one of those low order bits in the header to indicate the allocation status of the previous block.

+ Summary of Key Allocator Policies



- **Placement policy:**
 - *First-fit, next-fit, best-fit, etc.*
 - Trades off lower throughput for less fragmentation
- **Splitting policy:**
 - When do we go ahead and split free blocks?
 - How much internal fragmentation are we willing to tolerate?
- **Coalescing policy:**
 - *Immediate coalescing*: coalesce each time free is called
 - *Deferred coalescing*: improve performance of `free` by deferring coalescing.
 - Ex. Coalesce as you scan the free list for `malloc`
 - Ex. Coalesce when the amount of external fragmentation reaches some threshold

+ Implicit Lists: Summary



- **Implementation:** very simple
- **Allocate cost:**
 - linear time worst case
- **Free cost:**
 - constant time worst case
 - even with coalescing
- **Memory usage:**
 - will depend on placement policy
 - First-fit, next-fit or best-fit
- **Not used in practice for malloc/free because of linear-time allocation**
- **Concepts of splitting and boundary tag coalescing are general to all allocators**