




# Dynamic Memory Allocation *cont't*

# + Implicit List: Coalescing



- Join (coalesce) with next/previous blocks, if they are free
  - Coalescing with next block




- But how do we coalesce with previous block?


# + Implicit List: Boundary Tags (footers)



- **Boundary tags [Knuth '73]** [https://en.wikipedia.org/wiki/Donald\\_Knuth](https://en.wikipedia.org/wiki/Donald_Knuth)
  - Replicate size/allocated word at “bottom” (end) of free blocks
  - Allows us to traverse the “list” backwards, but requires extra space



*Format of  
allocated and  
free blocks*

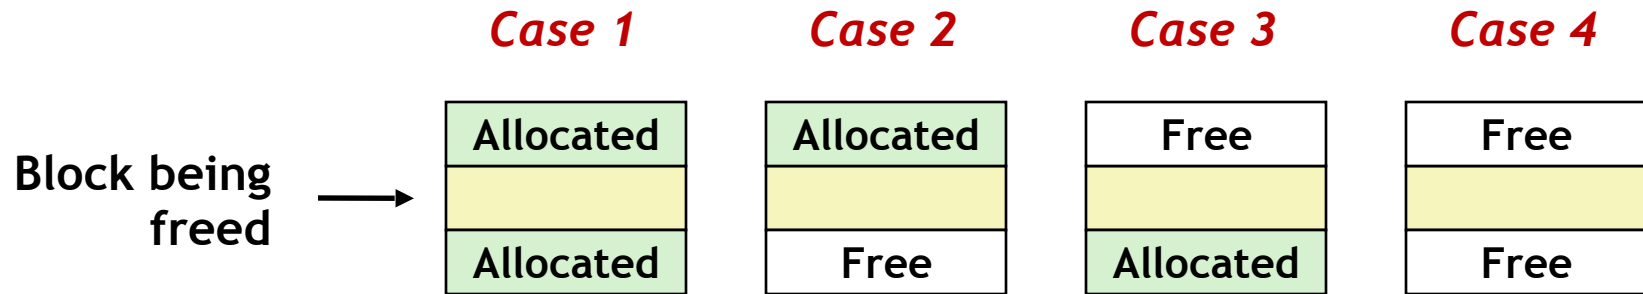


a = 1: Allocated block  
a = 0: Free block


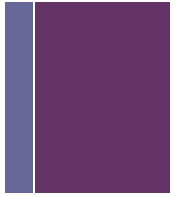
Size: Total block size

Payload: Application data  
(allocated blocks only)


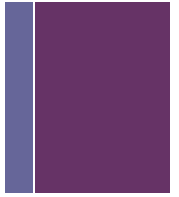
# + Constant Time Coalescing




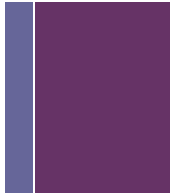
# + Constant Time Coalescing (Case 1)




## + Constant Time Coalescing (Case 2)



## + Constant Time Coalescing (Case 3)

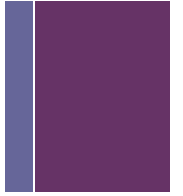


Frequency	Percentage
Daily	90%
Not daily	10%






# + Disadvantages of Boundary Tags



- **Internal fragmentation**
  - Again we are trading space for time, utilization for throughput
- **Can it be optimized?**
  - Which blocks need the footer tag?

# + Disadvantages of Boundary Tags




- **Internal fragmentation**
  - Again we are trading space for time, utilization for throughput
- **Can it be optimized?**
  - Which blocks need the footer tag?
  - Only free blocks!
- **So how do we know if the last word in the previous block is a boundary tag or not, after all its just bits back there!**

# + Disadvantages of Boundary Tags



- **Internal fragmentation**
  - Again we are trading space for time, utilization for throughput
- **Can it be optimized?**
  - Which blocks need the footer tag?
  - Only free blocks!
- **So how do we know if the last word in the previous block is a boundary tag or not, after all its just bits back there!**
  - We can use one of those low order bits in the header to indicate the allocation status of the previous block.

# + Implicit Lists: Summary



- **Implementation: very simple**
- **Allocate cost:**
  - linear time
- **Free cost:**
  - constant time (even with coalescing)
- **Memory usage:**
  - Will depend on placement policy
    - First-fit, next-fit or best-fit
- **Not used in practice for malloc/free because of linear-time allocation**
- **Concepts of splitting and boundary tag coalescing are general to all allocators**




# Explicit Free Lists


# + Keeping Track of Free Blocks



- **Method 1: Implicit free list using length—links all blocks**



- **Method 2: Explicit free list among the free blocks using pointers**




- **Method 3: Segregated free list**
  - Different free lists for different size classes


# + Explicit Free Lists



Allocated (as before)



Free



- Maintain list(s) of *free* blocks, not *all* blocks
  - The “next” free block could be anywhere
    - So we need to store forward/back pointers, not just sizes
  - Still need boundary tags for coalescing
  - Luckily we track only free blocks, so we can use payload area


# + Explicit Free Lists



- Logically:




- Physically: blocks can be in any order






# + Allocating From Explicit Free Lists




*Before*




*After*

*(with splitting)*

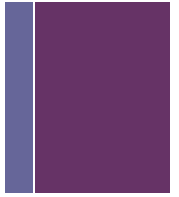


# + Freeing With Explicit Free Lists



- **Insertion policy: Where in the free list do you put a newly freed block?**
- **LIFO (last-in-first-out) policy**
  - Insert freed block at the beginning of the free list
  - **Pro:** simple and constant time
  - **Con:** studies suggest fragmentation is worse than address-ordered
- **Address-ordered policy**
  - Insert freed blocks so that free list blocks are always in address order:
$$\text{addr}(\text{prev}) < \text{addr}(\text{curr}) < \text{addr}(\text{next})$$
  - **Con:** requires search
  - **Pro:** studies suggest fragmentation is lower than LIFO

# + Explicit List Summary

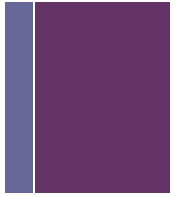


- **Comparison to implicit list:**
  - Allocate is linear time in number of *free* blocks instead of all
    - Much faster when good heap utilization is maintained
  - Slightly more complicated allocate and free since needs to splice blocks in and out of the list
  - Some extra space for the links (2 extra words needed per block)
    - Does this increase internal fragmentation?
- **Most common use of explicit free lists is in context of segregated free lists**




# Segregated Free Lists


# + Keeping Track of Free Blocks



- **Method 1: Implicit list using length—links all blocks**

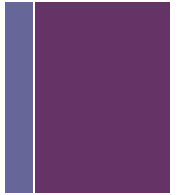


- **Method 2: Explicit list among the free blocks using pointers**




- **Method 3: Segregated free list**
  - Different free lists for different size classes

# + Size Classes

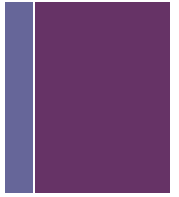


- Each *size class* of blocks has its own free list




- Often have separate classes for each small size
- For larger sizes: One class for each two-power size

# + Segregated List Allocator



- **Given an array of free lists, each one for some size class**
- **To allocate a block of size  $n$ :**
  - Search appropriate free list for block of size  $m > n$
  - If an appropriate block is found:
    - Split block and place fragment on appropriate list (optional)
  - If no block is found, try next larger class
  - Repeat until block is found
- **If no block is found:**
  - Request additional heap memory from OS (using `sbrk()`)
  - Allocate block of  $n$  bytes from this new memory
  - Place remainder as a single free block in some size class.

# + Segregated List Allocator *con't*



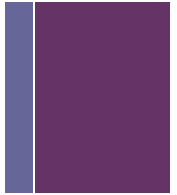
- **To free a block:**
  - Coalesce and place on appropriate list
    - Common technique is to place it at the end of the list, so constant time.
- **Advantages of segregated list allocators**
  - *Higher throughput*
    - Size class lists are a smaller search space
  - *Better memory utilization*
    - *First-fit* search of segregated free list approximates a *best-fit* search of entire heap.





# Garbage Collection

# + Garbage Collection



- Automatic free'ing of heap-allocated storage no longer in use

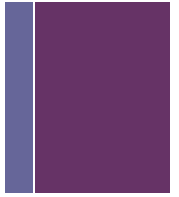
```
void foo() {  
    int* p = malloc(128);  
    return; /* p block is now garbage */  
}
```

- Common in many languages:
  - Python, Ruby, Java, Go, ML, Lisp....
- Variants exist for C and C++
  - However, pointer semantics make it impossible for it be perfect

```
void main() {  
    int* p = malloc(128);  
    int* q = p+32; //Confuses collector  
    // other code  
    return;  
}
```



# + Garbage



- **How does the memory manager know when memory can be freed?**
  - We can tell that certain blocks cannot be used if there are no pointers to them
- **Must make certain assumptions about pointers**
  - Memory manager can distinguish pointers from non-pointers
  - All pointers point to the start of a block
  - Cannot hide pointers  
(e.g., by coercing them to an int, and then back again)
  - All these things are true in Java but not in C

# + Classical GC Algorithms




- **Mark-and-sweep collection (McCarthy, 1960)**
  - Does not move blocks (unless you also “compact”)
- **Reference counting (Collins, 1960)**
  - Does not move blocks (not discussed)
- **Copying collection (Minsky, 1963)**
  - Moves blocks (not discussed)
- **Generational Collectors (Lieberman and Hewitt, 1983)**
  - Collection based on lifetimes
    - Most allocations become garbage very soon
    - So focus reclamation work on zones of memory recently allocated
- **For more information:**  
**Jones and Lin, “Garbage Collection: Algorithms for Automatic Dynamic Memory”, John Wiley & Sons, 1996.**

# + Memory as a Graph



- We view memory as a directed graph
  - Each block is a node in the graph , each pointer is an edge in the graph
  - Locations not in the heap that contain pointers into the heap are called root nodes (e.g. registers, locations on the stack, global variables)





A node (block) is **reachable** if there is a path from any root to that node.  
Non-reachable nodes are **garbage** (cannot be needed by the application)

# + Mark and Sweep Collecting




- Can build on top of malloc/free package
  - Allocate using `malloc` until you “run out of space”
- When out of space:
  - Use extra *mark bit* in the head of each block
  - *Mark*: Start at roots and set mark bit on each reachable block
  - *Sweep*: Scan all blocks and free blocks that are not marked

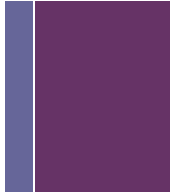


Note: arrows here denote memory refs, not free list ptrs.

 Mark bit set



# + Mark and Sweep (cont.)



## Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {  
    if (!is_ptr(p)) return;           // do nothing if not pointer  
    if (markBitSet(p)) return;        // check if already marked  
    setMarkBit(p);                    // set the mark bit  
    for (i=0; i < length(p); i++)    // call mark on all words  
        mark(p[i]);                  // in the block  
    return;  
}
```

## Sweep using lengths to find next block

```
ptr sweep(ptr p, ptr end) {  
    while (p < end) {  
        if markBitSet(p)  
            clearMarkBit();  
        else if (allocateBitSet(p))  
            free(p);  
        p += length(p);  
    }  
}
```

# + Conservative Collection in C



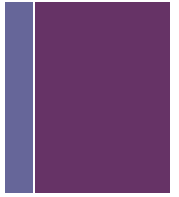
- A “conservative garbage collector” for C programs
  - `is_ptr()` determines if a word is a pointer by checking if it points to an allocated block of memory
  - But, in C pointers can point to the middle of a block



- So how to find the beginning of the block?
  - Keep a balance tree where the key is the starting address of every allocation.
  - Search tree and see if some pointer falls within the  $\text{key} + \text{size}$  range of addresses.
  - Could yield false positives! Hence “conservative”.


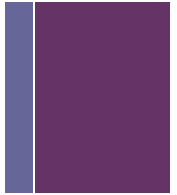


# + True Collection in Java: Copying

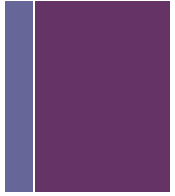


- Java uses a different technique than ‘Mark & Sweep’ called ‘Copying’
- The heap is split into two parts: FROM space and TO space
- Objects are allocated in the FROM space
  - When FROM is full, collection begins
- During traversal, each reachable object is copied to TO space
  - When traversal is done, all live objects in are in TO space
- Now the spaces are flipped, FROM becomes TO and vice versa
  - Everything in FROM gets freed.

# + True Collection in Java: Copying *con't*




# + True Collection in Java: Generational



- A variant of ‘Copying’.
- *Infant mortality* or the *generational hypothesis* is the observation that, in most cases, young objects are much more likely to die than old objects. Why?
  - Objects that live for a long time tend to make up core program data structures and will probably live until the end of the programs life.
- It turns out that the vast majority of data in typical programs (between 92 and 98 percent according to various studies), *die young*.
  - Moreover, most variables are short-lived.

# + True Collection in Java: Generational *con't*



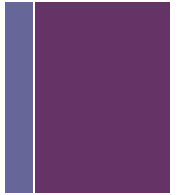
- Generational GCs exploit this ‘generational hypothesis’
- Instead of just two heaps (FROM and TO), we have several signifying ‘generations’ of objects.
  - Younger generations collected more frequently than older generations (because younger generations will have more garbage to collect)
  - When a generation is traversed, live objects are copied to the next-older generation
  - When a generation fills up, it is garbage collected.

# + Urban Performance Legends



- **"Garbage collection will never be as efficient as manual memory management." - Snooty C Programmer**
  - In a way, those statements are right -- automatic memory management is not as fast -- it's often considerably faster.
  - Explicit allocators deal with blocks of memory *one at a time*, whereas the garbage collection approach tends to deal with memory management in large batches, yielding more opportunities for optimization.

# + Urban Performance Legends *con't*



- **Allocation in JVMs was not always so fast -- early JVMs indeed had poor allocation and garbage collection performance.**
  - A lot has happened since the JDK 1.0 days; the introduction of generational collectors in JDK 1.2 has greatly improved performance.
- **As a result, for most objects, the garbage collection cost is -- zero.**
  - This is because a copying collector does not need to visit or copy dead objects, only live ones. So objects that become garbage shortly after allocation contribute no workload to the collection cycle.