



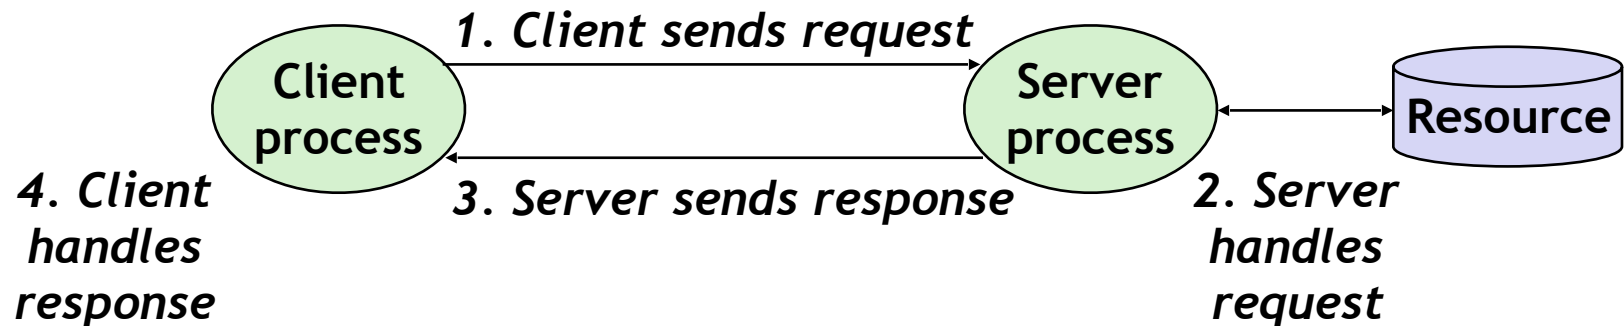
+

Network Programming

+ A Client-Server Transaction

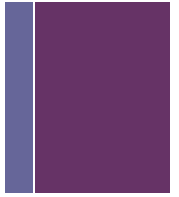


- Many network applications are based on the client-server model:
 - A server *process* and one or more client processes
 - Server manages some *resource*
 - Server provides service by manipulating resource for clients
 - Server activated by *request* from client



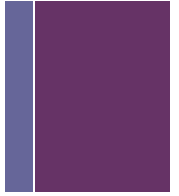
Note: clients and servers are processes running on hosts (can be the same or different hosts)

+ Computer Networks



- **A network is a group of connected systems that are able to communicate in order to exchange data.**
- **There are many kinds of networks, some examples..**
 - *LAN* (Local Area Network) spans a building or campus
 - *Ethernet* is most prominent example
 - *WAN* (Wide Area Network) spans country or world
 - Typically high-speed *point-to-point telecom lines*
- **An internetwork (internet) is an interconnected set of networks**
 - The Global IP Internet (uppercase “I”) is the most famous example of an internet (lowercase “i”)

+ Computer Networks *con't*




- Network devices that originate, route and terminate data are called '*nodes*'.
 - Nodes or '*hosts*' can be personal computers, phones, servers as well as special networking hardware.
- Two such devices can be said to be '*networked*' together when one device is able to exchange information with the other device.
- A '*router*' is an networking device that forwards data between networks in an internet.
- A '*link*' is the means of connecting one location to another for the purpose of transmitting and receiving data across networks.
 - phone lines, fiberoptic cables, bluetooth, ethernet, wireless, etc...

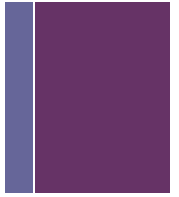
+ Logical Structure of an internet



- **Ad hoc interconnection of networks**
 - No particular topology
 - Vastly different router & link capacities
- **Send packets from source to destination by hopping through networks**
 - Router forms bridge from one network to another
 - Different *packets* may take different routes

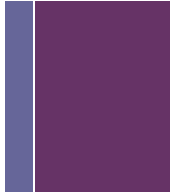


+ The Notion of an internet Protocol



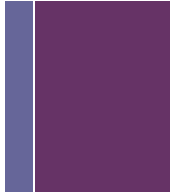
- **How is it possible to send bits across incompatible LANs and WANs?**
- **Solution: *protocol* software running on each host and router**
 - A protocol is a set of rules that governs how hosts and routers should cooperate when they transfer data from network to network.
 - The rules define the syntax, semantics and synchronization of communication and possible error recovery methods.
 - Protocols may be implemented by hardware, software, or a combination of both.

+ What Does an internet Protocol Do?



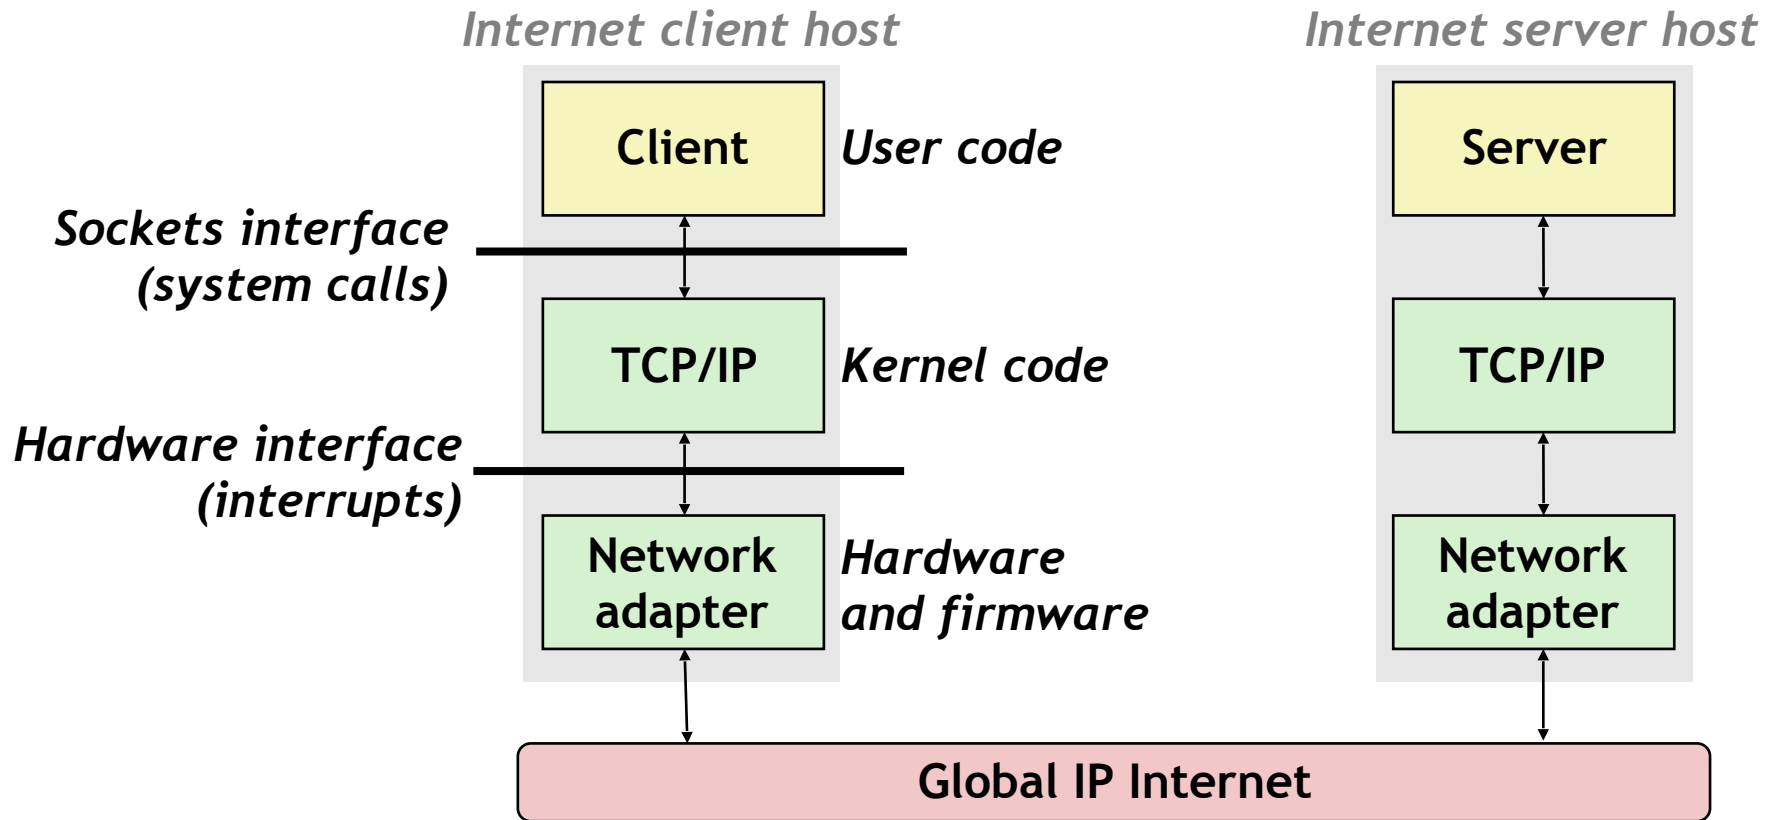
- **Provides a naming scheme**
 - An internet protocol defines a uniform format for *host addresses*
 - Each host (and router) is assigned at least one of these internet addresses that uniquely identifies it
- **Provides a delivery mechanism**
 - An internet protocol defines a standard transfer unit (*packet*)
 - Packet consists of *header* and *payload*
 - *Header*: contains info such as packet size, source and destination addresses
 - *Payload*: contains data bits sent from source host
- **Lots of other things...**
 - Example, what order do the bytes go on the wire? Little-endian or Big-endian?

+ Global IP Internet (upper case)



- **Most famous example of an internet**
- **Based on the TCP/IP protocol family**
 - *IP (Internet Protocol) :*
 - Provides basic *naming scheme* and *unreliable delivery capability* of packets (*datagrams*) from *host-to-host*
 - *UDP (Unreliable Datagram Protocol)*
 - Uses IP to provide *unreliable* datagram delivery from *process-to-process*
 - *TCP (Transmission Control Protocol)*
 - Uses IP to provide *reliable* byte streams from *process-to-process*
- **Accessed via a mix of Unix file I/O and functions from the *sockets interface***

+ Organization of an Internet Application



+ A Programmer's View of the Internet



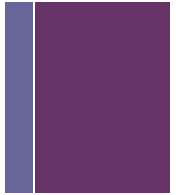
- Hosts are mapped to a set of 32-bit *IP addresses*
 - 128.122.49.30
- The set of IP addresses is mapped to a set of identifiers called *Internet domain names*
 - 128.122.49.30 is mapped to `cs.nyu.edu`
- A process on one Internet host can communicate with a process on another Internet host over a *connection*

+ Domain Naming System (DNS)



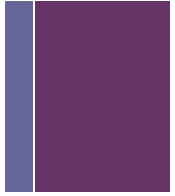
- **The Internet maintains a mapping between IP addresses and domain names in a huge worldwide distributed database called DNS**
- **Conceptually, programmers can view the DNS database as a collection of millions of host entries.**
 - Each host entry defines the mapping between a set of domain names and IP addresses.
- **These names can be resolved on the command line using `nslookup` command.**
- **Many names can be mapped to the same IP or multiple IPs mapped to same name.**
 - Why? Geolocation.

+ Internet Connections



- **Clients and servers communicate by sending streams of bytes over *connections*. Each connection is:**
 - *Point-to-point*: connects a pair of processes.
 - *Full-duplex*: data can flow in both directions at the same time,
 - *Reliable*: stream of bytes sent by the source is eventually received by the destination in the same order it was sent.
- **A *port* is a 16-bit integer that identifies a *process*:**
 - *Ephemeral port*: Assigned automatically by client kernel when client makes a connection request. (Also used by server in some cases)
 - *Well-known port*: Associated with some service provided by a server (e.g., port 80 is associated with Web servers)
- **A *socket* is an endpoint of a connection**
 - Socket address is an `IPAddress:port` pair

+ Well-known Ports and Service Names

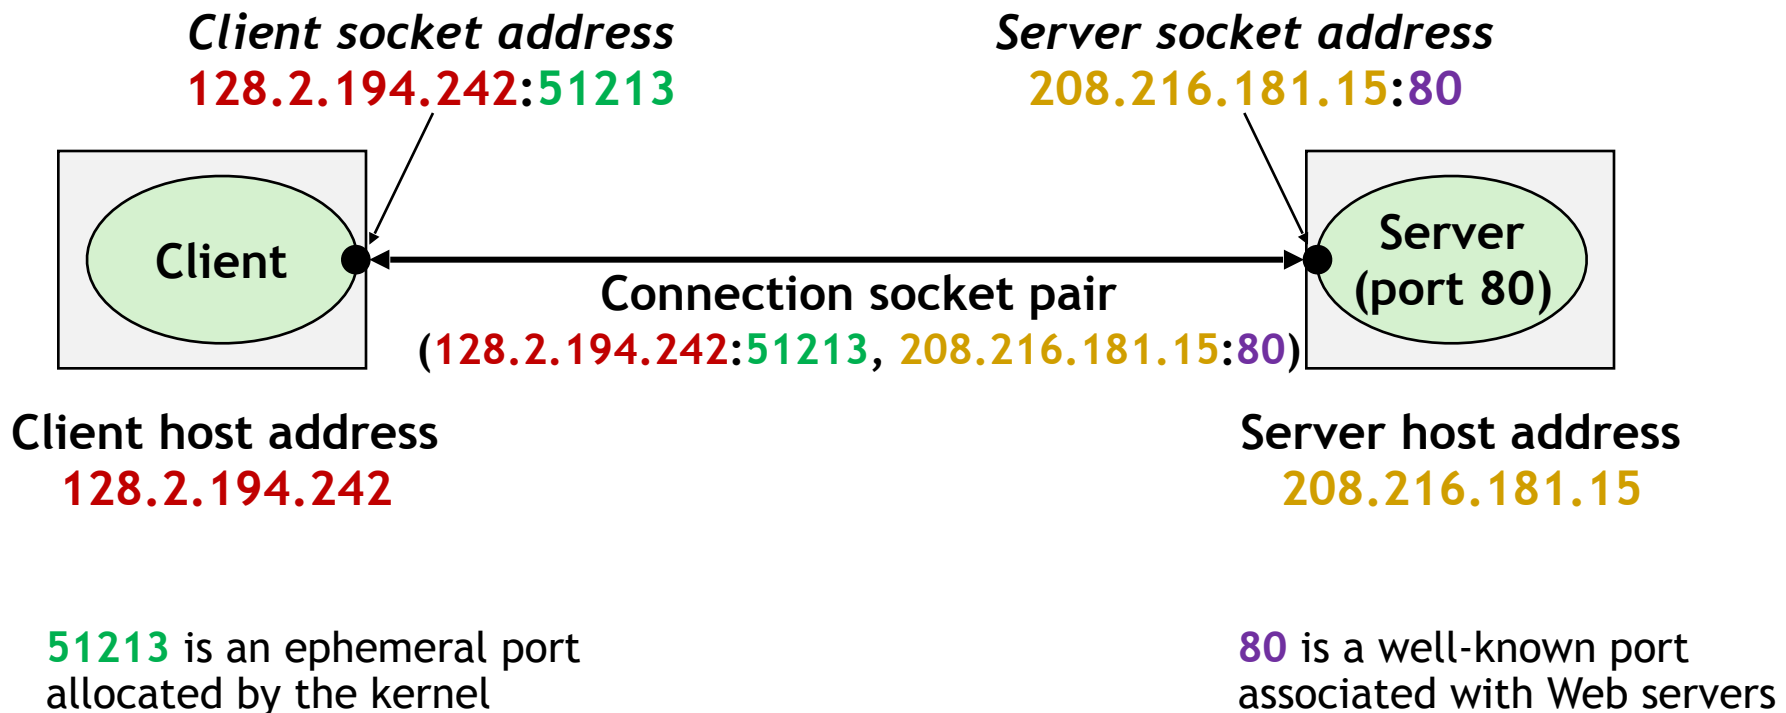


- Popular services have permanently assigned well-known ports and corresponding well-known service names:
 - Examples
 - echo server: 7/echo
 - ssh servers: 22/ssh
 - email server: 25/smtp
 - web servers: 80/http
 - 0-1023 have special semantics
 - https://en.wikipedia.org/wiki/List_of_TCP_and_UDP_port_numbers
- Mappings between well-known ports and service names is contained in the file /etc/services on each Linux machine.

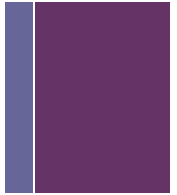
+ Anatomy of a Connection



- A connection is uniquely identified by the socket addresses of its endpoints (socket pair)
 - `(cliaddr:cliport, servaddr:servport)`



+ Sockets



- **What is a socket?**
 - To the kernel, a socket is an endpoint of communication
 - To an application, a socket is a file descriptor that lets the application read/write from/to the network
 - All Unix I/O devices, including networks, are modeled as files
- **Clients and servers communicate with each other by reading from and writing to socket descriptors**



- **The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors**

+ Socket Address Structures



- **Generic socket address:**

- For address arguments to `connect`, `bind`, and `accept`
- Necessary only because C did not have generic (`void*`) pointers when the sockets interface was designed

```
struct sockaddr {  
    uint16_t    sa_family;    /* Protocol family */  
    char        sa_data[14]; /* Address data. */  
};
```

sa_family



Family Specific

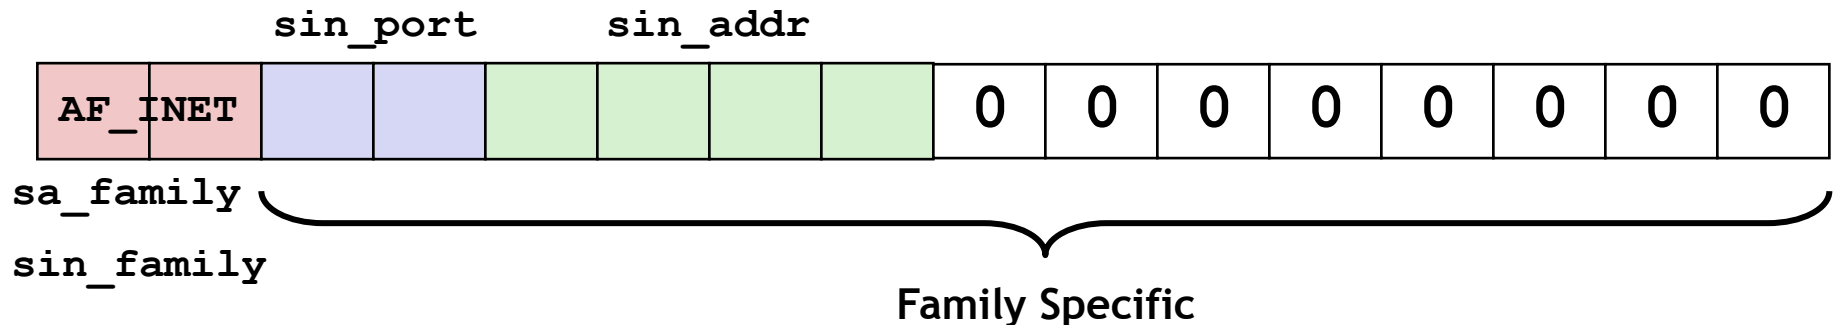
+ Socket Address Structures *con't*



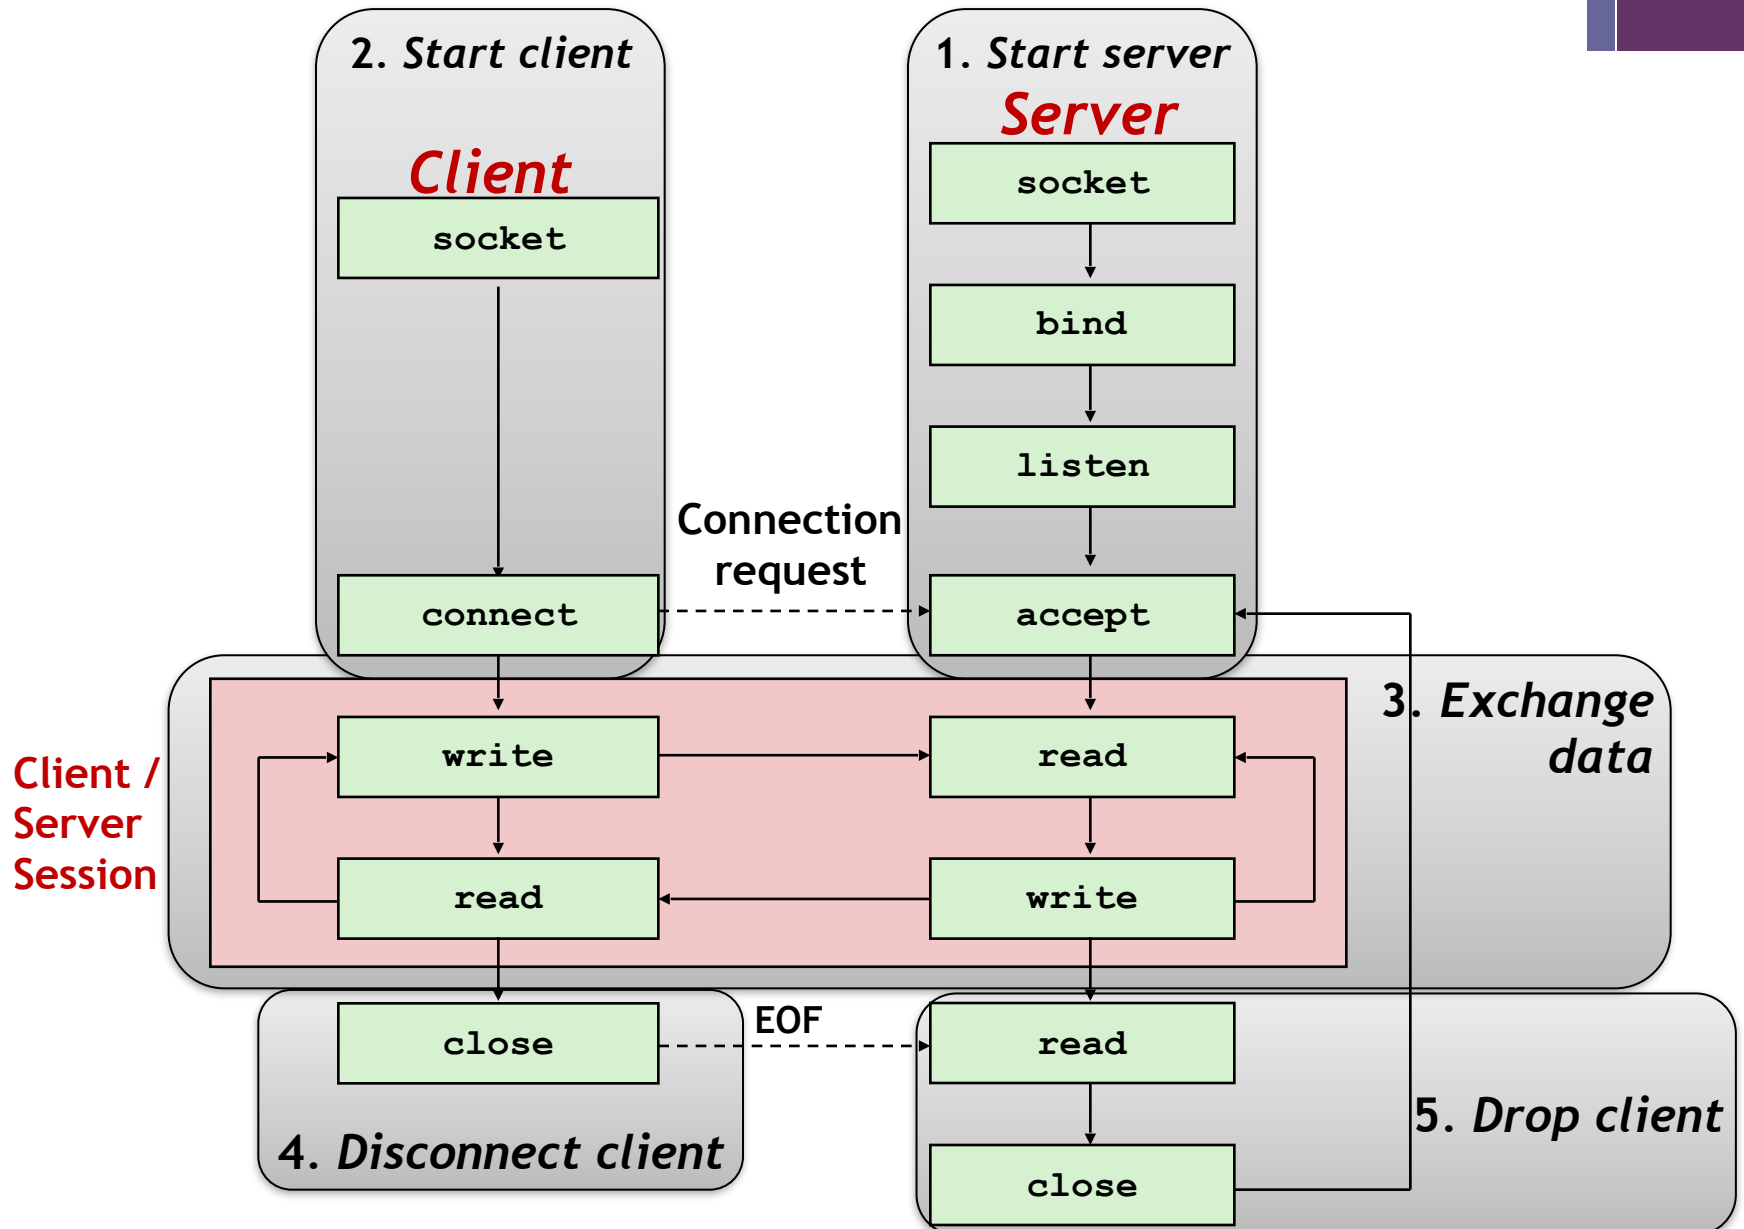
- Internet-specific socket address:

- Must cast `struct sockaddr_in*` to `struct sockaddr*` for functions that take socket address arguments.

```
struct sockaddr_in {  
    uint16_t      sin_family; /* Protocol family (always AF_INET) */  
    uint16_t      sin_port;   /* Port num in network byte order */  
    struct in_addr sin_addr;   /* IP addr in network byte order */  
    unsigned char  sin_zero[8]; /* Pad to sizeof(struct sockaddr) */  
};
```



+ Client/Server



+ Sockets Interface: `socket`



- Clients and servers use the `socket` function to create a socket descriptor:

```
int socket(int domain, int type, int protocol)
```


- Example:

```
int clientfd = Socket(AF_INET, SOCK_STREAM, 0);
```

Indicates that we are
using 32-bit IPV4
addresses

Indicates that the socket
will be the end point of a
TCP connection

+ Sockets Interface: `bind`

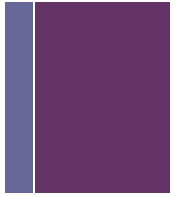


- A server uses `bind` to ask the kernel to associate the server's socket address with a socket descriptor:

```
int bind(int sockfd, sockaddr* addr, socklen_t addrlen);
```

- The process can read bytes that arrive on the connection whose endpoint is `addr` by reading from descriptor `sockfd`.
- Similarly, writes to `sockfd` are transferred along connection whose endpoint is `addr`.

+ Sockets Interface: `listen`

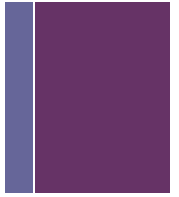


- By default, kernel assumes that descriptor from `socket` function is an active socket that will be on the client end of a connection.
- A server calls `listen` to tell the kernel that a descriptor will be used by a server rather than a client:

```
int listen(int sockfd, int backlog);
```

- Converts `sockfd` from an active socket to a *listening* socket that can accept connection requests from clients.
- `backlog` is a hint about the number of outstanding connection requests that the kernel should queue up before starting to refuse requests.

+ Sockets Interface: `accept`

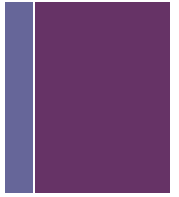


- Servers wait for connection requests from clients by calling `accept`:

```
int accept(int listenfd, sockaddr* addr, int* addrlen);
```

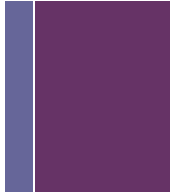
- Waits for connection request to arrive on the connection bound to `listenfd`, then fills in client's socket address in `addr` and size of the socket address in `addrlen`.
- Returns a connected descriptor that can be used to communicate with the client via Unix I/O routines.

+ Connected vs. Listening Descriptors



- **Listening descriptor**
 - End point for client connection requests
 - Created once and exists for lifetime of the server
- **Connected descriptor**
 - End point of the connection between client and server
 - A new descriptor is created each time the server accepts a connection request from a client
 - Exists only as long as it takes to service client
- **Why the distinction?**
 - Allows for concurrent servers that can communicate over many client connections simultaneously
 - E.g., Each time we receive a new request, we fork a child to handle the request

+ Sockets Interface: `connect`




- A client establishes a connection with a server by calling `connect`:


```
int connect(int clientfd, sockaddr* addr, socklen_t addrlen);
```

- Attempts to establish a connection with server at socket address `addr`
 - If successful, then `clientfd` is now ready for reading and writing.
 - `addrlen` is `sizeof(sockaddr_in)`


+ accept & connect Illustrated



1. Server blocks in `accept`, waiting for connection request on listening descriptor `listenfd`

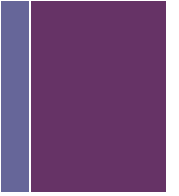


2. Client makes connection request by calling and blocking in `connect`



3. Server returns `connfd` from `accept`. Client returns from `connect`. Connection is now established between `clientfd` and `connfd`

+ Socket Client & Server



- Lets take a look at some code...
- See `lecture24/client.c` and `lecture24/server.c`

