



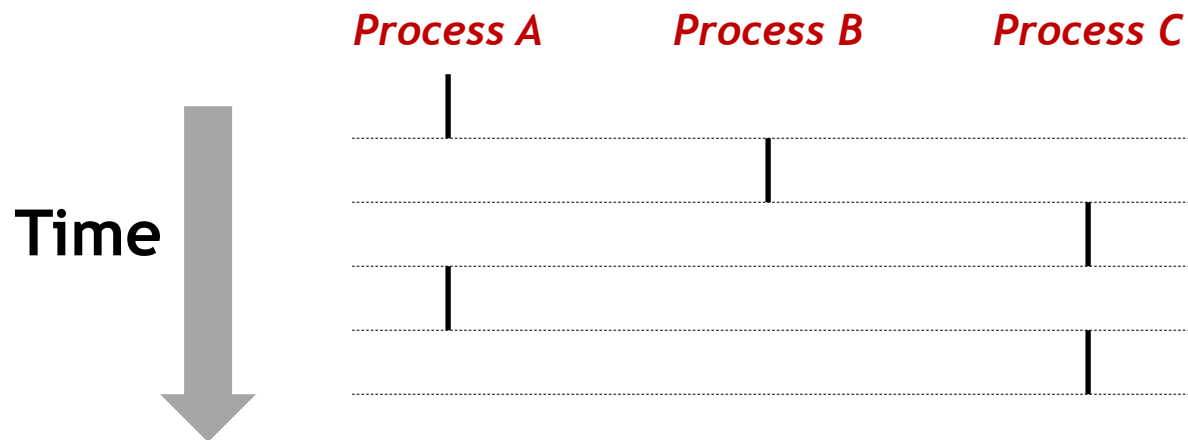
+

Concurrent Programming

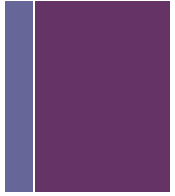
# + Concurrency (Review)



- Multiple logical control flows.
- Flows run concurrently if they overlap in time
  - Otherwise, they are sequential
- Examples (running on single core):
  - Concurrent: A & B, A & C
  - Sequential: B & C

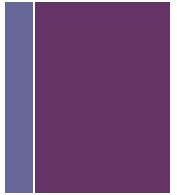


# + What & Why is Concurrency?



- **What: things happening “simultaneously”**
  - e.g. single CPU interleaving instructions from two flows
  - e.g. multiple CPU cores concurrently executing instructions
  - e.g. CPU and network card concurrently doing processing
- **Why: efficiency**
  - Due to ‘power wall’ cores not getting faster, just more numerous
    - To speed up programs using multiple CPUs we have to write concurrent code.
  - From systems perspective, don’t idle CPU while IO is performed
    - To speed up programs the system interleaves CPU processing and I/O.

# + Concurrent Programming is Hard!



- The human mind tends to be sequential
- Thinking about all possible sequences of events in a computer system is at least error prone and frequently impossible
  - Imagine two control flows of 2 instructions each A, B and C, D
    - Possible interleaved execution orders
      - A,B,C,D
      - A,C,B,D
      - A,C,D,B
      - C,D,A,B
      - C,A,D,B
      - C,A,B,D
  - Some orderings might yield unexpected results.

# + Approaches for Writing Concurrent Programs



- **Process-based**

- Kernel automatically interleaves multiple logical flows
- Each flow has its *own private address space*

- **Thread-based**

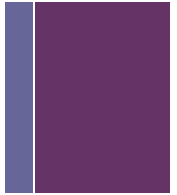
- Kernel automatically interleaves multiple logical flows
- Each flow *shares the same address space*
- Threads are less expensive for the system!



+

Process-based Concurrency


# + Process-Based Concurrent Program



- What does this program do?
- What would be printed from line 18?

```
1  int numbers[1000];
2  int sum1 = 0, sum2 = 0;
3
4  int main() {
5      for (int i = 0; i < 1000; i++)
6          numbers[i] = 1;
7
8      int pid = fork();
9      if (pid != 0) {
10         for (int i = 0; i < 500; i++)
11             sum1 += numbers[i];
12     } else {
13         for (int i = 0; i < 500; i++)
14             sum2 += numbers[500+i];
15         return 0;
16     }
17     waitpid(pid, NULL, 0);
18     printf("sum is %d\n", sum1 + sum2);
19
20     return 0;
21 }
```

# + Process-Based Concurrent Program




- What does this program do?
- What would be printed from line 18?
- Two processes concurrently sum the numbers.
- However, it is not simple to share data between them because they have *separate address spaces*.

```
1  int numbers[1000];
2  int sum1 = 0, sum2 = 0;
3
4  int main() {
5      for (int i = 0; i < 1000; i++)
6          numbers[i] = 1;
7
8      int pid = fork();
9      if (pid != 0) {
10         for (int i = 0; i < 500; i++)
11             sum1 += numbers[i];
12     } else {
13         for (int i = 0; i < 500; i++)
14             sum2 += numbers[500+i];
15         return 0;
16     }
17     waitpid(pid, NULL, 0);
18     printf("sum is %d\n", sum1 + sum2);
19
20     return 0;
21 }
```

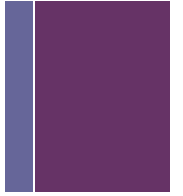


# + Interprocess Communication



- **How to communicate across processes? (*inter-process communication or IPC*)**
  - via *sockets*
  - via *pipes* (file system)
  - via *shared memory objects*

# + Unix Pipes



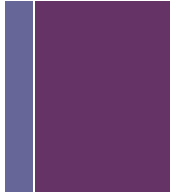
- Unlike other forms of interprocess communication, a pipe is one-way communication only
- Via a pipe, output of one process is the input to another process.
- A limitation of pipes is that the processes using pipes must have a common parent process

# + Pipes in C

- The `pipe` system call is called with a pointer to an array of two integers.
- The first element of the array contains the file descriptor that corresponds to the output of the pipe
- The second element of the array contains the file descriptor that corresponds to the input of the pipe.

```
1  int main()
2  {
3      int fd[2];
4      pipe(fd);
5
6      int pid = fork();
7      if (pid != 0) { // parent
8          write(fd[1], "This is a message!", 18);
9      }
10     else // child
11     {
12         int n;
13         char buf[1025];
14         if ((n = read(fd[0], buf, 1024)) >= 0)
15         {
16             buf[n] = 0; // null terminate string
17             printf("Child -> %s \n", buf);
18         }
19         return 0;
20     }
21
22     waitpid(pid, NULL, 0);
23     return 0;
24 }
```

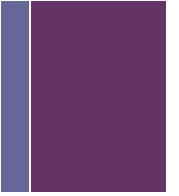
# + Pipes from the Shell



- Using the terminal you can **two commands together** so that the output from one program becomes the input of the next program.
- When you pipe commands together in the terminal in this way, it is called a *pipeline*
- **Example...**

```
ls | grep ".c" | sort -r | cut -c 1-5
```

# + Pros & Cons of Pipes



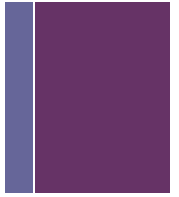
- **Pros**

- Efficient use of memory and CPU time
- Easy to create

- **Cons**


- Can be confusing quickly in non-trivial programs
- Uni-directional
- Little capability for error handling
- Requires system call

# + Shared Memory



- Shared Memory is an efficient means of passing data between programs.
- Allow two or more *processes* access to the same address space for reading and writing.
- A process creates or accesses a shared memory segment using `shmget()`
- You can review an example of two processes sharing some memory in the following files
  - `lecture25/shared_memory_server.c`
  - `lecture25/shared_memory_client.c`

# + Pros & Cons of Shared Memory



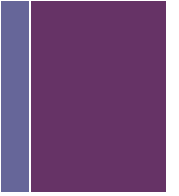
- **Pros**

- Highly performant, bidirectional communication

- **Cons**

- Error prone, difficult to debug
- Requires system call
- All the same *synchronization* problems as threads (which we will understand soon!)

# + Pros & Cons of Process-based Concurrency



- **Pros**

- Clean sharing model
  - descriptors (no)
  - file tables (yes)
  - global variables (no)

- **Cons**

- Systems calls necessary
- Nontrivial to share data between processes
  - Requires IPC (interprocess communication) mechanisms

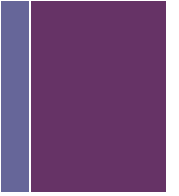




+

Thread-based Concurrency

# + What is a Thread?




- **A thread is...**
  - a unit of execution, associated with a process.
  - the smallest sequence of instructions that can be managed independently by the OS scheduler
- **Multiple threads can..**
  - exist within one process
  - be executing concurrently
  - *share resources* such as memory


# + Traditional View of a Process

- **Process = process context + code, data & stack**

## Process context



## Code, data, and stack

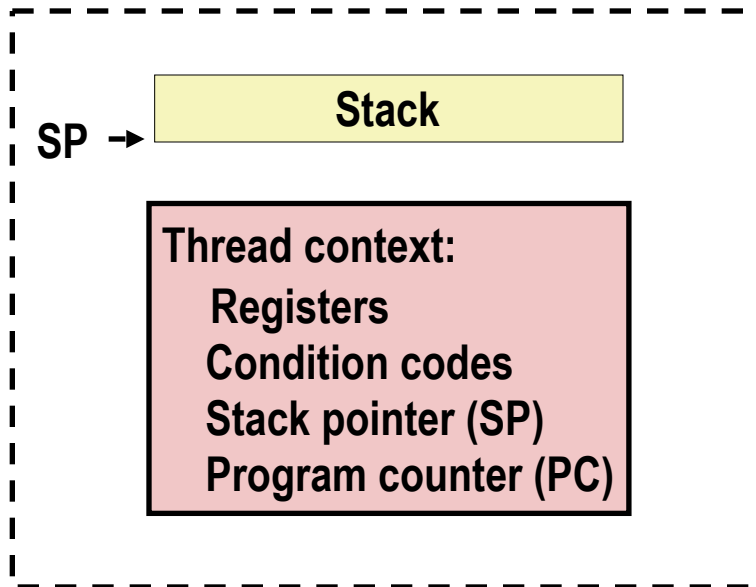


# + Alternate View of a Process

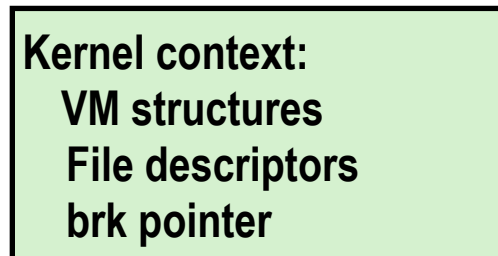
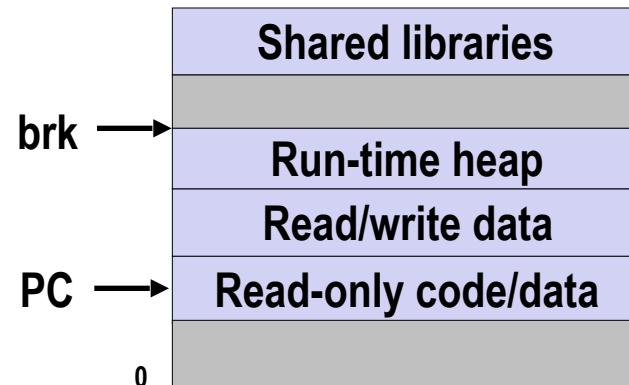


- **Process = thread context(s) + code, data & kernel context**

## Thread context



## Code, data, and kernel context



# + A Process With Multiple Threads



- **Multiple threads can be associated with a process**
  - Each thread has its own logical control flow
  - Each thread shares the same code, data, and kernel context
  - Each thread has its own stack for local variables
  - Each thread has its own thread id (TID)

## Thread 1 (main thread)

## Thread 2 (peer thread)


## Shared code and data

**stack 1**

**stack 2**

Thread 1 context:  
Data registers  
Condition codes  
SP1  
PC1

Thread 2 context:  
Data registers  
Condition codes  
SP2  
PC2




Kernel context:  
VM structures  
Descriptor table  
brk pointer


# + Logical View of Threads

- Threads associated with process form a pool of peers
  - Unlike processes which form a tree hierarchy

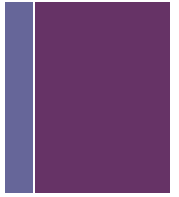
Threads associated with some process




Process hierarchy



# + Concurrent Threads



- Two threads are concurrent if their flows overlap in time
- Otherwise, they are sequential
- Examples:
  - Concurrent: A & B, A&C
  - Sequential: B & C

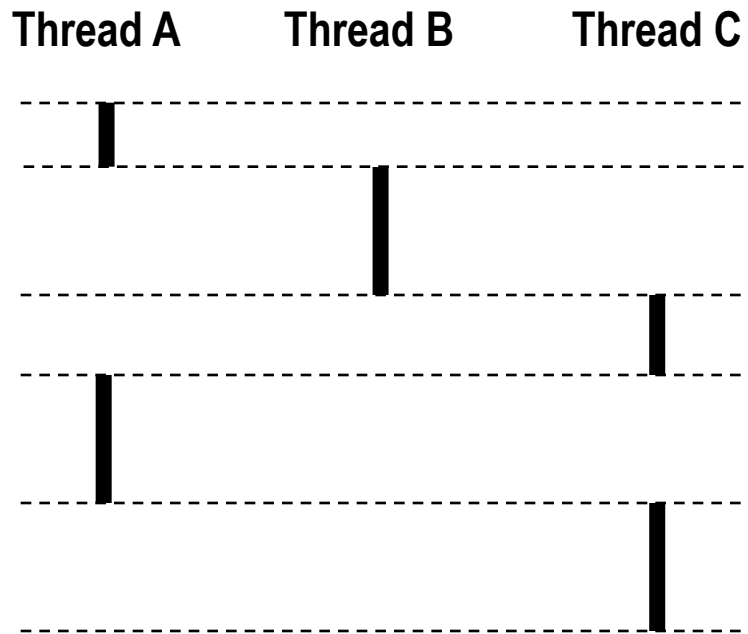


# + Concurrent vs Parallel Thread Execution



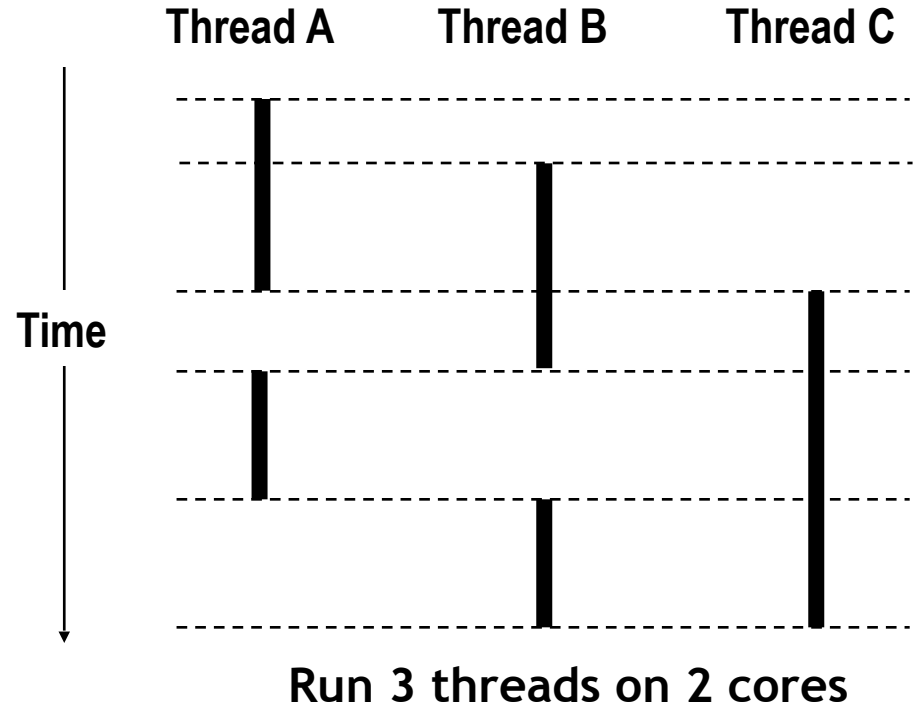
- **Single Core Processor**

- Simulate parallelism by time slicing



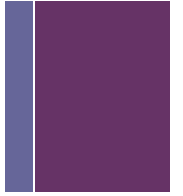
- **Multi Core Processor**

- Can have true parallelism





# + Threads vs. Processes




- **Similarities**

- Each has its own logical control flow
- Each can run concurrently (possibly on different cores)
- Each is context switched

- **Differences**

- Threads share all code and data (except local stacks)
  - Processes do not
- Threads are somewhat less expensive than processes
  - Process control (creating/reaping) **2x** as expensive as thread control

# + Posix Threads (Pthreads) Interface



- **Pthreads: Standard interface for ~60 functions that manipulate threads from C programs**
  - Creating and reaping threads
    - `pthread_create()`
    - `pthread_join()`
  - Determining your thread ID
    - `pthread_self()`
  - Terminating threads
    - `pthread_cancel()`
    - `pthread_exit()`
    - `exit()` [terminates all threads] , `RET` [terminates current thread]
  - Synchronizing access to shared variables
    - `pthread_mutex_init`
    - `pthread_mutex_[un]lock`

# + The Pthreads "hello, world" Program



```
/*
 * hello.c - pthreads "hello, world" program
 */
void *thread(void *vargp);

int main()
{
    pthread_t tid;
    pthread_create(&tid, NULL, thread, NULL);
    pthread_join(tid, NULL);
    exit(0);
}
```

Thread ID

Thread attributes  
(usually NULL)

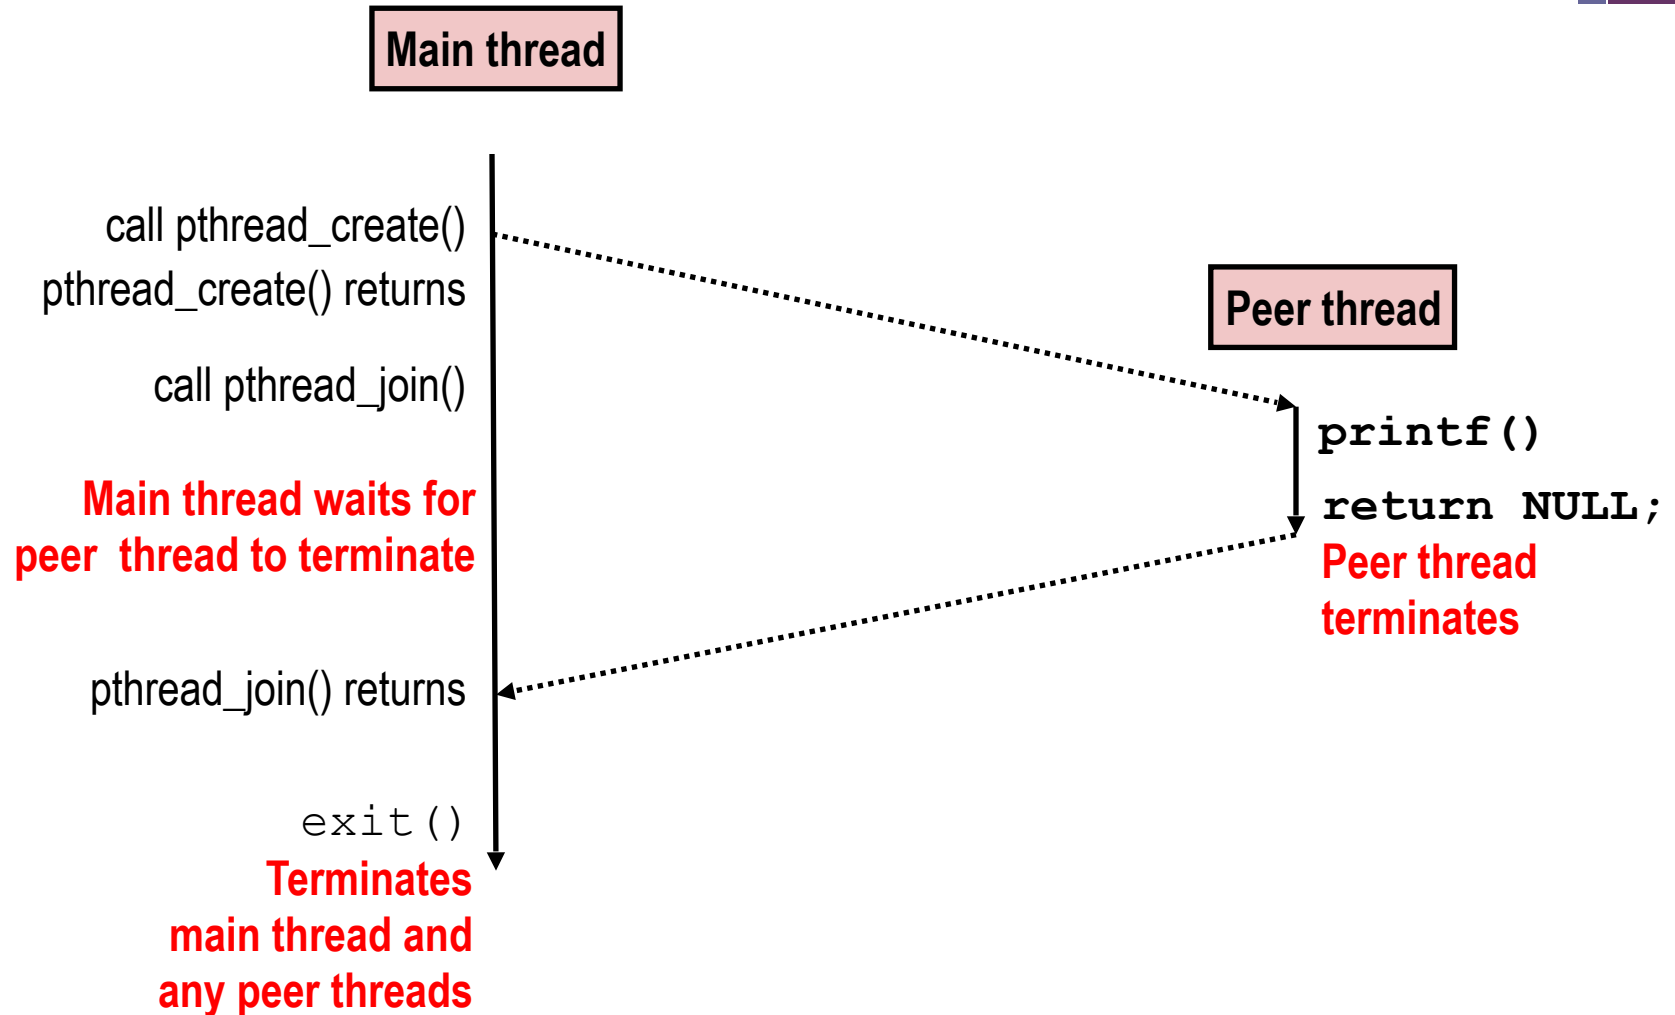
Thread routine

Thread arguments  
(void \*p)

```
void *thread(void *vargp) /* thread routine */
{
    printf("Hello, world!\n");
    return NULL;
}
```

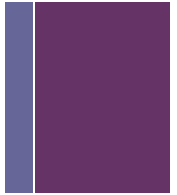
Return value  
(void \*\*p)

# + Execution of Threaded “hello, world”



- See `lecture25/thread_sum.c` for another example.

# + Concurrent Programming is Hard!



- The ease with which threads share data and resources also makes them vulnerable to subtle and baffling errors.
- Classical problem classes of concurrent programs:
  - *Races*: outcome depends on arbitrary scheduling decisions elsewhere in the system
  - *Deadlock*: improper resource allocation prevents forward progress
  - *Livelock / Starvation / Fairness*: external events and/or system scheduling decisions can prevent sub-task progress

# + Race Example

- What's the expected output on line 11?
  - 2
- Possible output...
  - 1
- ???????

```
1  int numbers[2] = { 1, 1 };
2  int sum = 0;
3
4  int main() {
5      pthread_t tid;
6      pthread_create(&tid, NULL, run, numbers[1]);
7      for (int i = 0; i < 1; i++) {
8          sum += numbers[i];
9      }
10     pthread_join(tid, NULL);
11     printf("sum is %d\n", sum);
12 }
13
14 void* run(void* arg) {
15     int* numbers = (int*) arg;
16     for (int i = 0; i < 1; i++) {
17         sum += numbers[i];
18     }
19     return NULL;
20 }
```

# + Race Example *con't*



- Why can the outcome be 1? This line is the culprit.

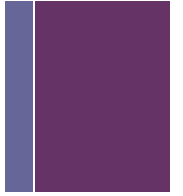
```
sum += numbers[i];
```

- What does this look like in assembly?

```
1  movq ...(%rdi,4), %rcx
2  movq ...(%rsi), %rdx
3  addq %rcx, %rdx
4  movq %rdx, ...(%rsi)
```

- Two threads T, T' have combinatorial number of interleavings
  - **OK:** T1, T2, T3, T4, T'1, T'2, T'3, T'4
  - **BAD:** T1, T'1, T2, T'2, T3, T'3, T4, T'4
- Global variable sum is written as 1 by both threads at T4 & T'4

# + Pros and Cons of Thread-Based Designs



- **Pros**

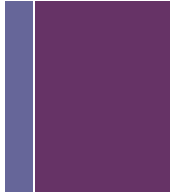
- Easy to share data structures between threads
- Threads are more efficient than processes

- **Cons**

- Unintentional sharing can introduce subtle and hard-to-reproduce errors.
  - Hard to detect by testing since probability of bad outcome can be low



# + Summary: Approaches to Concurrency



- **Process-based**

- Hard to share resources: easy to avoid unintended sharing
- High overhead in adding/removing clients

- **Thread-based**

- Easy to share resources: perhaps too easy
- Medium overhead
- Not much control over scheduling policies
- Difficult to debug - event orderings not repeatable