




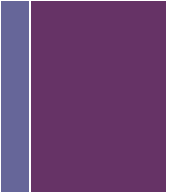
Final Exam

+ Logistics



- **Exam: 100 minutes (half an hour longer than midterm)**
 - Covers all materials, emphasizes second half
 - You may bring 1 page (8.5 by 11, two-sided) of notes. It must be hand-written.
- **Disclaimer: this review is not complete.**
 - Not all exam materials are mentioned by this review!
 - We are not even mentioning things from the first half of the semester.
 - This is a survey of topics and some of the core concepts.
 - There are about 55 slides.. so we will be doing on average a slide a minute!!

+ General Tips



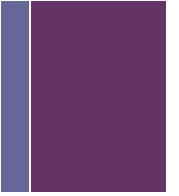
- Read through all questions, do what you think are the easier ones first.
- Make sure you have *completed and understood* all the homeworks, recitation exercises, labs and exams.
- Use Piazza! Have discussions there with your classmates.



+

Machine-Level Programming

+ Machine-Level Programs



- **Basic information about x84-64**
- **Source, assembly, object & machine code**
- **General purpose registers**
- **Data movement instructions.**
 - `movq %rdi, %rsi`
 - `movq (%rdi), %rsi`

- D(Rb, Ri, S)** **Mem[D + Reg[Rb] + Reg[Ri] * S]**

 - D: Constant “displacement”
 - Rb: Base register: Any of 16 integer registers
 - Ri: Index register: Any, except for **%rsp**
 - S: Scale: 1, 2, 4, or 8

▪ Special cases: you can omit certain arguments if not needed.

D(Rb,Ri) Mem[D + Reg[Rb] + Reg[Ri]]

(Rb,Ri,S) **Mem[Reg[Rb] + Reg[Ri] * S]**

+ Address Computation

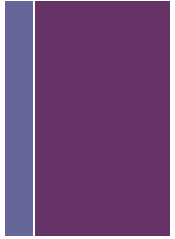
- `leaq src, dest`
 - `src` is an address computation expression
 - set `dest` to address denoted by expression
- use case 1
 - Computing addresses without a memory reference
 - E.g., translation of `p = &x[i];`
- Example

```
char* a2(char* x){  
    return &x[2];  
}
```

```
leaq 2(%rdi), %rax    # return &x[2]  
ret
```



+ Arithmetic Operations



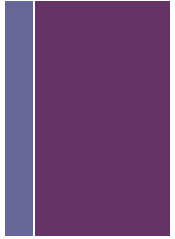
- Two Operand Instructions:

<u>Format</u>	<u>Computation</u>
▪ addq src, dest	dest = dest + src
▪ subq src, dest	dest = dest - src
▪ imulq src, dest	dest = dest * src
▪ salq src, dest	dest = dest << src (also called shlq)
▪ sarq src, dest	dest = dest >> src (arithmetic)
▪ shrq src, dest	dest = dest >> src (logical)
▪ xorg src, dest	dest = dest ^ src
▪ andq src, dest	dest = dest & src
▪ orq src, dest	dest = dest src

- Watch out for argument order!



Arithmetic Expression Example



```
long arith (long x, long y, long z){  
    long t1 = x + y;  
    long t2 = z + t1;  
    long t3 = x + 4;  
    long t4 = y * 48;  
    long t5 = t3 + t4;  
    long rval = t2 * t5;  
    return rval;  
}
```

```
arith:  
    leaq    (%rdi,%rsi), %rax    #t1  
    addq    %rdx, %rax          #t2  
    leaq    (%rsi,%rsi,2), %rdx  
    salq    $4, %rdx            #t4  
    leaq    4(%rdi,%rdx), %rcx   #t5  
    imulq   %rcx, %rax          #rval  
    ret
```

- Noteworthy instructions:
 - **leaq**: “address” computation
 - **salq**: shift
 - **imulq**: integer multiplication

Register	Use(s)
%rdi	Argument x
%rsi	Argument y
%rdx	Argument z
%rax	t1, t2, rval
%rdx	t4
%rcx	t5

+ Control

- Information about currently executing program...
 - temporary data
(`%rax`, ...)
 - location of runtime stack
(`%rsp`)
 - location of current code point
(`%rip`)
 - status of recent tests
(`CF`, `ZF`, `SF`, `OF`)

Registers

<code>%rax</code>	<code>%r8</code>
<code>%rbx</code>	<code>%r9</code>
<code>%rcx</code>	<code>%r10</code>
<code>%rdx</code>	<code>%r11</code>
<code>%rsi</code>	<code>%r12</code>
<code>%rdi</code>	<code>%r13</code>
<code>%rsp</code>	<code>%r14</code>
<code>%rbp</code>	<code>%r15</code>

`%rip`

Instruction

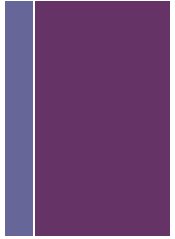
<code>CF</code>	<code>ZF</code>	<code>SF</code>	<code>OF</code>
-----------------	-----------------	-----------------	-----------------

Condition

Current stack 'top' Current instruction



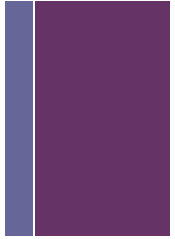
Condition Codes (Implicit Setting)



- **Single bit registers**
 - **CF** Carry Flag (for unsigned)
 - **SF** Sign Flag (for signed)
 - **ZF** Zero Flag
 - **OF** Overflow Flag (for signed)
- **Implicitly set (think of it as a side effect) by arithmetic operations**
 - Example: **addq *src, dest*** $\leftrightarrow b = a + b$
 - **CF** set if carry out from most significant bit (unsigned overflow)
 - **ZF** set if $t == 0$
 - **SF** set if $t < 0$ (as signed)
 - **OF** set if two's-complement (signed) overflow
 $(a > 0 \ \&\& \ b > 0 \ \&\& \ t < 0) \ || \ (a < 0 \ \&\& \ b < 0 \ \&\& \ t > 0)$
- **Not set by leaq instruction (!!!)**



Conditional Branching by Jumping



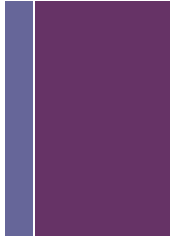
```
long absdiff(long x, long y)
{
    long result;
    if (x > y)
        result = x - y;
    else
        result = y - x;
    return result;
}
```

```
absdiff:
    cmpq    %rsi, %rdi # y, x
    jle     .L4
    movq    %rdi, %rax
    subq    %rsi, %rax
    ret
.L4:       # x <= y
    movq    %rsi, %rax
    subq    %rdi, %rax
    ret
```

Register	Use(s)
%rdi	Argument x
%rsi	Argument y
%rax	Return value



Conditional Move Example




```
long absdiff(long x, long y){
    long result;
    if (x > y)
        result = x-y;
    else
        result = y-x;
    return result;
}
```

Register	Use(s)
%rdi	Argument x
%rsi	Argument y
%rax	Return value
%rdx	Temp variable

```
absdiff:
    movq    %rdi, %rax    # x
    subq    %rsi, %rax    # if-val = x-y
    movq    %rsi, %rdx
    subq    %rdi, %rdx    # else-val = y-x
    cmpq    %rsi, %rdi    # %rsi = y, %rdi = x
    cmovle  %rdx, %rax    # if y <= x, result = else-val
    ret
```



“Do-While” Loop Compilation



```
long pcount_goto(unsigned long x){
    long result = 0;
loop:
    result += x & 0x1;
    x >>= 1;
    if(x) goto loop;
    return result;
}
```

Register	Use(s)
%rdi	Argument x
%rax	result

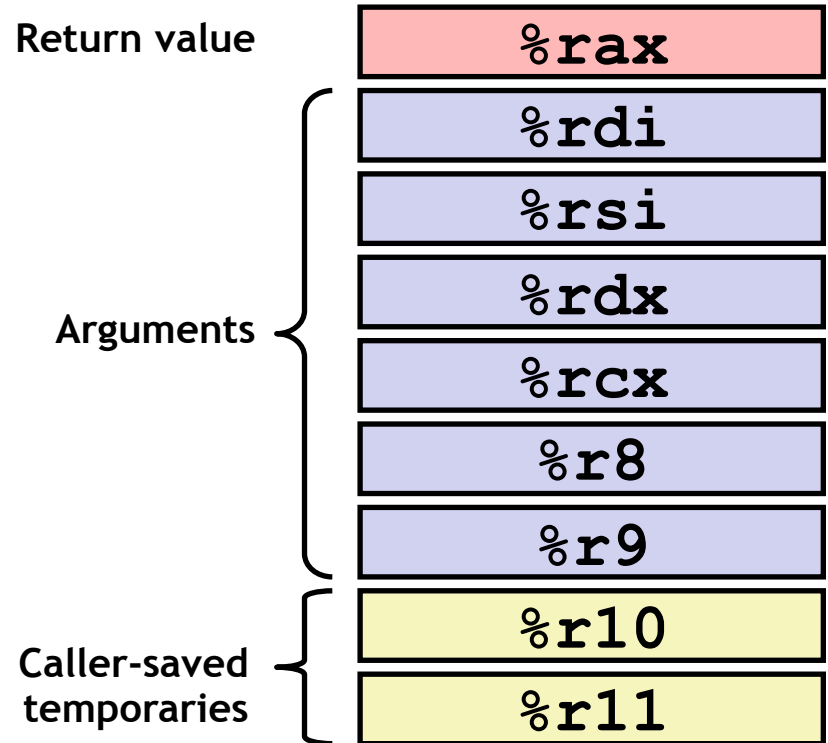
```
    movl    $0, %rax           # result = 0
.L2:                                     # loop:
    movq    %rdi, %rdx
    andq    $1, %rdx           # t = x & 0x1
    addq    %rdx, %rax          # result += t
    shrq    %rdi               # x >>= 1
    jne     .L2                # if (x) goto loop
    rep; ret
```

- Note: some processors' branch predictors behave badly when a branch's target or fall-through is a **ret** instruction, and adding the **rep;** prefix avoids this.

+ x86-64 Caller-saved Registers



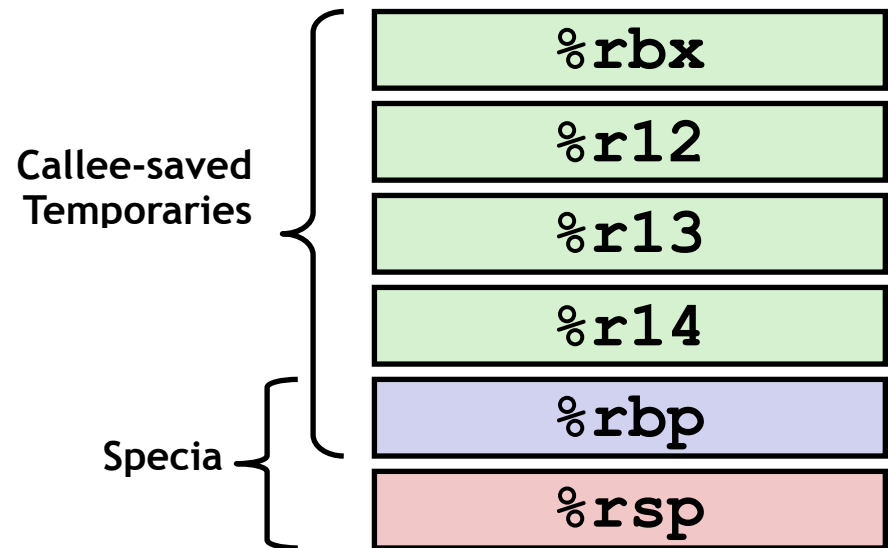
- **%rax**
 - Return value
 - Also caller-saved
 - Can be modified by procedure
- **%rdi, ..., %r9**
 - Arguments
 - Also caller-saved
 - Can be modified by procedure
- **%r10, %r11**
 - Caller-saved
 - Can be modified by procedure



+ x86-64 Callee-saved Registers



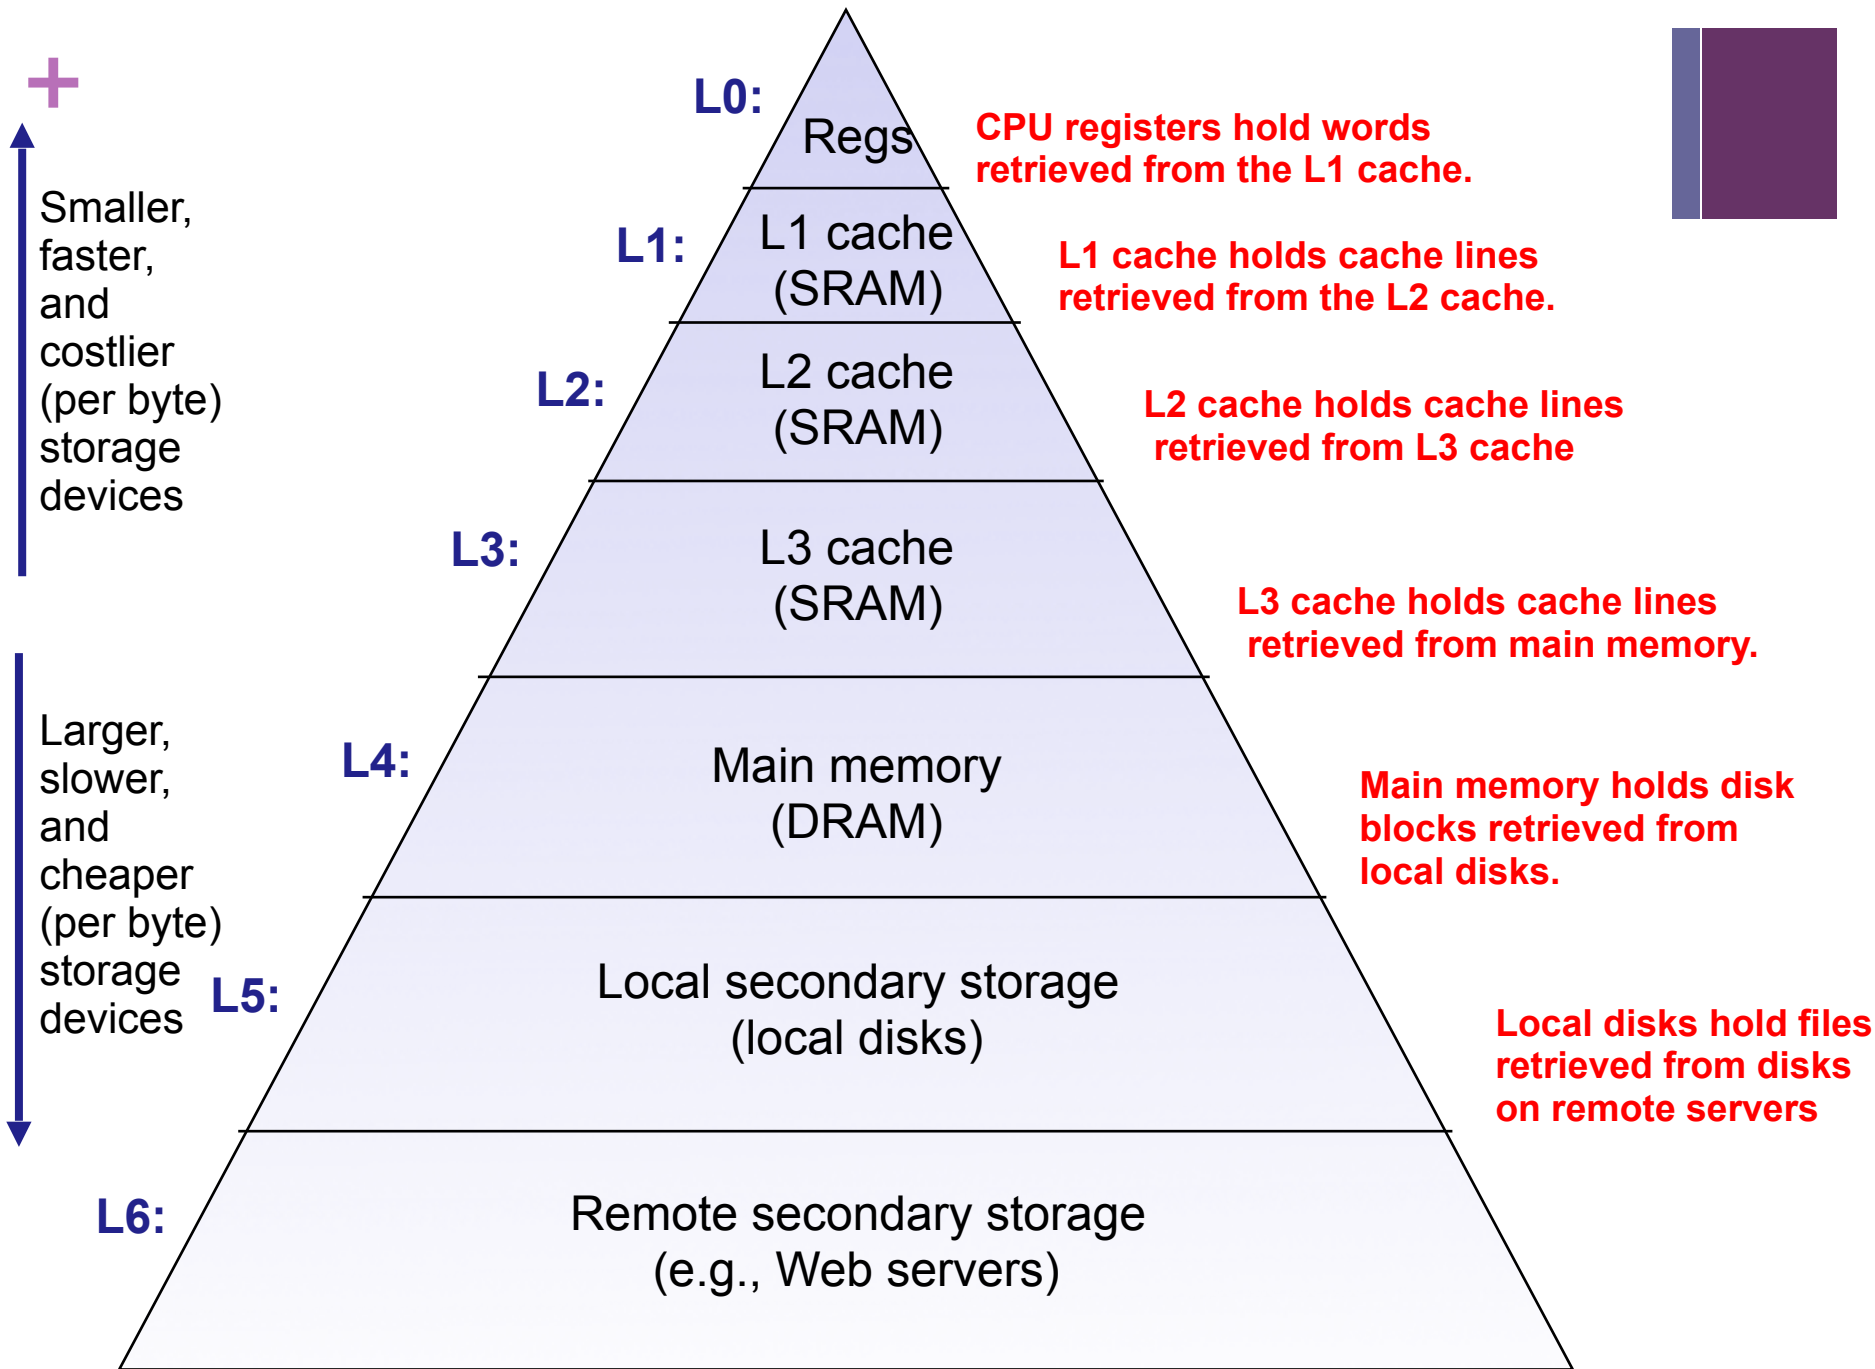
- **%rbx, %r12, %r13, %r14**
 - Callee-saved
 - Callee must save & restore
- **%rbp**
 - Callee-saved
 - Callee must save & restore
 - May be used as frame pointer
- **%rsp**
 - Special form of callee save
 - Restored to original value upon exit from procedure





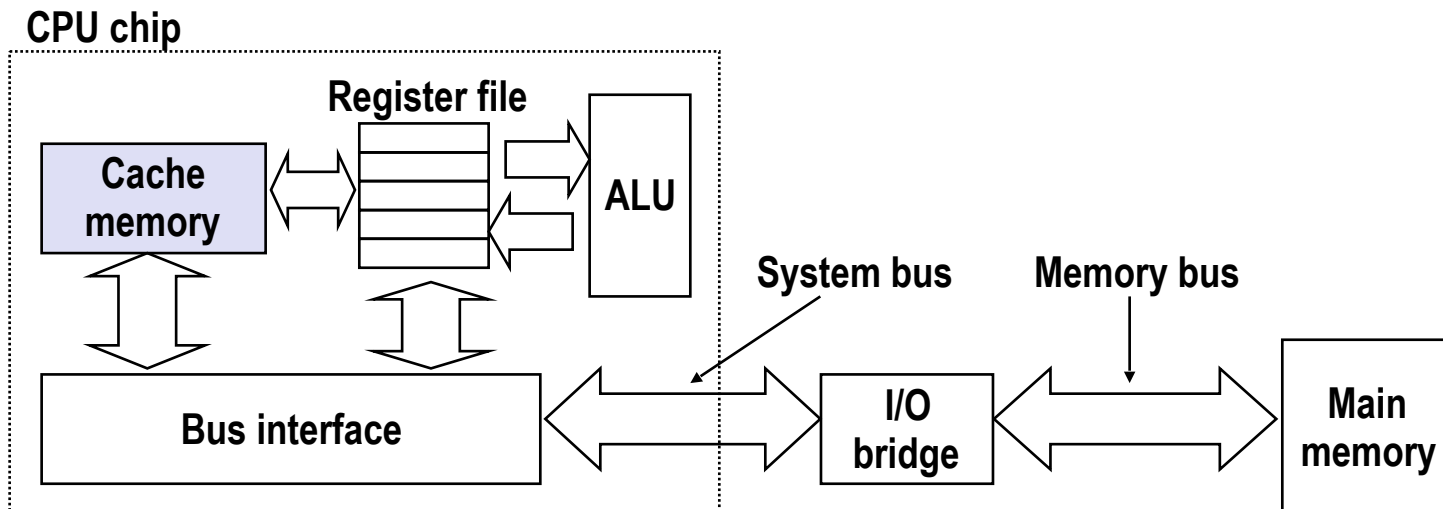
+

Memory



+ Cache Memories

- Cache memories are small, fast SRAM-based memories managed automatically in hardware
 - Hold frequently accessed blocks of main memory
- CPU looks first for data in cache
- Simplified system structure:



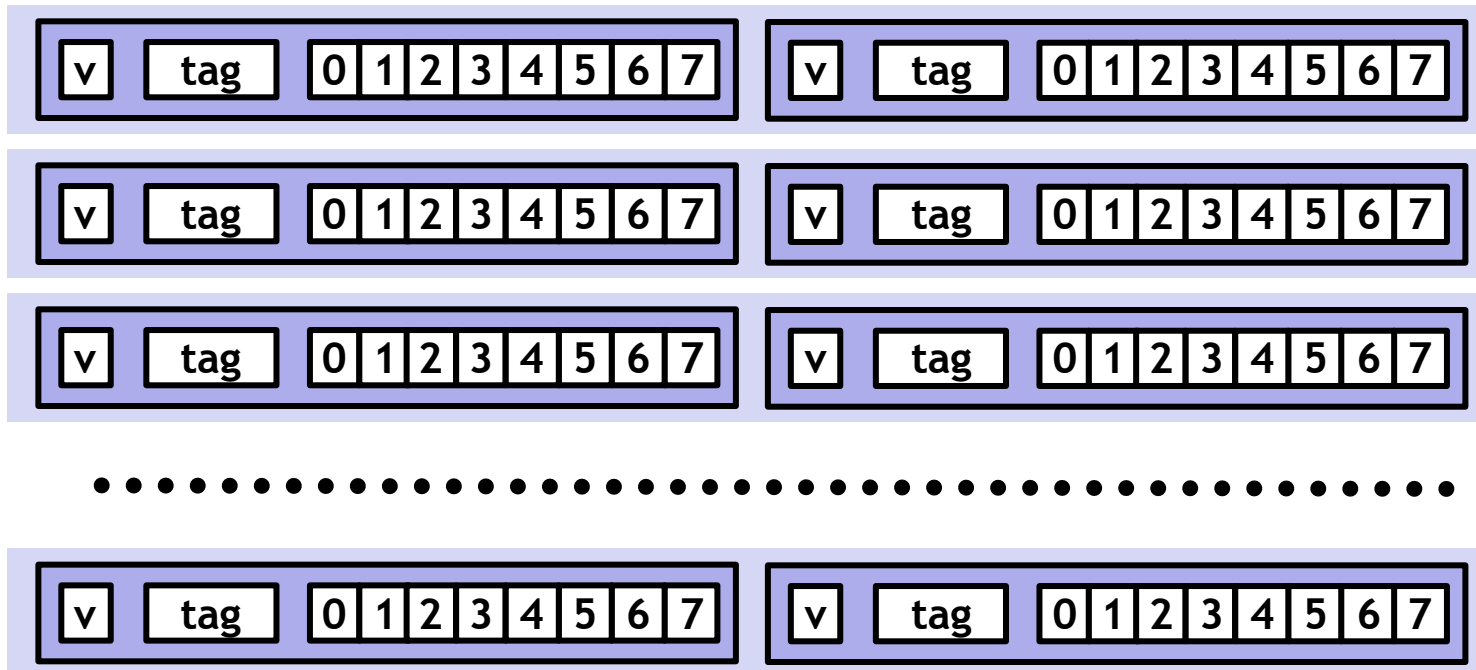
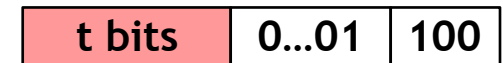
+ E-way Set Associative Cache (E = 2)



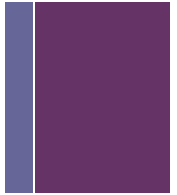
E = 2: Two lines per set

Assume: cache block size 8 bytes

Address of short int:



+ Cache Performance Metrics



- **Miss Rate**

- Fraction of memory references not found in cache (misses / accesses) = $1 - \text{hit rate}$
- Typical numbers:
 - 3-10% for L1
 - $< 1\%$ for L2, depending on size, etc.

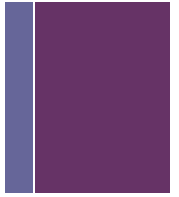
- **Hit Time**

- Time to deliver a line in the cache to the processor
 - includes time to determine whether the line is in the cache
- Typical numbers:
 - 4 clock cycle for L1
 - 10 clock cycles for L2

- **Miss Penalty**

- Additional time required because of a miss
 - typically 50-200 cycles for main memory


+ Locality



- **Principle of Locality:** Programs tend to use data and instructions with addresses near or equal to those they have used recently


- **Temporal locality:**

- Recently referenced items are likely to be referenced again in the near future

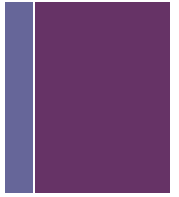


- **Spatial locality:**

- Items with nearby addresses tend to be referenced close together in time




+ Stride-n Access Patterns



- **C arrays allocated in row-major order, contiguously**
- **Stepping through columns in one row:**
 - `for (i = 0; i < N; i++)`
 `sum += a[0][i];`
 - accesses successive, contiguous elements
- **Stepping through rows in one column:**
 - `for (i = 0; i < n; i++)`
 `sum += a[i][0];`
 - accesses distant elements (stride-rowsize pattern)
 - no spatial locality!

+ Array Accessing Example

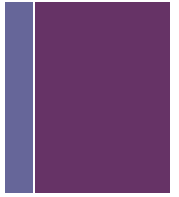


```
int get_digit(int[] z, int digit){  
    return z[digit];  
}
```


```
# %rdi = z  
# %rsi = digit  
movq (%rdi,%rsi,4), %rax # z[digit]
```

- Register `%rdi` contains starting address of array
- Register `%rsi` contains array index
- Desired digit at `%rdi + 4*%rsi`
- Use memory reference `(%rdi,%rsi,4)`

+ Structure Representation



```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};
```

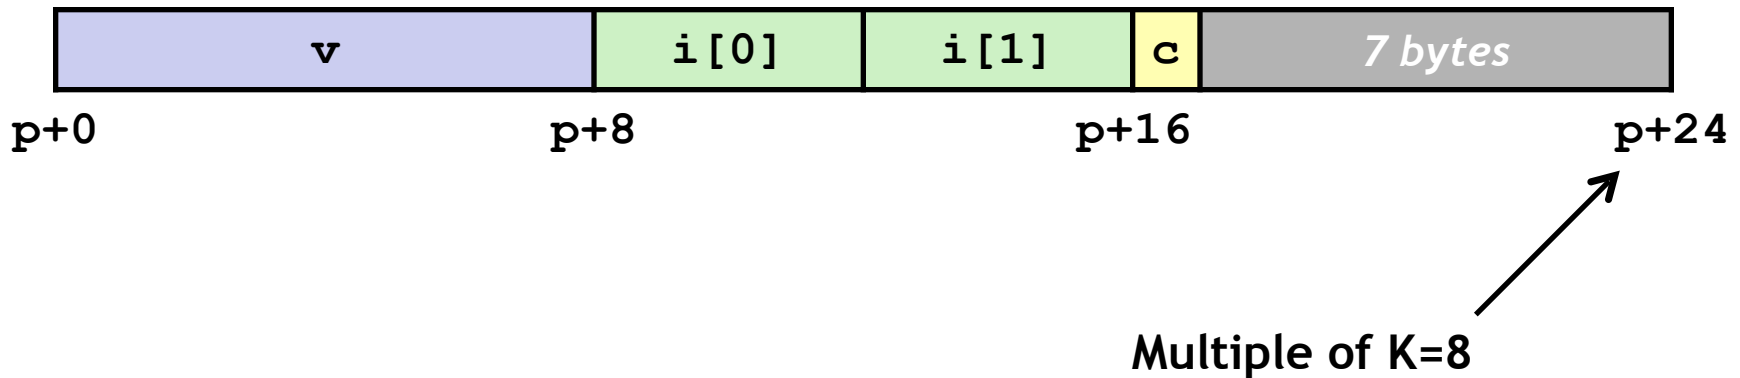


- **Structure represented as block of memory**
 - Big enough to hold all of the fields
- **Fields ordered according to declaration**
 - Even if another ordering could yield a more compact representation
 - Machine-level program has no understanding of the structures in the source code

+ Alignment Requirement

- For largest alignment requirement K
- Overall structure must be multiple of K


```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



+ Saving Space (Struct packing)



- Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

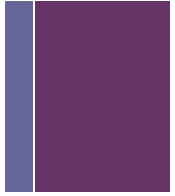
- Effect (K=12 -> K=4)






Exceptions & Processes

+ Control Flow




- **Processors do only one thing:**
 - From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
 - This sequence is the CPU's control flow (or flow of control)

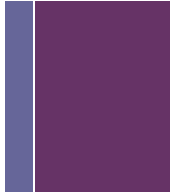


+ Exceptions

- An exception is a transfer of control to the OS kernel in response to some event (i.e., change in processor state)
 - Examples of events: Divide by 0, page fault, I/O request completes, typing Ctrl-C

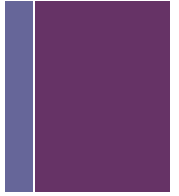


+ Asynchronous Exceptions (Interrupts)



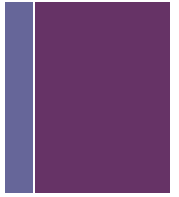
- **Caused by events external to the processor**
 - Indicated by setting the processor's *interrupt pin*
 - Handler returns to “next” instruction
- **Examples:**
 - Timer interrupt
 - Every few ms, an external timer chip triggers an interrupt
 - Used by the kernel to take back control from user programs
 - I/O interrupt from external device
 - Arrival of a packet from a network
 - Arrival of data from a disk

+ Synchronous Exceptions




- **Caused by events that occur as a result of executing an instruction:**
 - Traps
 - Intentional
 - Example: *system calls*
 - Returns control to “next” instruction
 - Faults
 - Unintentional but possibly recoverable
 - Example: *page fault*
 - Either re-executes faulting (“current”) instruction or aborts
 - Aborts
 - Unintentional and unrecoverable
 - Example: *illegal memory access*
 - Aborts current program

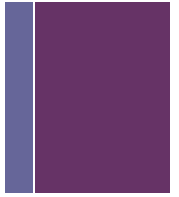
+ Processes



- A *process* is an instance of a running program.
 - One of the most successful ideas in computer science
 - Not the same as “program”
- Process provided with two key abstractions by OS:
 - *Logical control flow*
 - Each program seems to have exclusive use of the CPU
 - Provided by kernel mechanism called *context switching*
 - *Private address space*
 - Each program seems to have exclusive use of main memory.
 - Provided by kernel mechanism called *virtual memory*

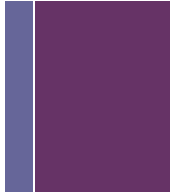


+ Creating and Terminating Processes



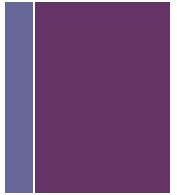
- From a programmer's perspective, we can think of a process as being in one of three states
- **Running**
 - Process is either executing, or waiting to be executed and will eventually be scheduled (i.e., chosen to execute) by the kernel
- **Stopped**
 - Process execution is suspended and will not be scheduled until further notice
- **Terminated**
 - Process is stopped permanently
 - Zombie!

+ Process Management




- **Creating processes**
 - Call `fork`
 - One call, *two* returns
- **Process completion**
 - Call `exit`
 - One call, no return
- **Reaping and waiting for processes**
 - Call `wait` or `waitpid`
- **Loading and running programs**
 - Call `execve`
 - One call, no return

+ Process Graph Example



```
int main() {  
    pid_t pid;  
    int x = 1;  
  
    pid = fork();  
    if (pid == 0) { /* Child */  
        printf("child : x=%d\n", ++x);  
        exit(0);  
    }  
  
    /* Parent */  
    printf("parent: x=%d\n", --x);  
    exit(0);  
}
```



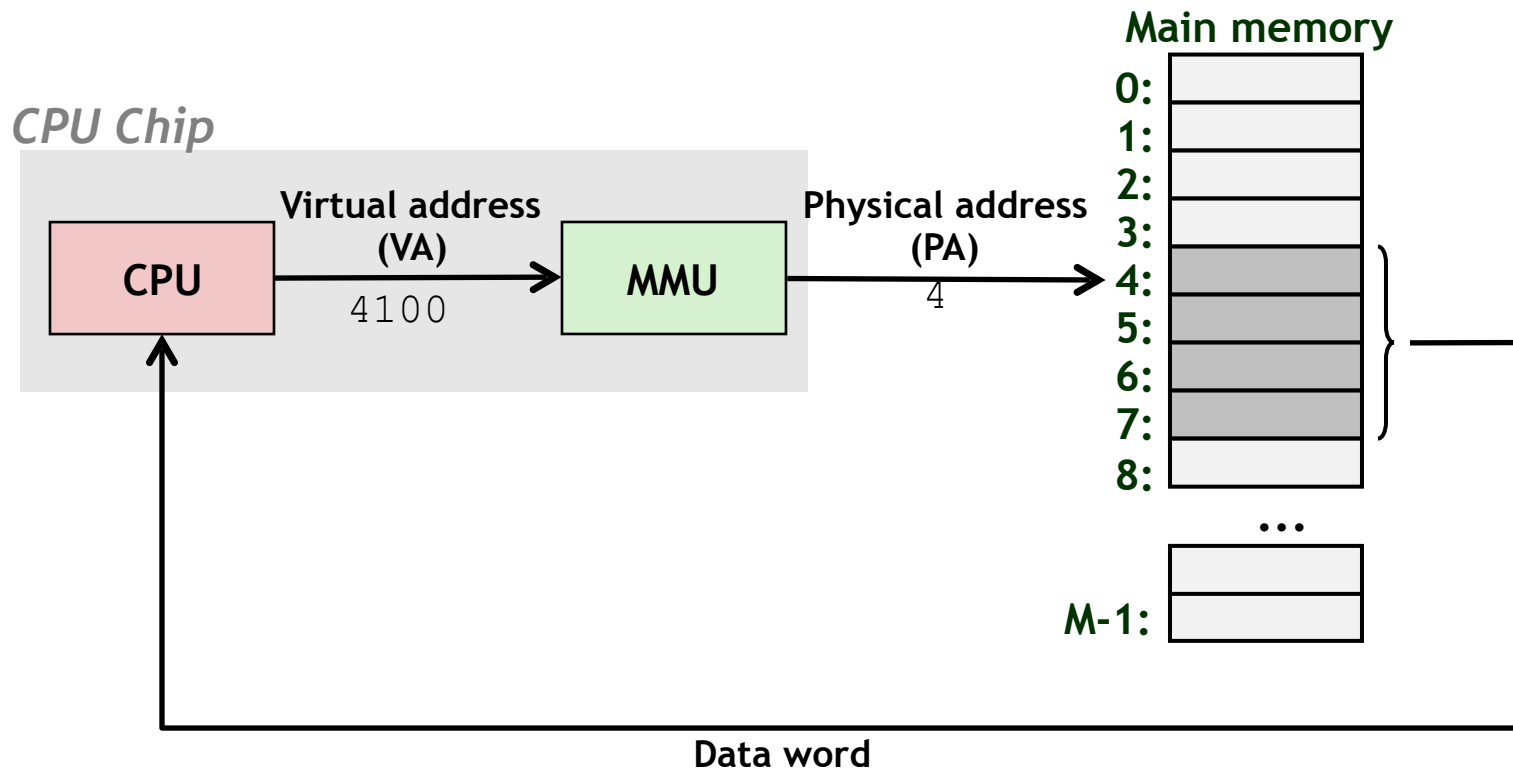


+

Virtual Memory

+ Virtual Addressing

- Creates illusion to process that it has total address space.
- Used in all modern, non-trivial systems



+ Why Virtual Memory (VM)?




- **Uses main memory efficiently**
 - Use DRAM as a cache for parts of a virtual address space
- **Simplifies memory management**
 - Each process gets the same uniform linear address space
- **Isolates address spaces**
 - One process can't interfere with another's memory
 - User program cannot access privileged kernel information and code

+ VM as Caching

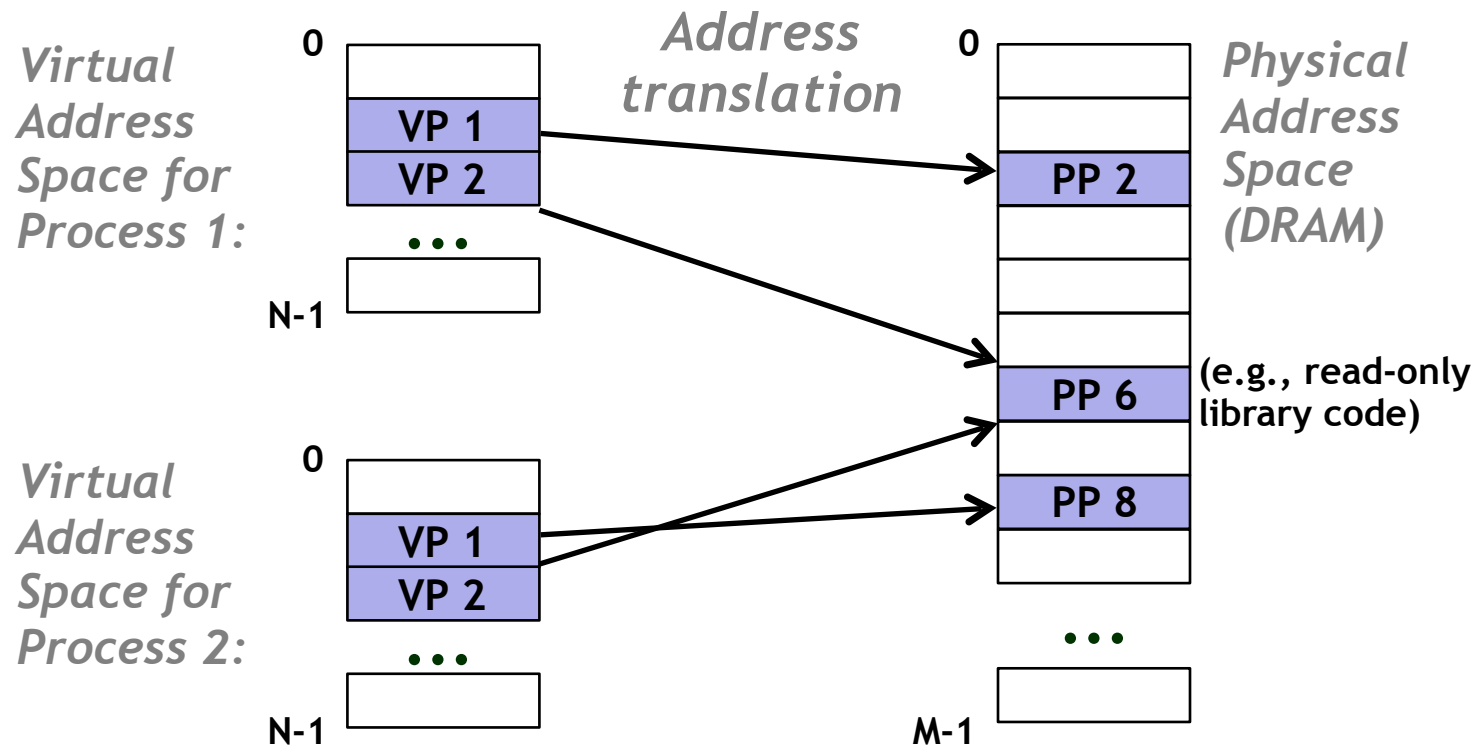


- Conceptually, virtual memory is an array of N contiguous bytes stored on disk.
- The contents of the array on disk are cached in physical memory
 - These cache blocks are called *pages*



+ VM as a Tool for Memory Management

- **Key idea: each process has its own virtual address space**
 - It can view memory as a simple linear array
 - Mapping function scatters addresses through physical memory



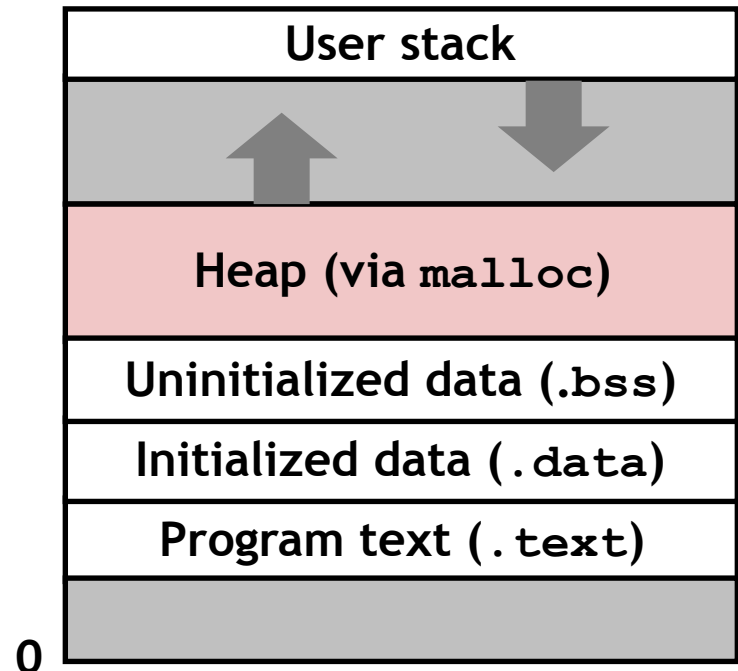



Dynamic Allocation

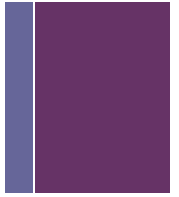
+ Dynamic Memory Allocators



- **Programmers use dynamic memory allocators to acquire virtual memory at run time.**
 - For data structures whose size is only known at runtime.
- **Dynamic memory allocators manage an area of process virtual memory known as the heap.**



+ Constraints




- **Applications**

- Can issue arbitrary sequence of `malloc` and `free` requests
- `free` request must be to a `malloc`'d block

- **Allocators**

- Can't control number or size of allocated blocks
- Can't reorder or buffer requests for memory
- Must allocate blocks from *free* memory
- Can manipulate and modify only *free* memory
- Can't move the allocated blocks once they are `malloc`'d

+ Performance Goals




- **Goal: Increase utilization (by reducing fragmentation)**
 - internal fragmentation
 - external fragmentation
- **Goal: Maximize throughput:**
 - Number of completed requests per unit time
 - Example:
 - 5,000 `malloc` calls and 5,000 `free` calls in 10 seconds
 - Throughput is 1,000 operations/second


+ Implicit Lists



- For each block we need both size and allocation status
 - Could store this information in two words: wasteful!
- Standard trick
 - If blocks are word-aligned, some low-order address bits are always 0
 - Instead of storing an always-0 bit, use it as a allocated/free flag
 - *(Note: when reading size, must mask out any extra bits)*



+ Detailed Implicit Free List Example



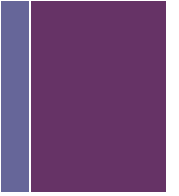
8-byte word
aligned

Allocated blocks: shaded

Free blocks: unshaded

Headers: labeled with size in
bytes/allocated bit

+ Ideas to remember...



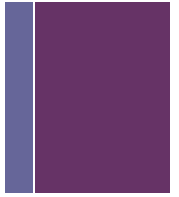
- **Allocation**
- **Freeing**
- **Fragmentation**
 - External
 - Internal
- **Coalescing**




+

Concurrent Programming

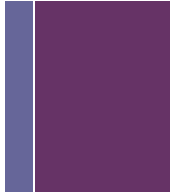
+ Concurrency



- Multiple logical control flows.
- Flows run concurrently if they overlap in time
 - Otherwise, they are sequential
- Examples (running on single core):
 - Concurrent: A & B, A & C
 - Sequential: B & C

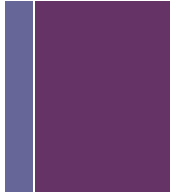


+ Process Concurrency



- Use fork to launch multiple processes to do work.
 - Hard for them to share data. Must use IPC.
- How to communicate across processes? (*inter-process communication or IPC*)
 - via *sockets*
 - via *pipes*
 - via *shared memory objects*

+ Thread-based Concurrency



- **A thread is...**
 - a unit of execution, associated with a process.
 - the smallest sequence of instructions that can be managed independently by the OS scheduler
- **Multiple threads can..**
 - exist within one process
 - be executing concurrently
 - *share resources* such as memory

+ Threads vs. Processes



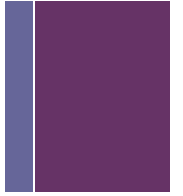
- **Similarities**

- Each has its own logical control flow
- Each can run concurrently (possibly on different cores)
- Each is context switched

- **Differences**

- Threads share all code and data (except local stacks)
 - Processes do not
- Threads are somewhat less expensive than processes
 - Process control (creating/reaping) **2x** as expensive as thread control

+ Concurrent Programming is Hard!



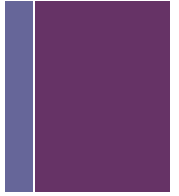
- The ease with which threads share data and resources also makes them vulnerable to subtle and baffling errors.
- Classical problem classes of concurrent programs:
 - *Races*: outcome depends on arbitrary scheduling decisions elsewhere in the system

+ Race Example

- What's the expected output on line 11?
 - 2
- Possible output...
 - 1

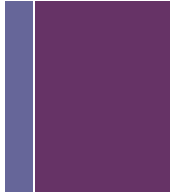
```
1  int numbers[2] = { 1, 1 };
2  int sum = 0;
3
4  int main() {
5      pthread_t tid;
6      pthread_create(&tid, NULL, run, numbers[1]);
7      for (int i = 0; i < 1; i++) {
8          sum += numbers[i];
9      }
10     pthread_join(tid, NULL);
11     printf("sum is %d\n", sum);
12 }
13
14 void* run(void* arg) {
15     int* numbers = (int*) arg;
16     for (int i = 0; i < 1; i++) {
17         sum += numbers[i];
18     }
19     return NULL;
20 }
```

+ Synchronizing Threads



- Shared variables are sometimes useful but they introduce the possibility of *synchronization* errors.
 - Like the one we saw last time
- How do we prevent such things?
 - We need to make sure that only one thread is executing “critical sections” (usually mutating shared variables!)
 - This is known as *mutual exclusion*
- Moreover, we must protect *critical sections*.
- **AGAIN, MULTIPLE THREADS WRITING TO THE EXACT SAME LOCATION IN MEMORY CONCURRENTLY IS A BAD THING!**

+ Mutual Exclusion



- A *mutex*...
 - is synchronization variable that is used to protect the access to shared variables.
 - surrounds critical sections so that one threads is allowed inside at a time.
- **However we have to be careful with these because they can lead to code that performs poorly, or in some cases not at all...**